

Parallel Rendering and Large Data Visualization

PhD Defense Talk
Stefan Eilemann

Outline

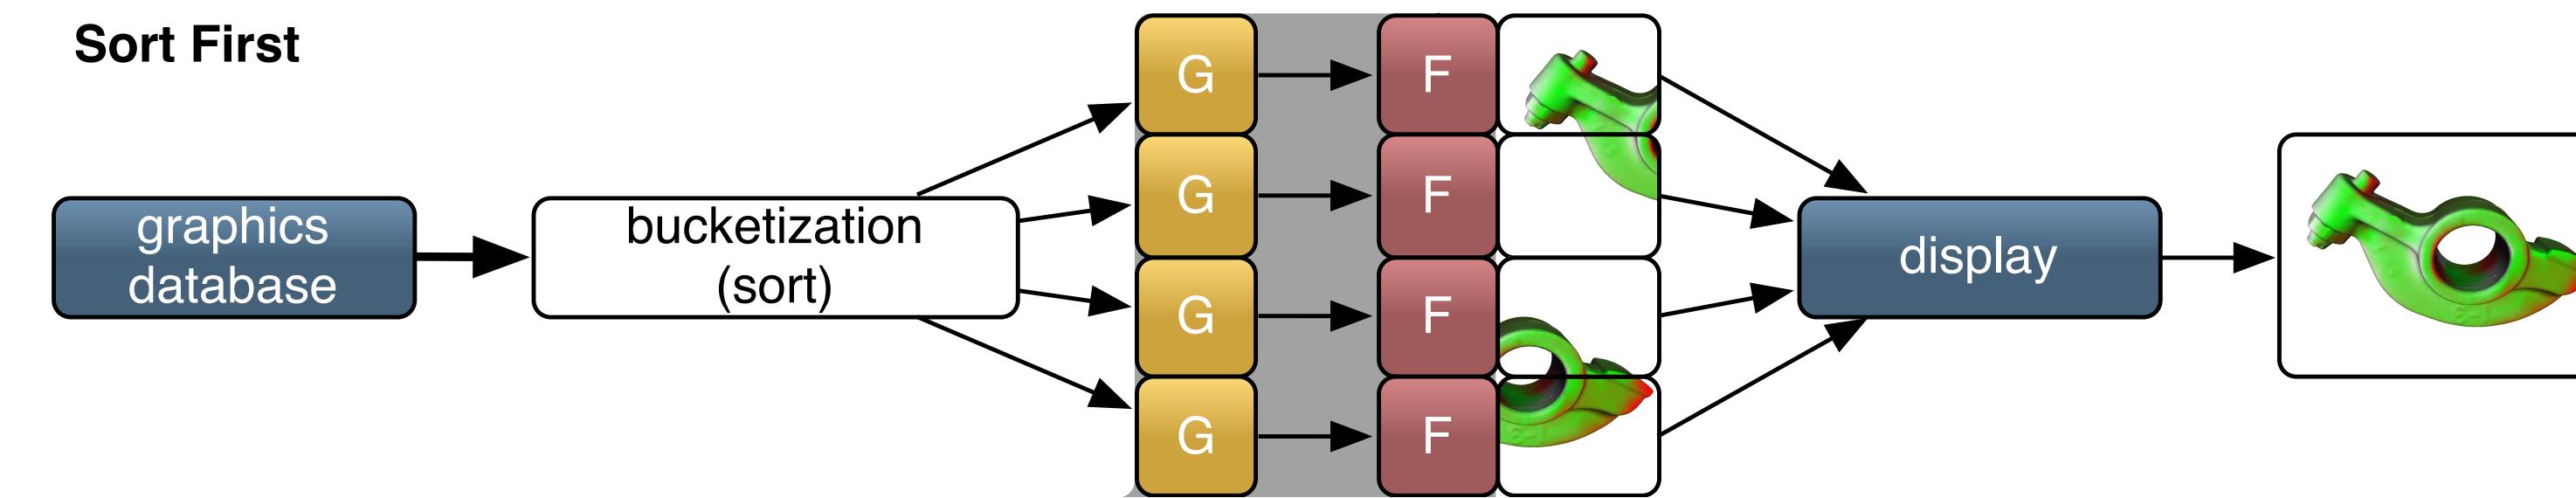
- Parallel and Scalable Rendering
- Load Balancing
- Equalizer Framework Architecture

Parallel and Scalable Rendering

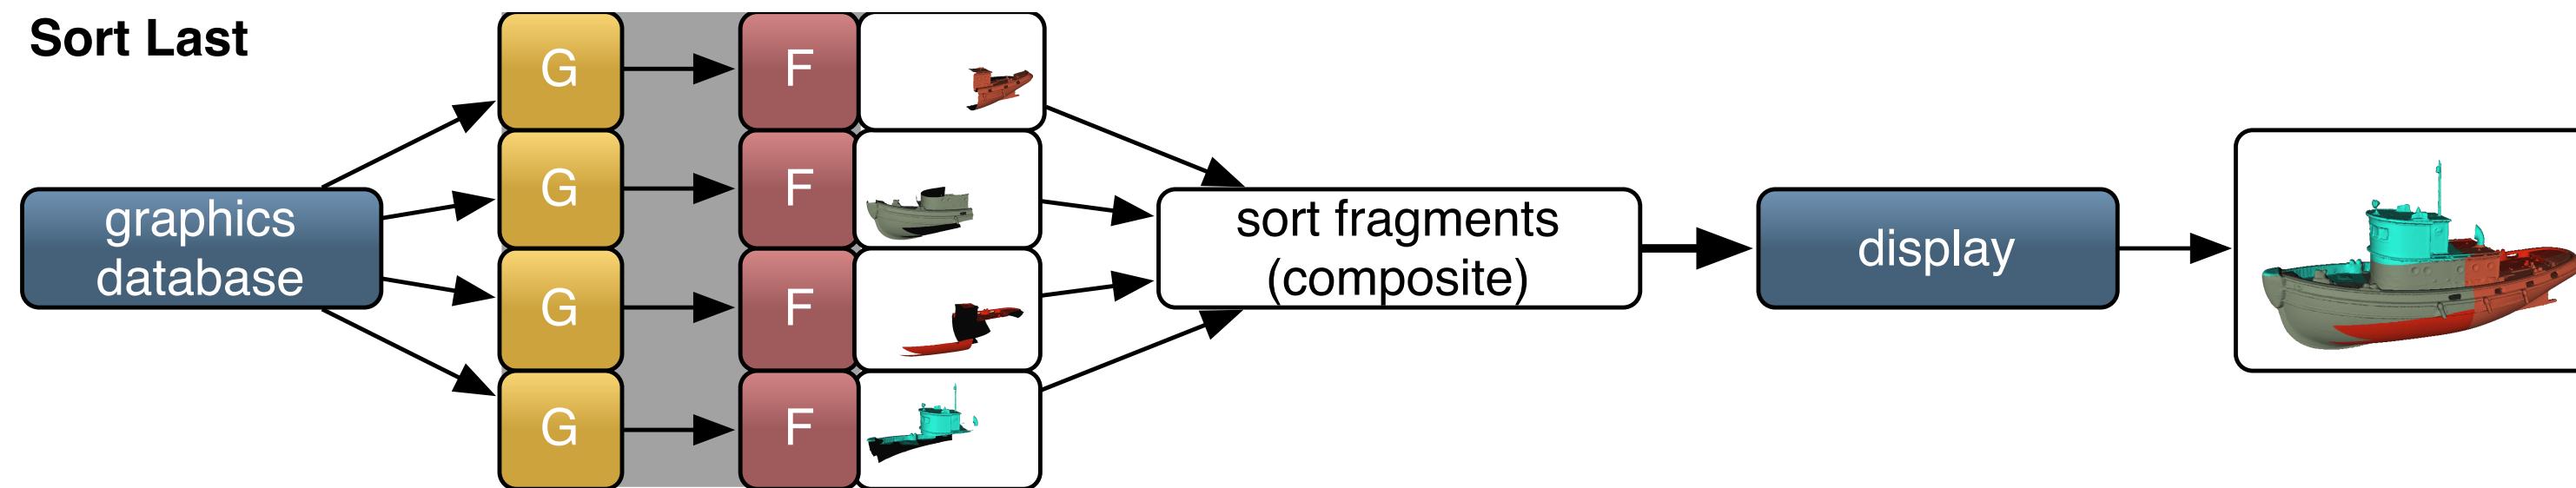
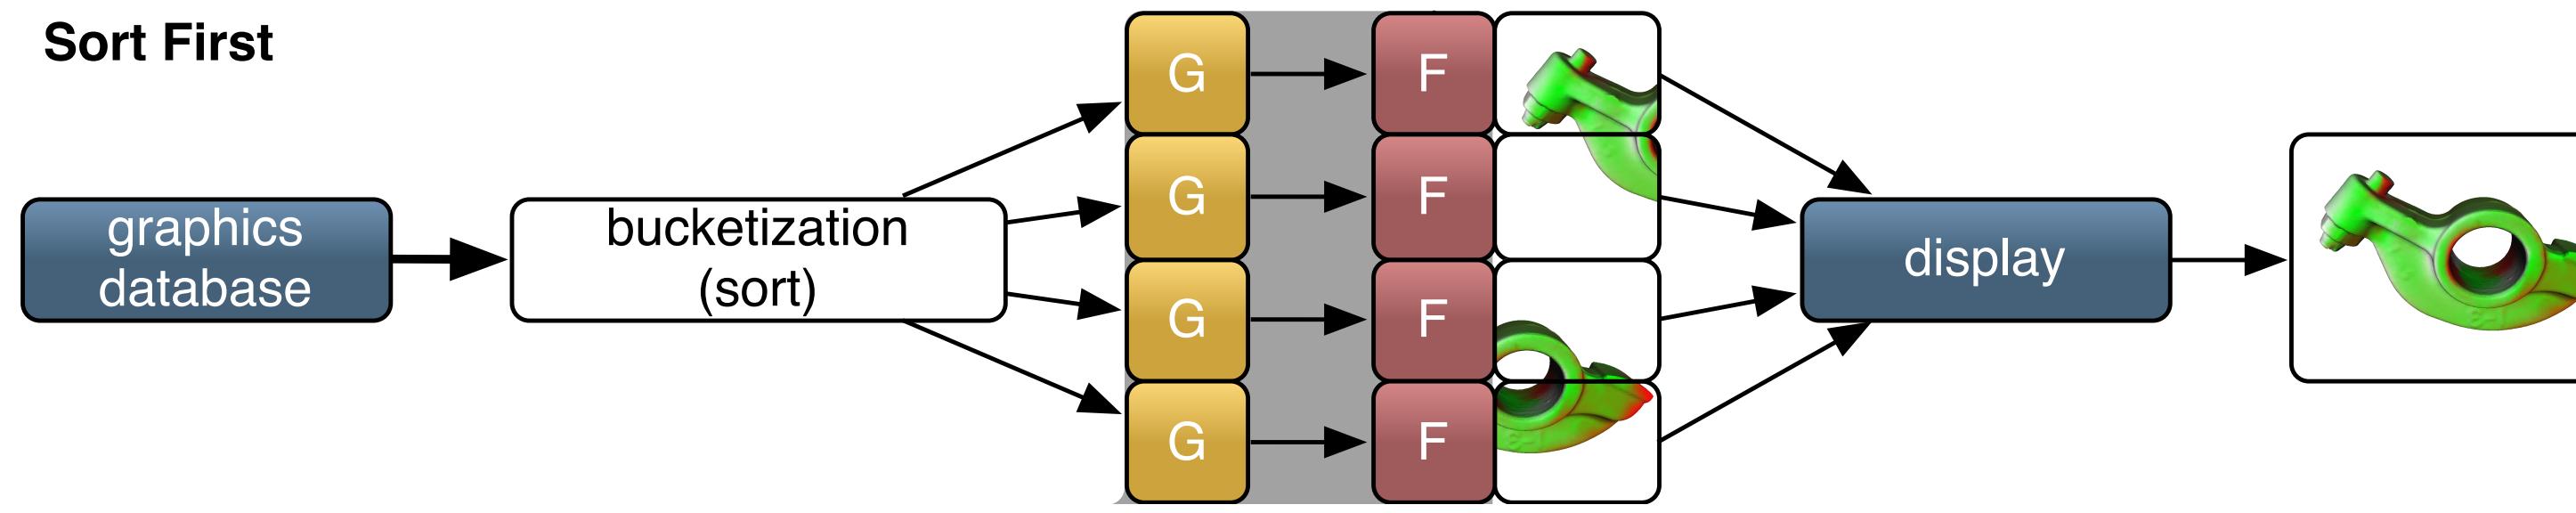
Parallel Rendering



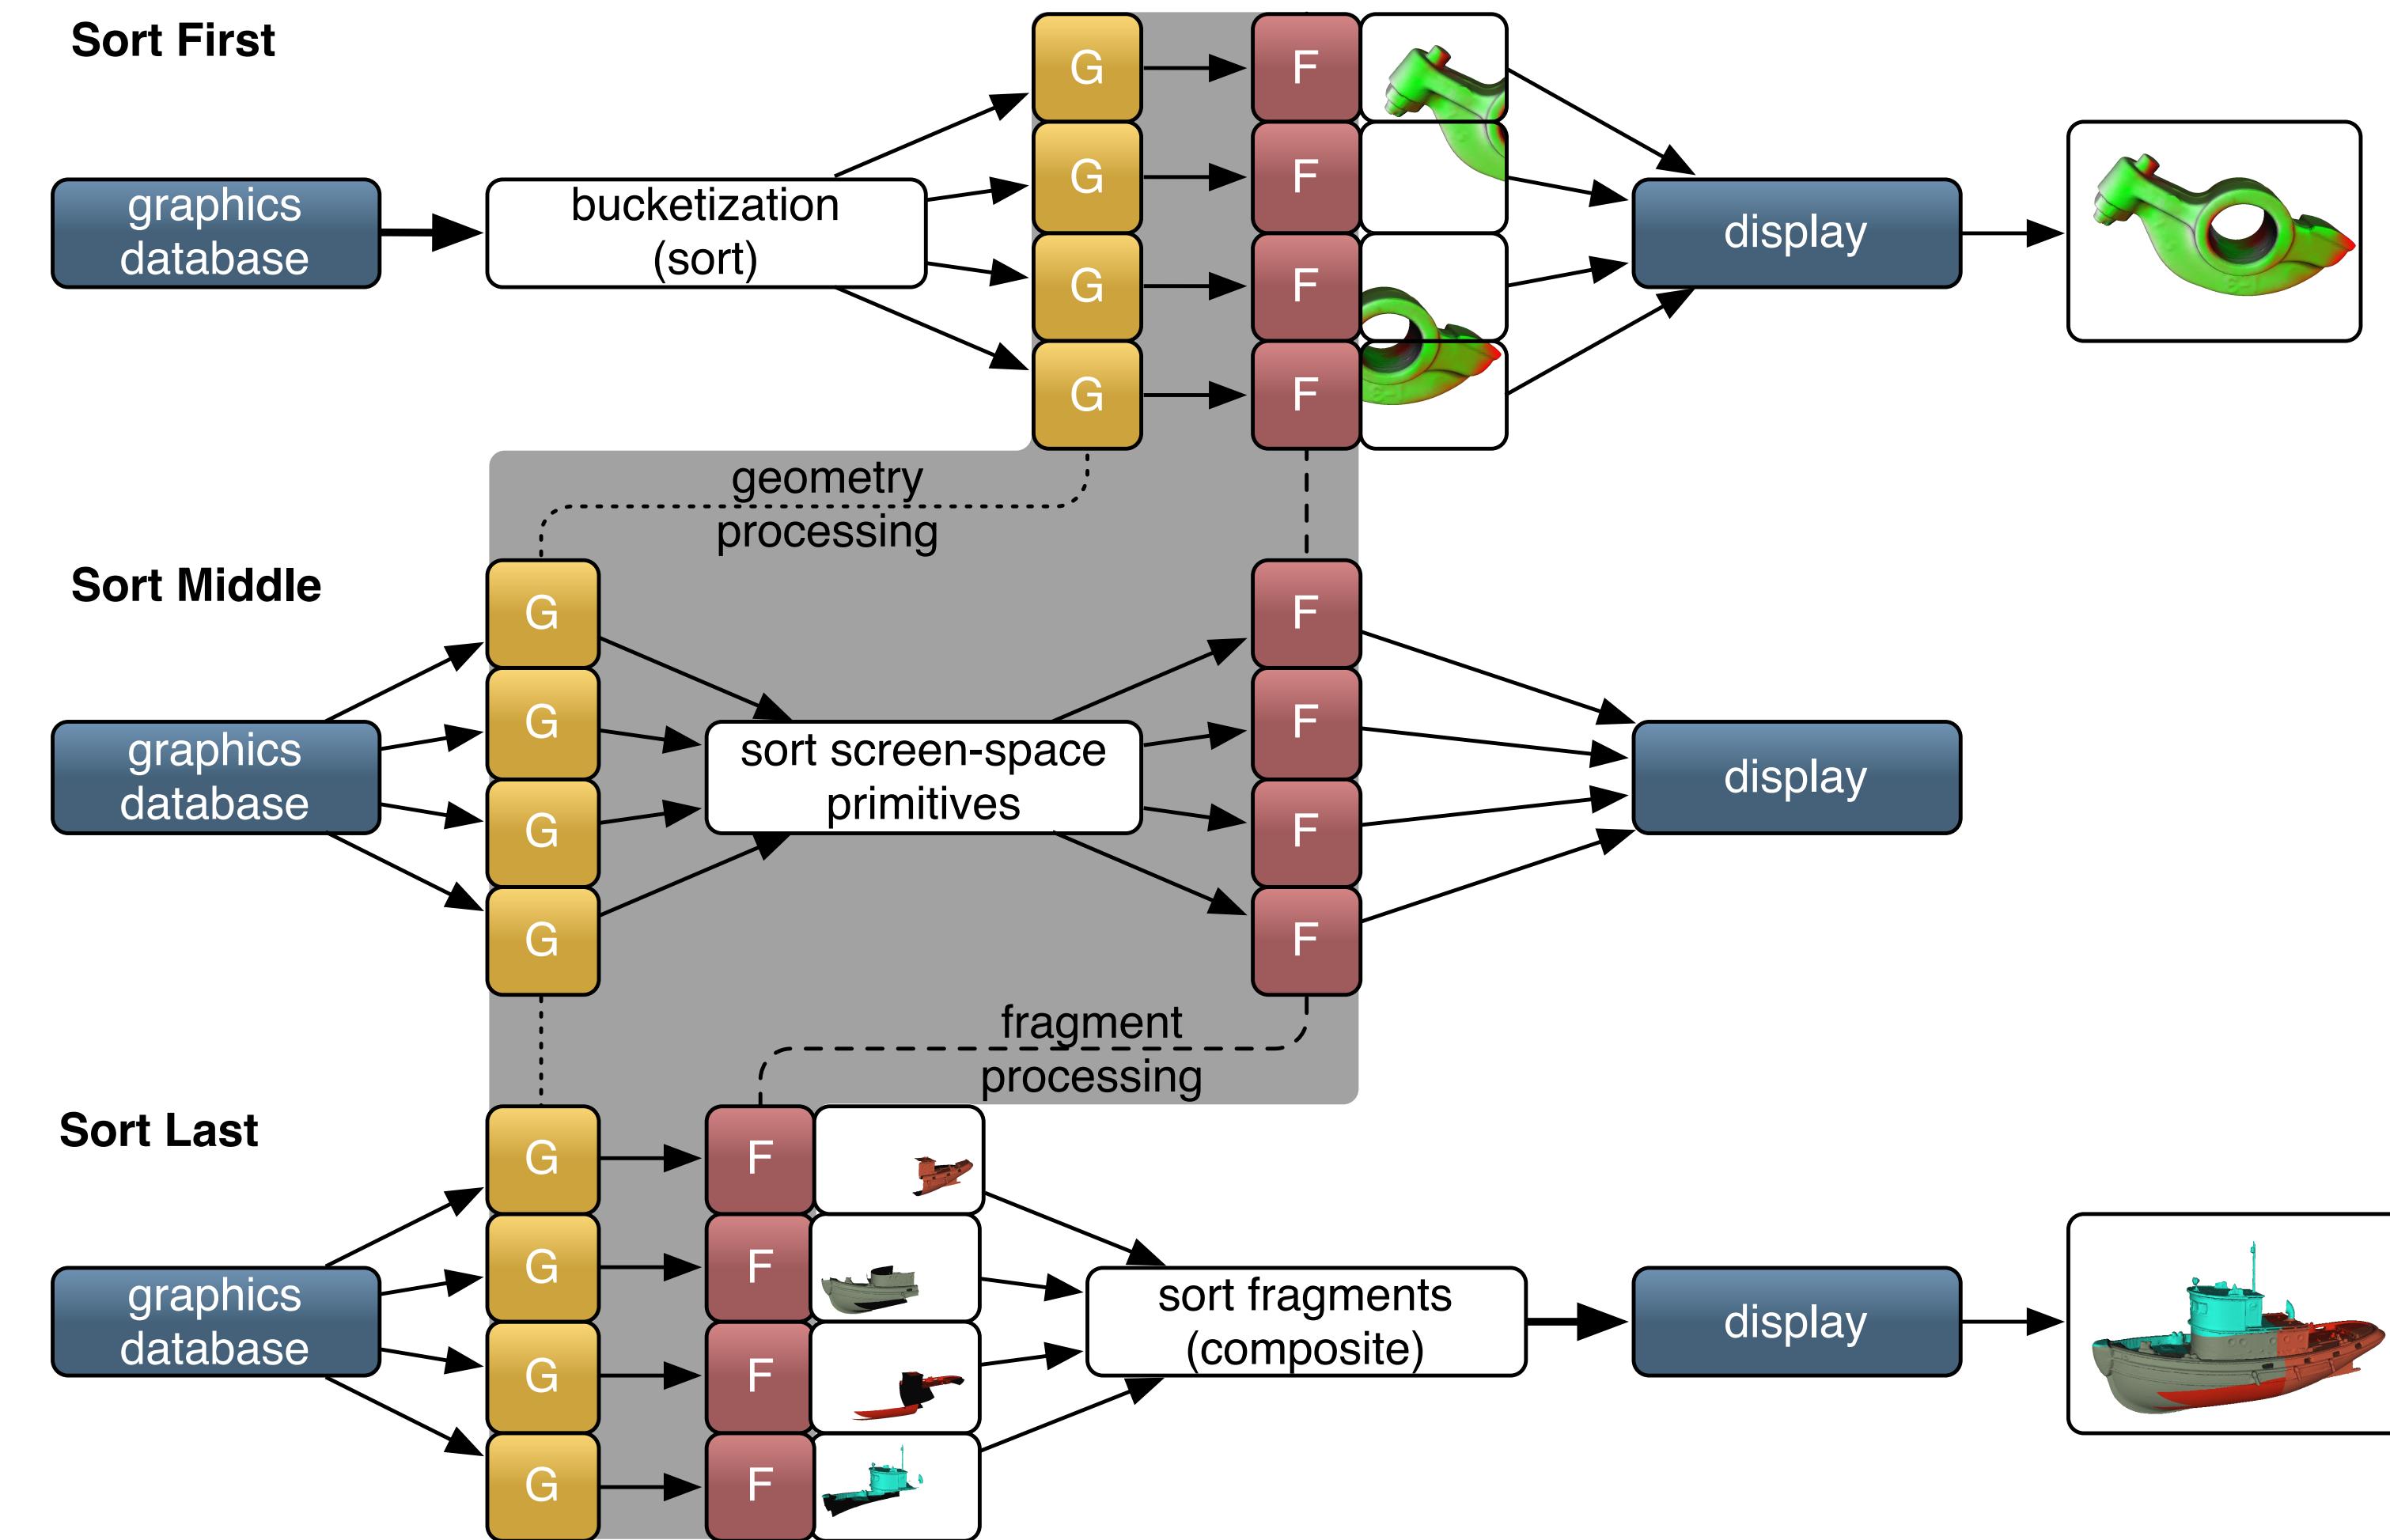
Scalable Rendering



Scalable Rendering



Scalable Rendering



Load Balancing

How we can improve load-balancing for sort-first rendering, in particular for large display systems?

Single Display Load Balancing

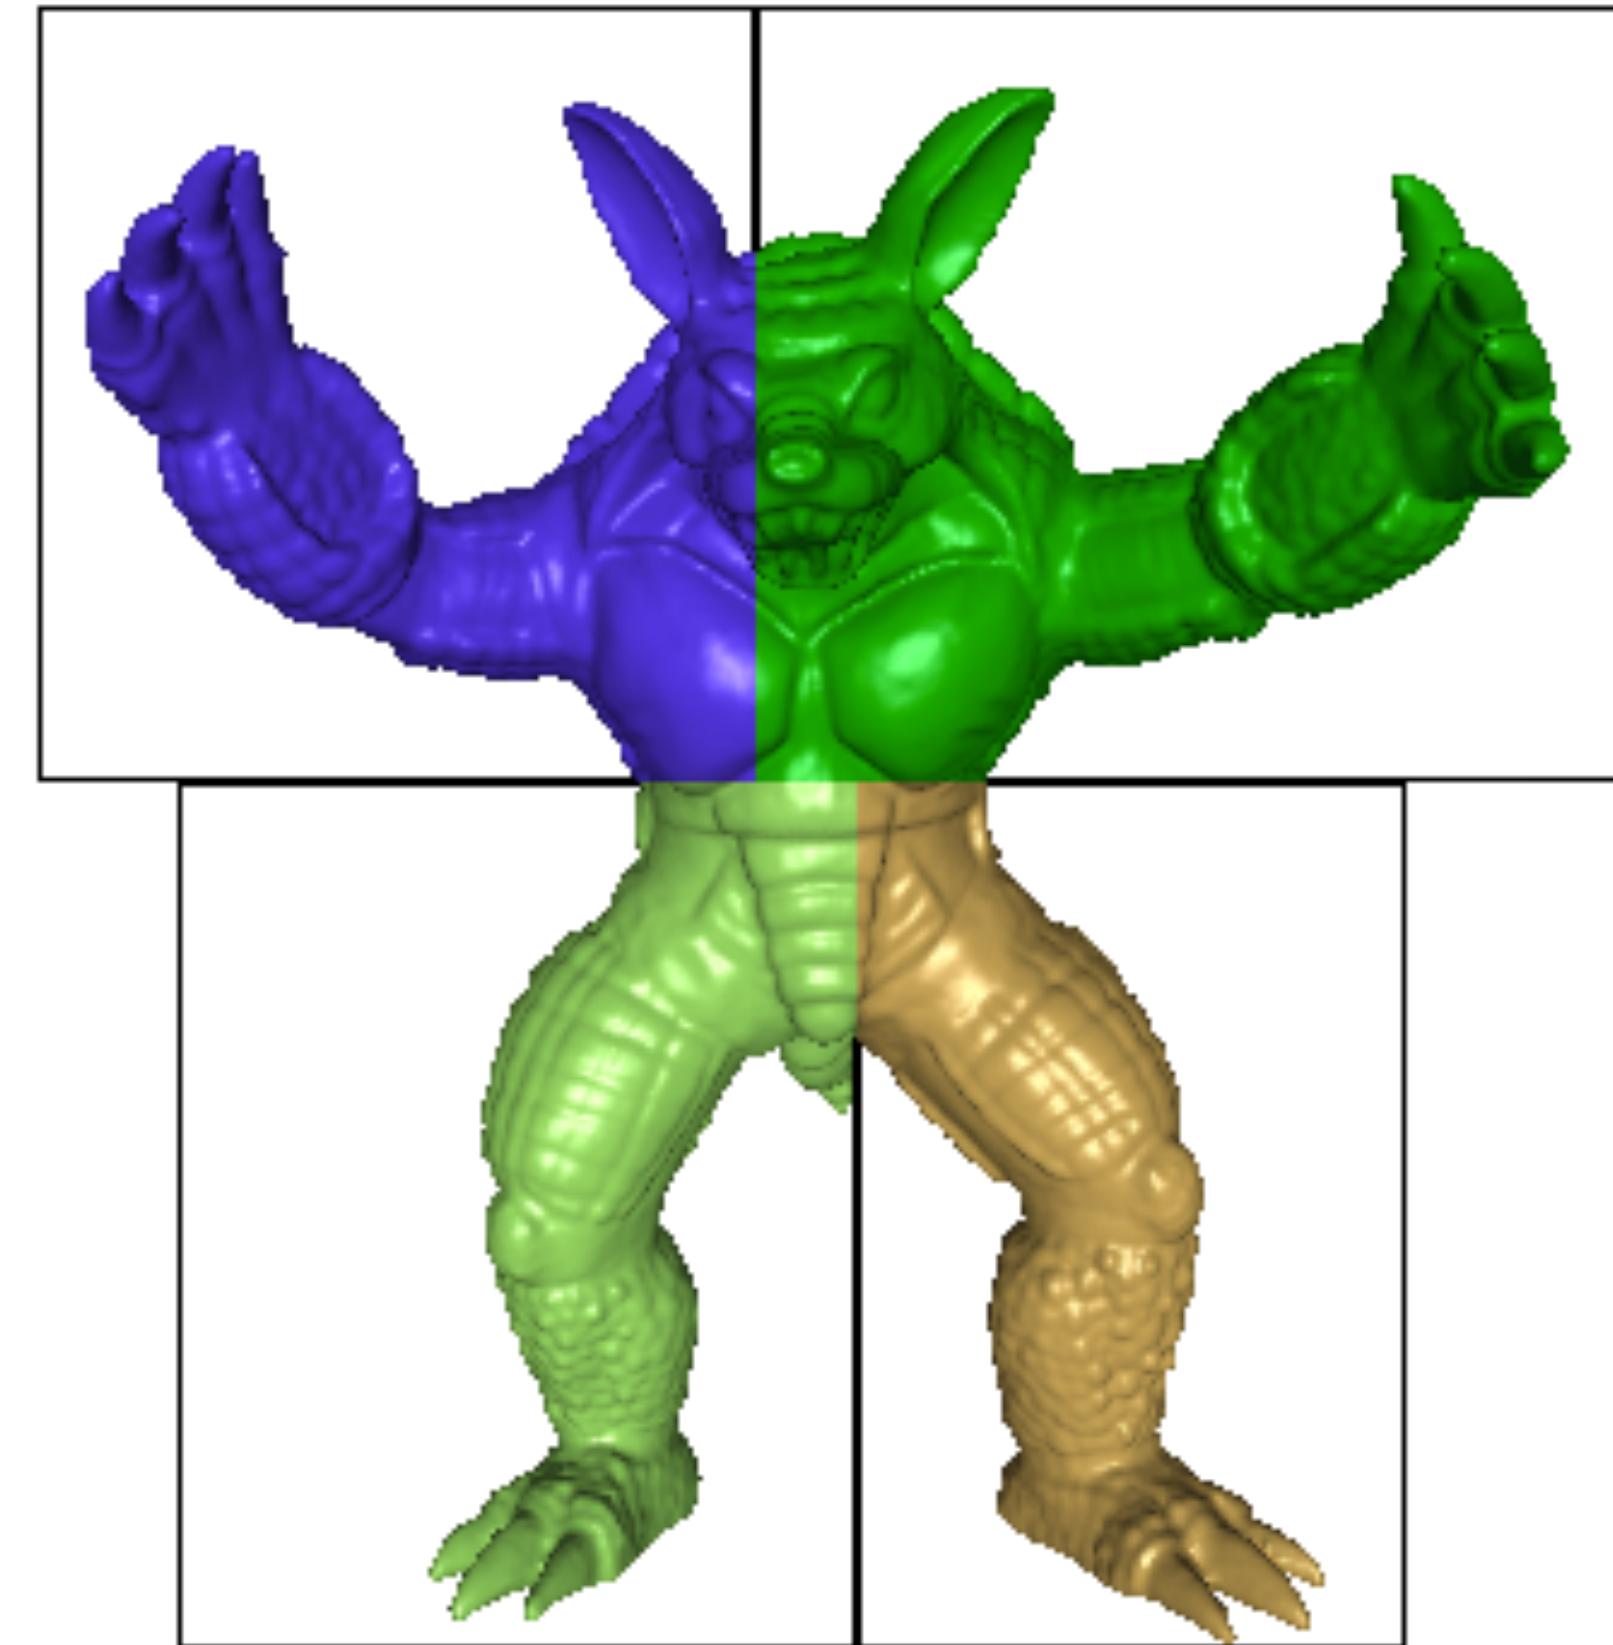
- Static, Reactive,
Predictive, Implicit

Single Display Load Balancing

- Static, Reactive,
Predictive, Implicit
- Equalizers for
runtime adaptations

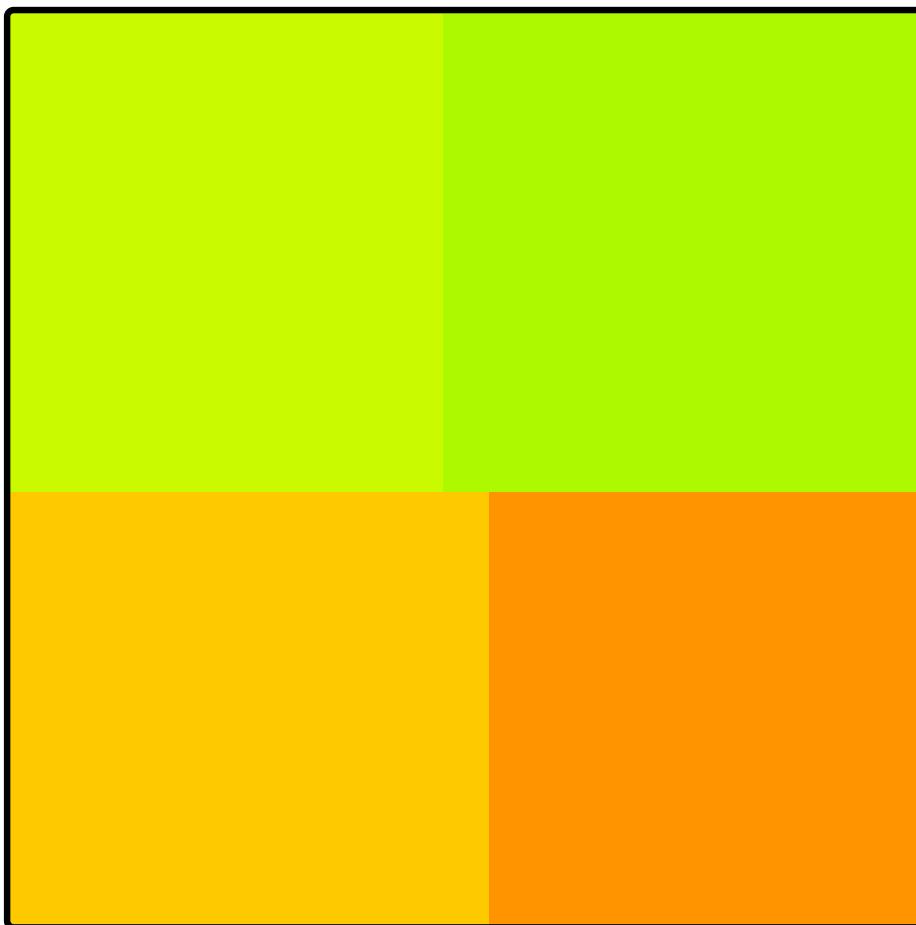
Single Display Load Balancing

- Static, Reactive,
Predictive, Implicit
- Equalizers for
runtime adaptations
- Binary split tree



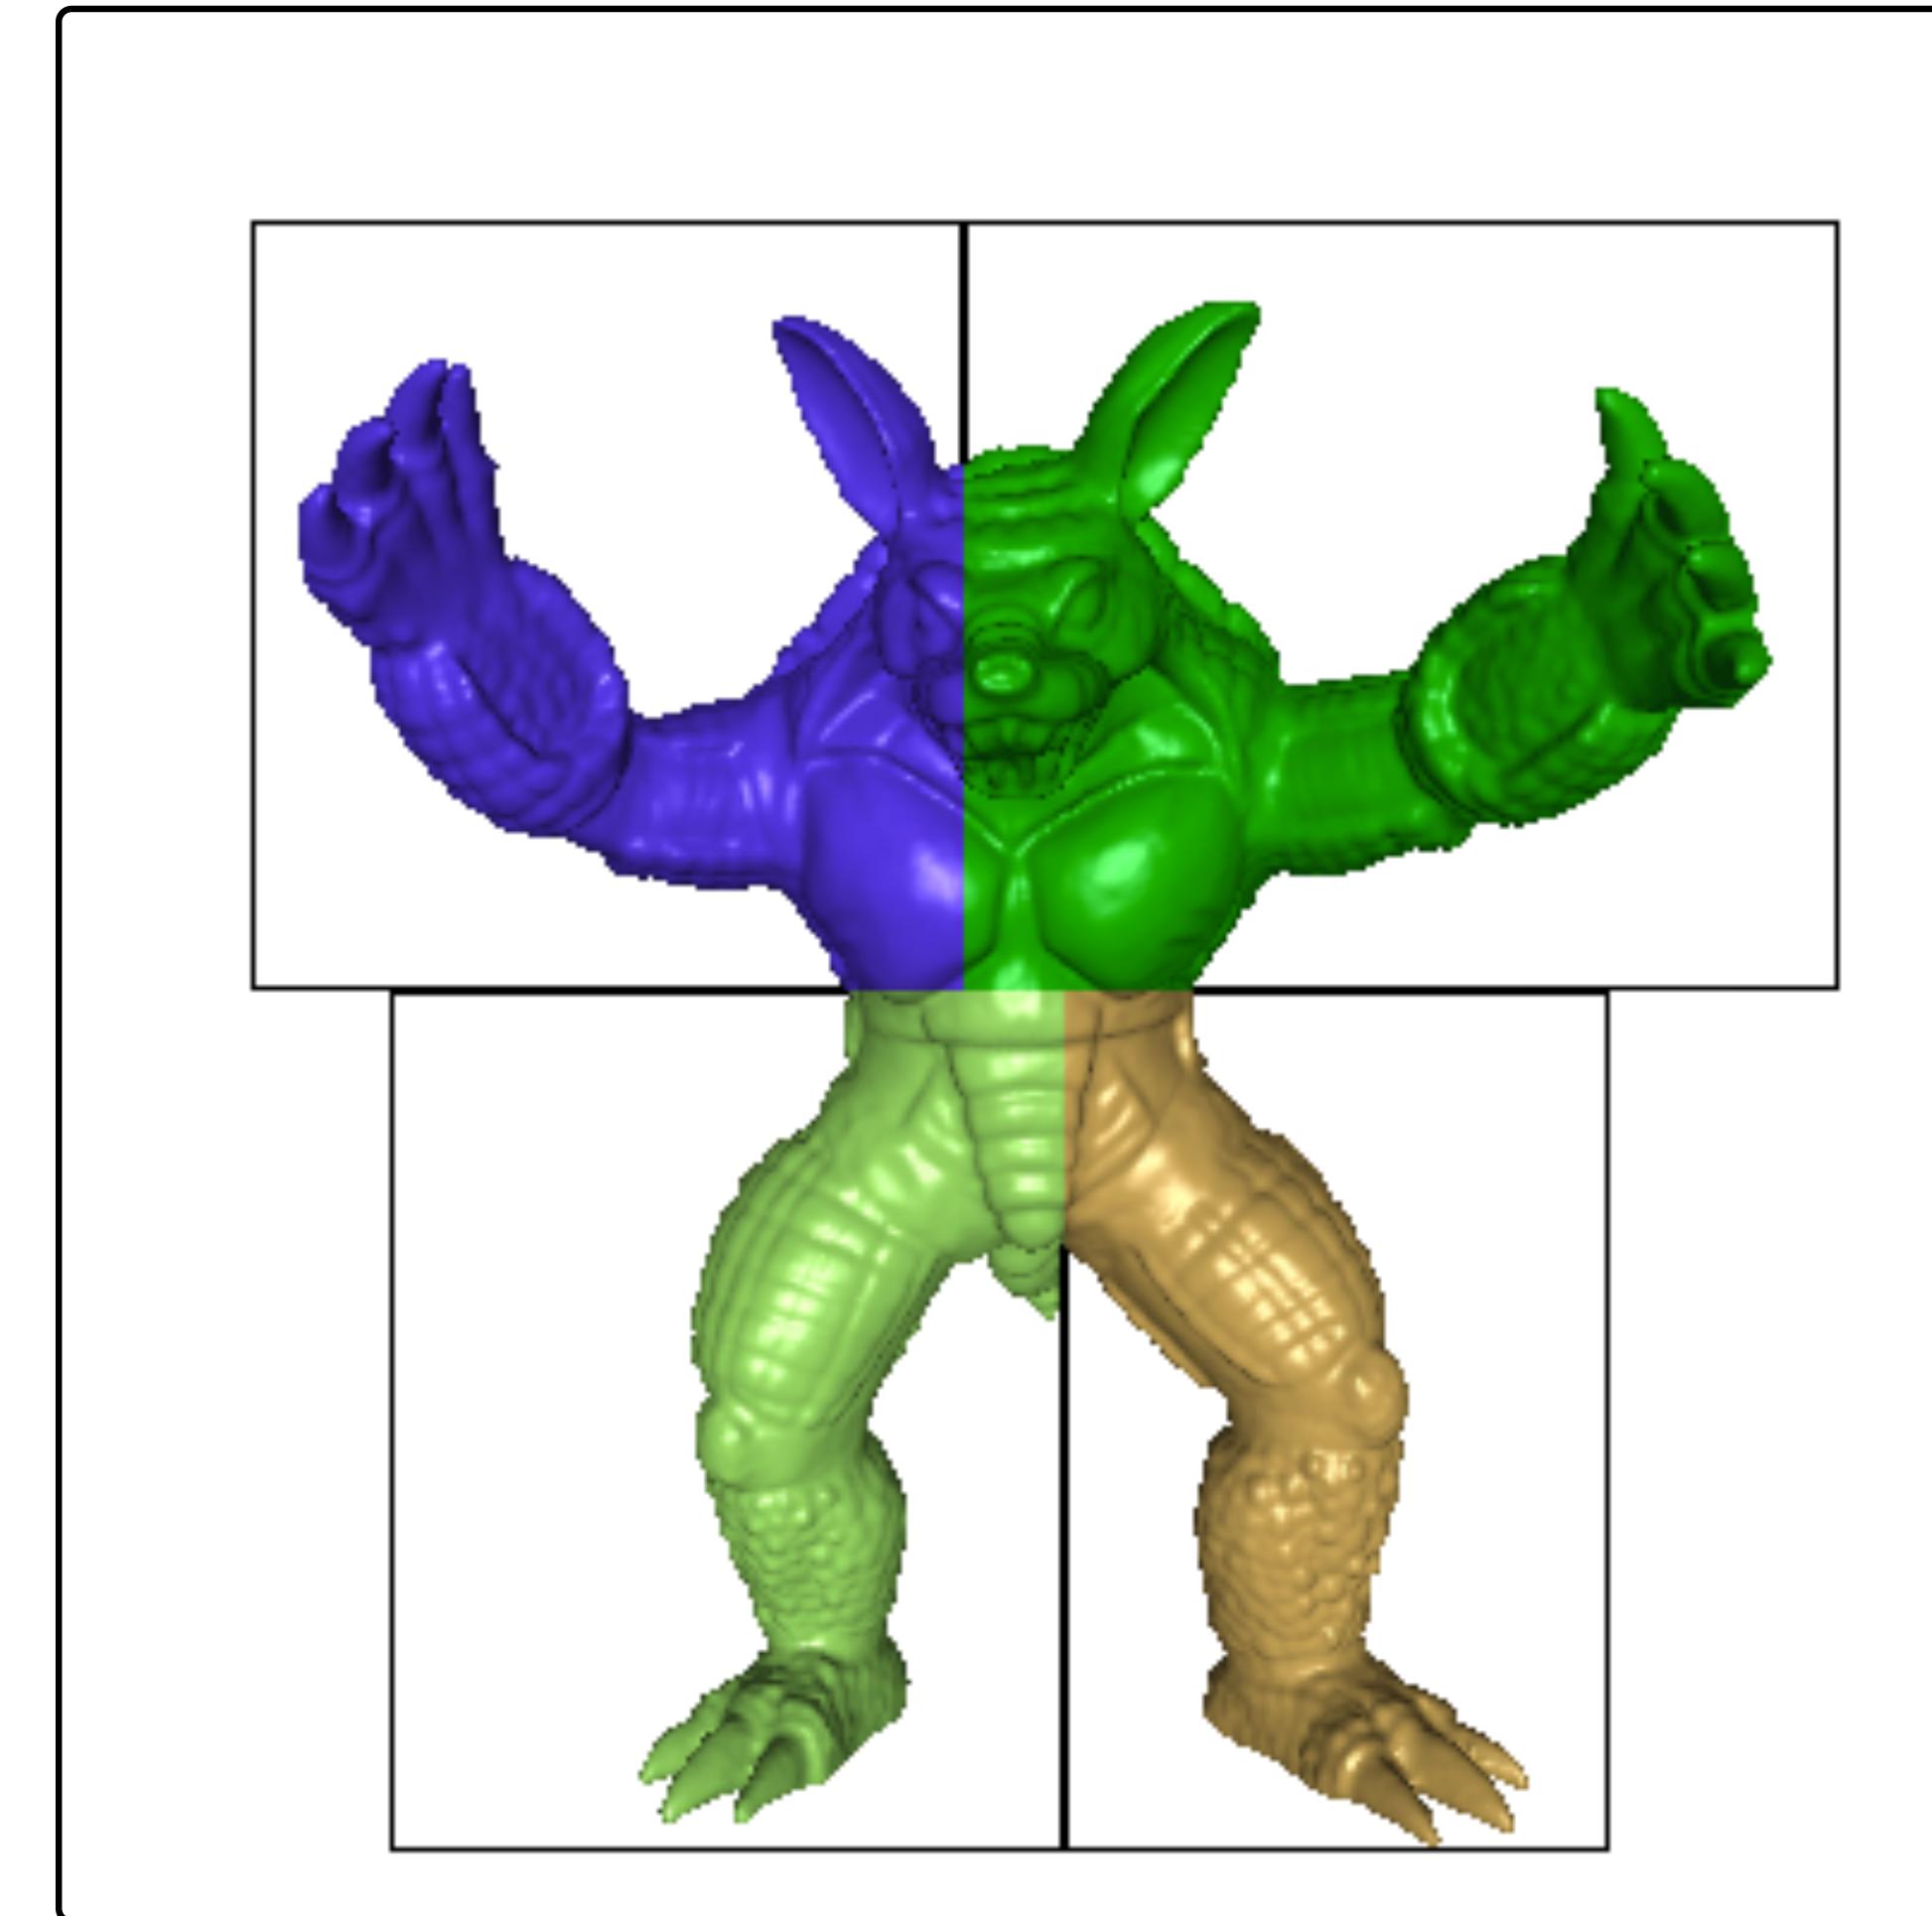
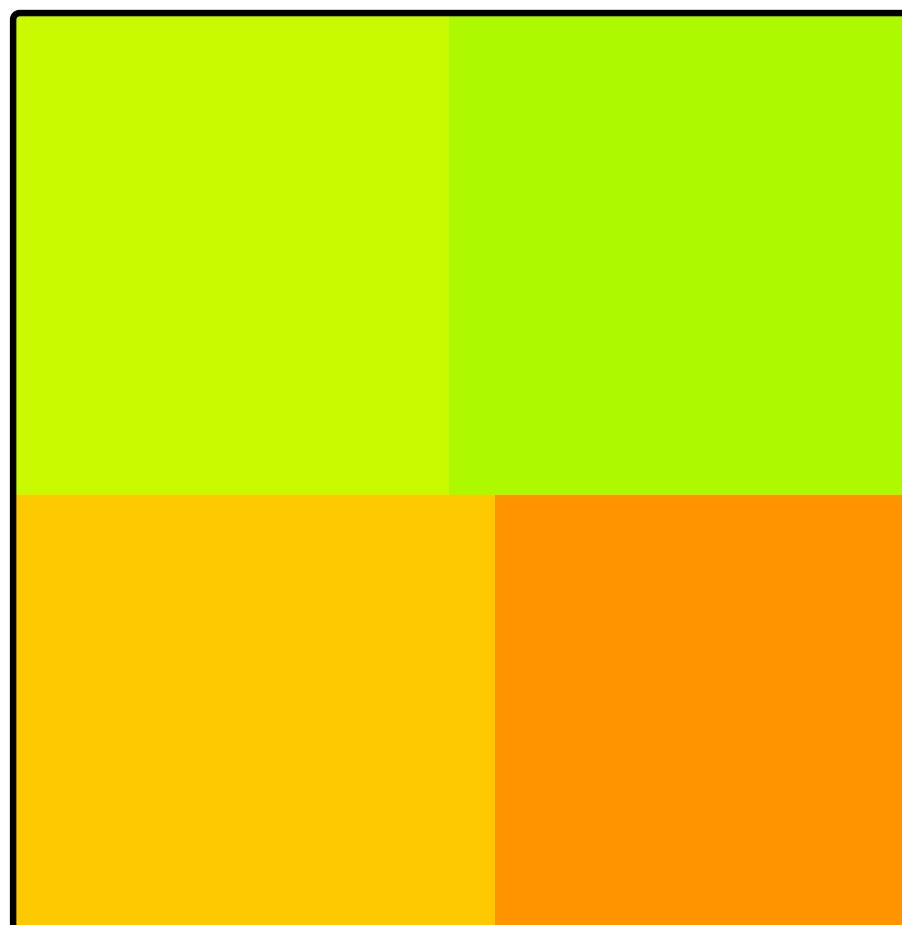
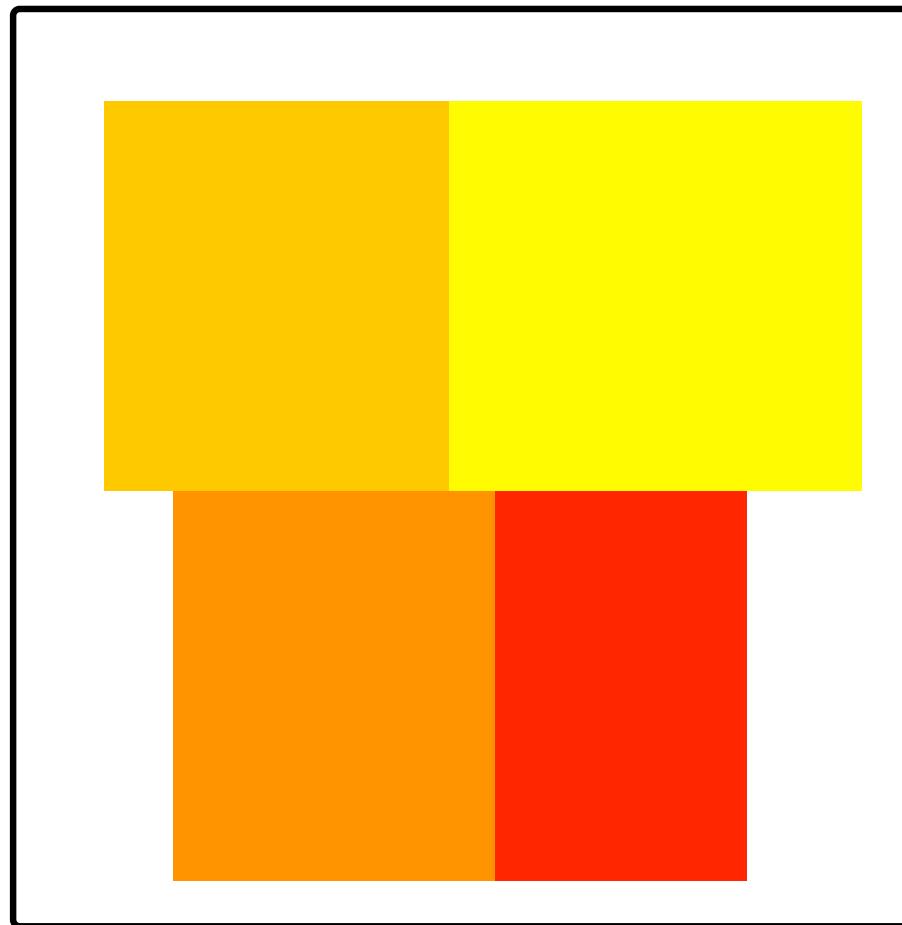
Single Display Load Balancing

- Static, Reactive, Predictive, Implicit
- Equalizers for runtime adaptations
- Binary split tree
 - ▶ Load Equalizer: Integrate over load grid



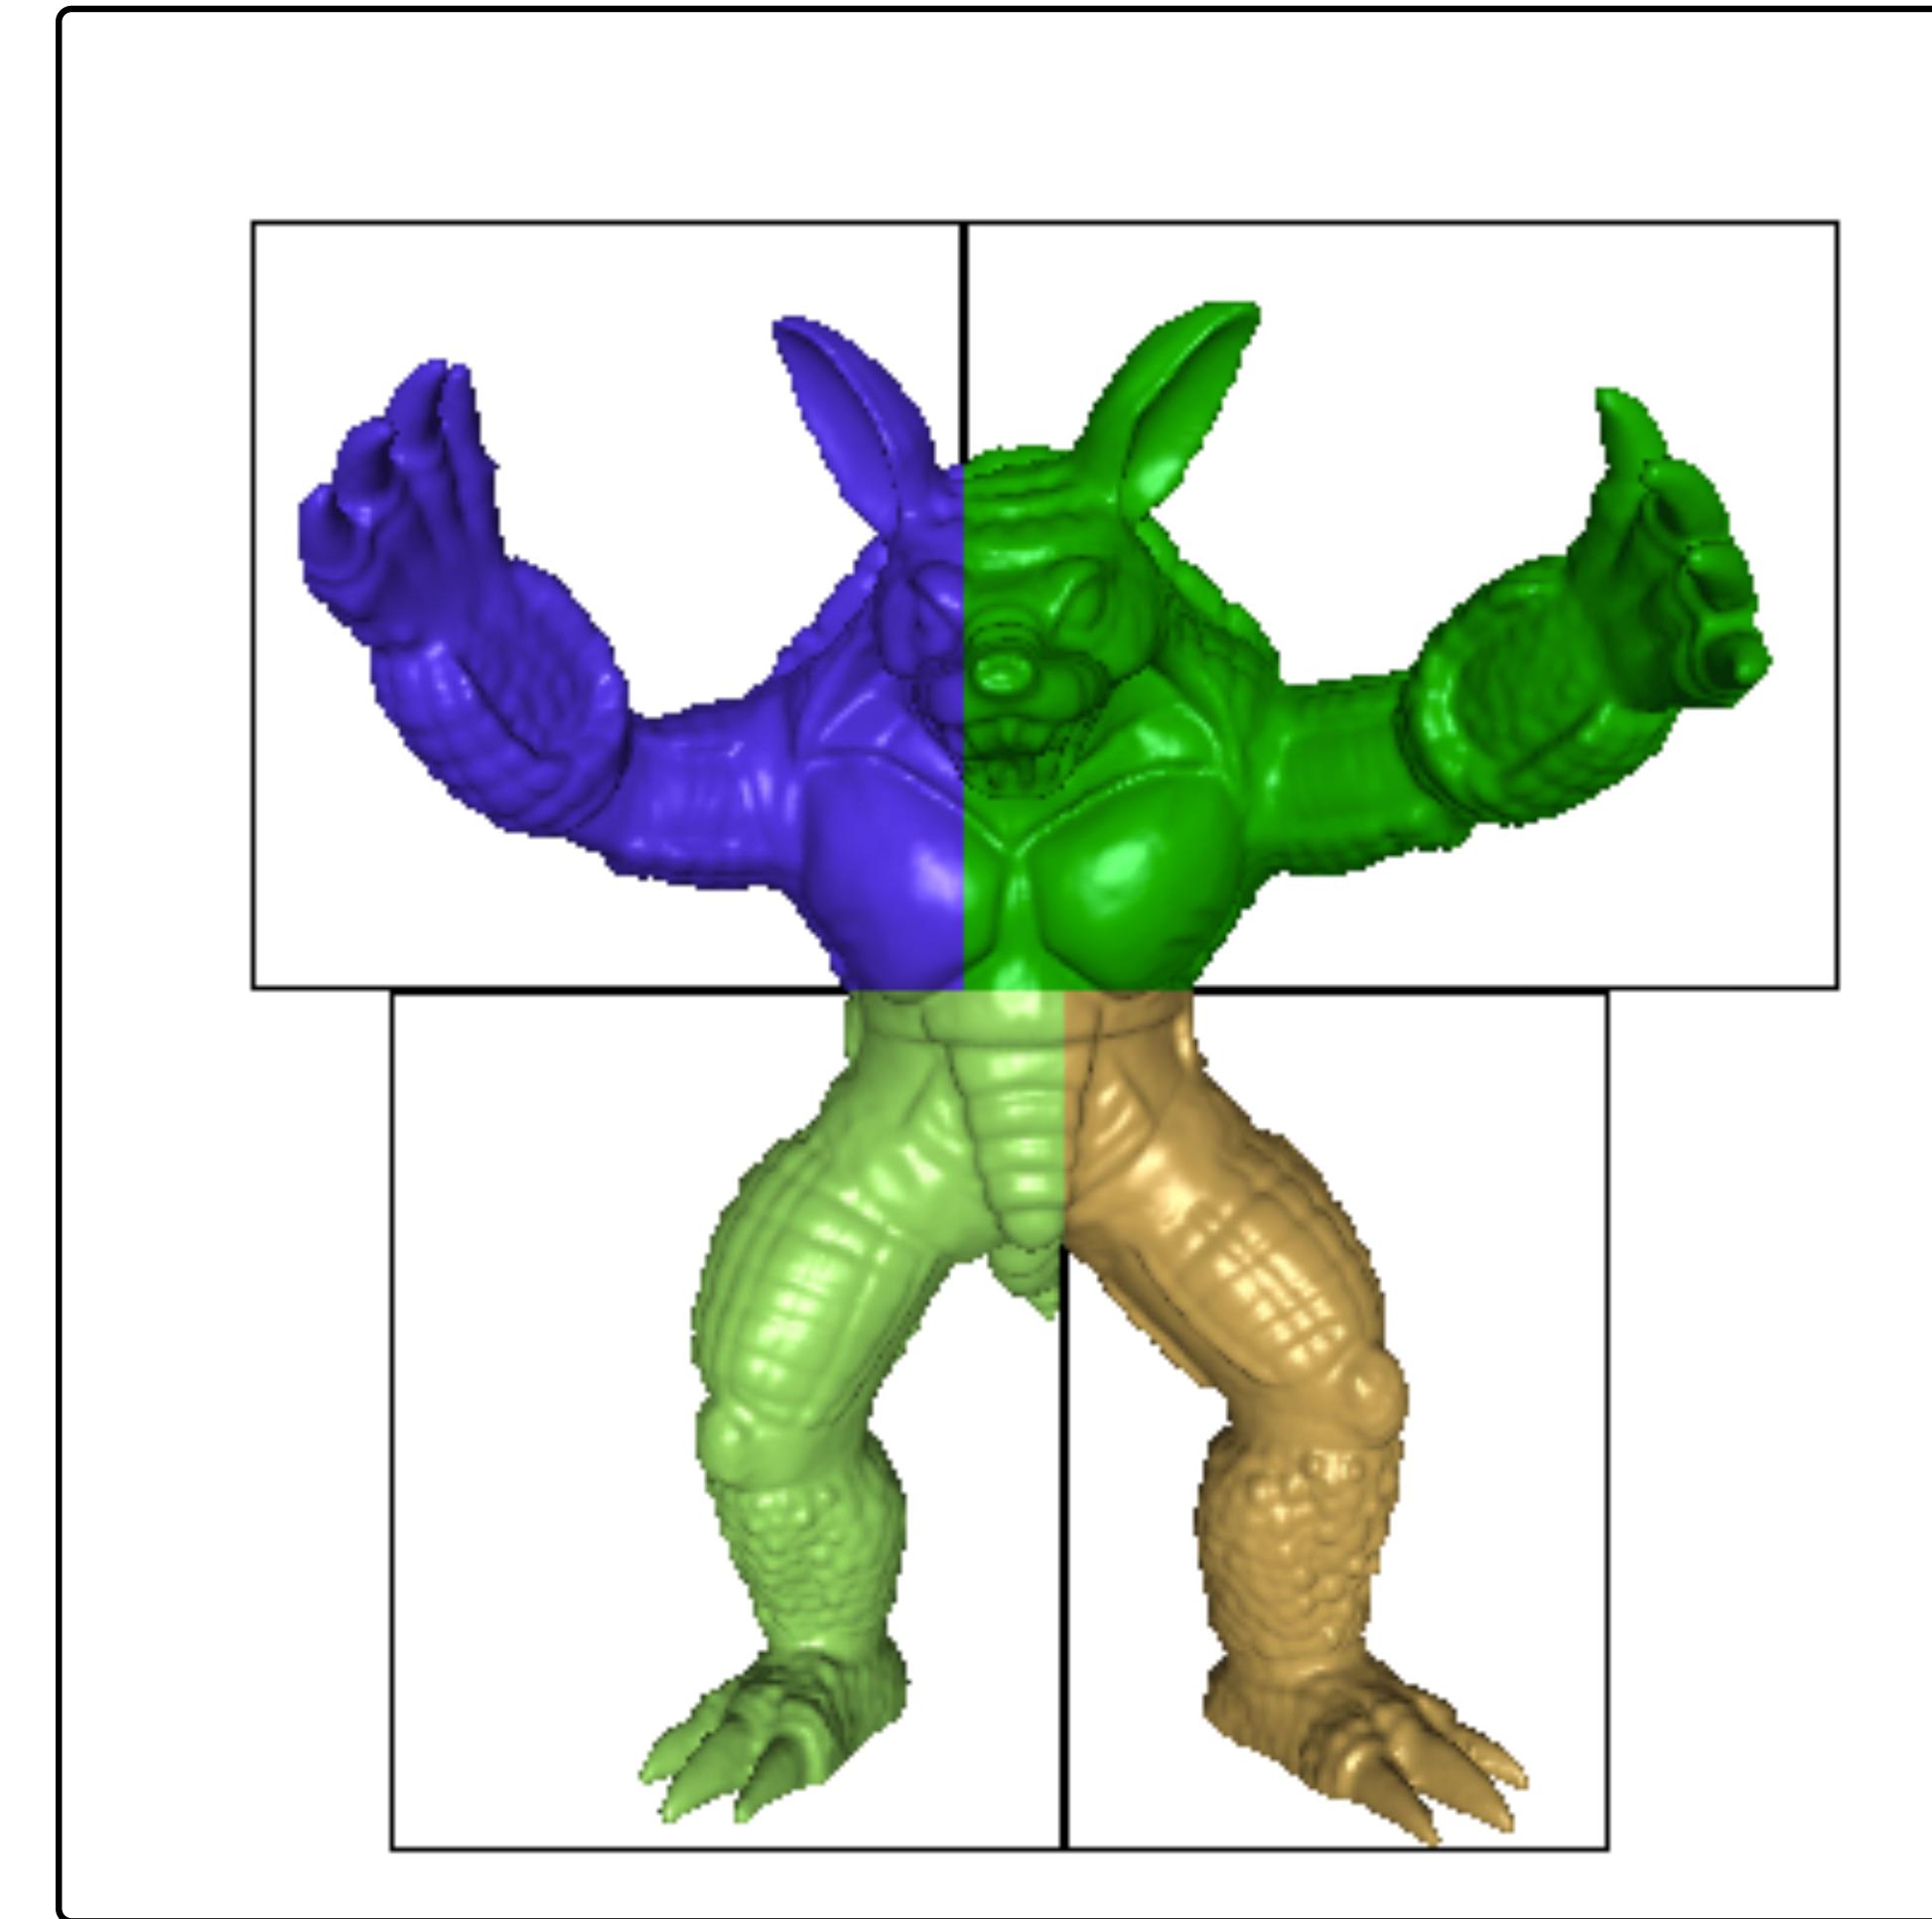
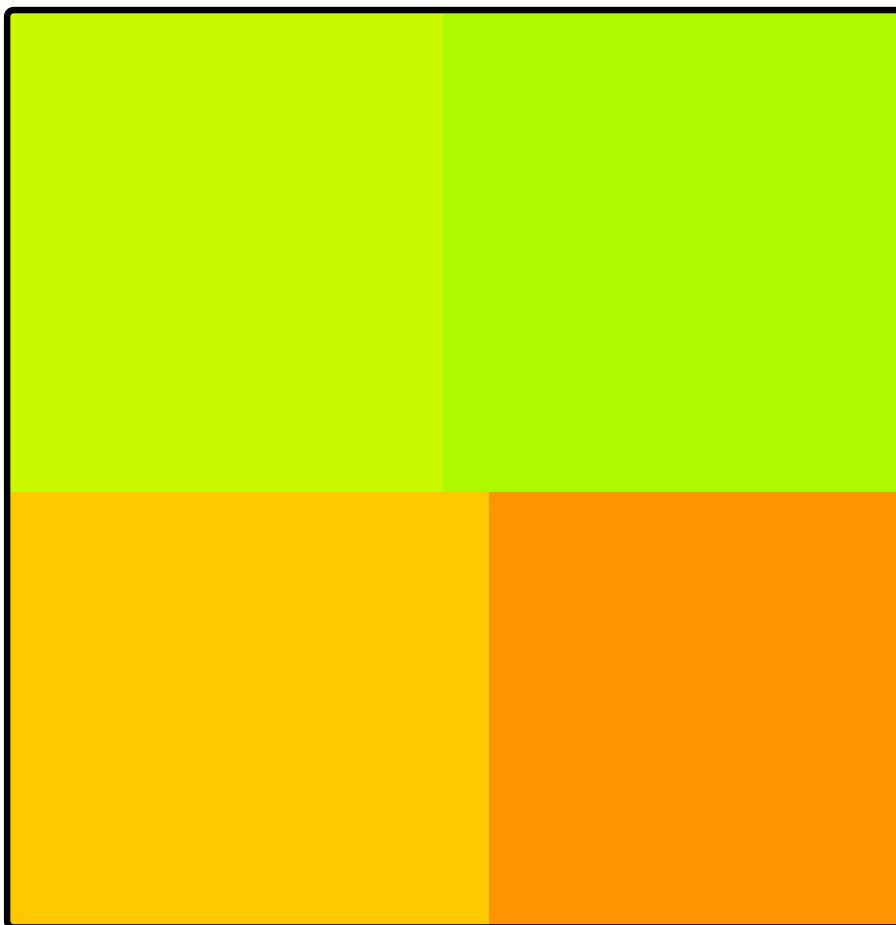
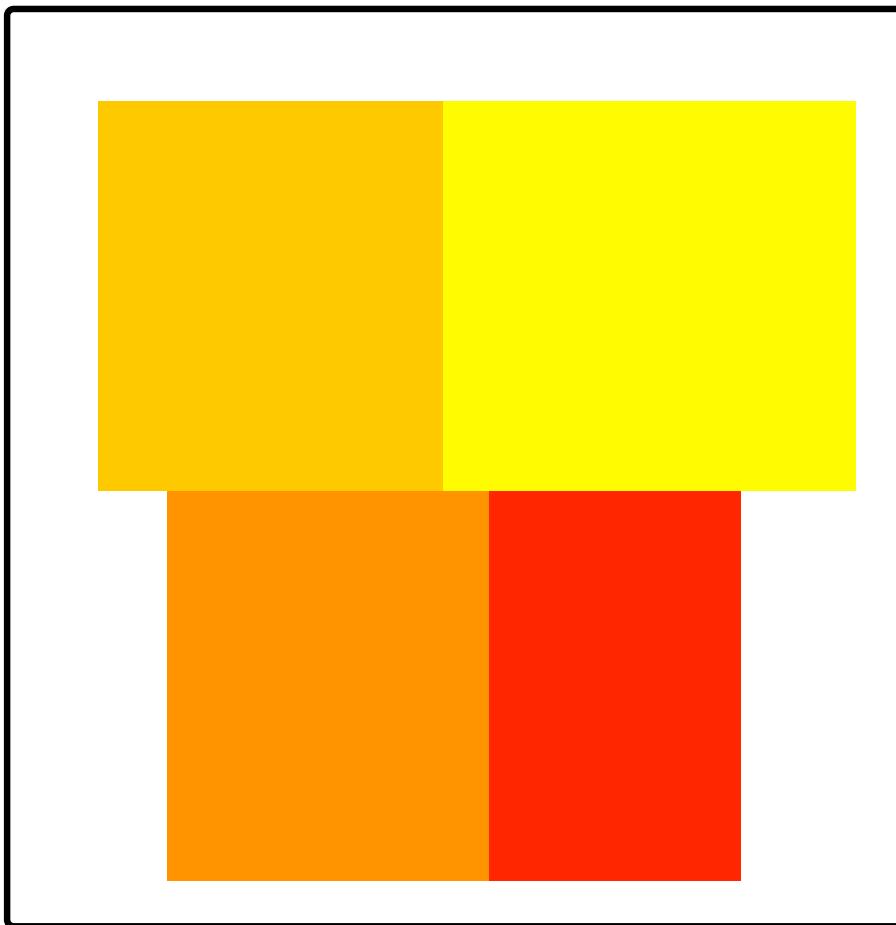
Single Display Load Balancing

- Static, Reactive, Predictive, Implicit
- Equalizers for runtime adaptations
- Binary split tree
 - ▶ Load Equalizer: Integrate over load grid



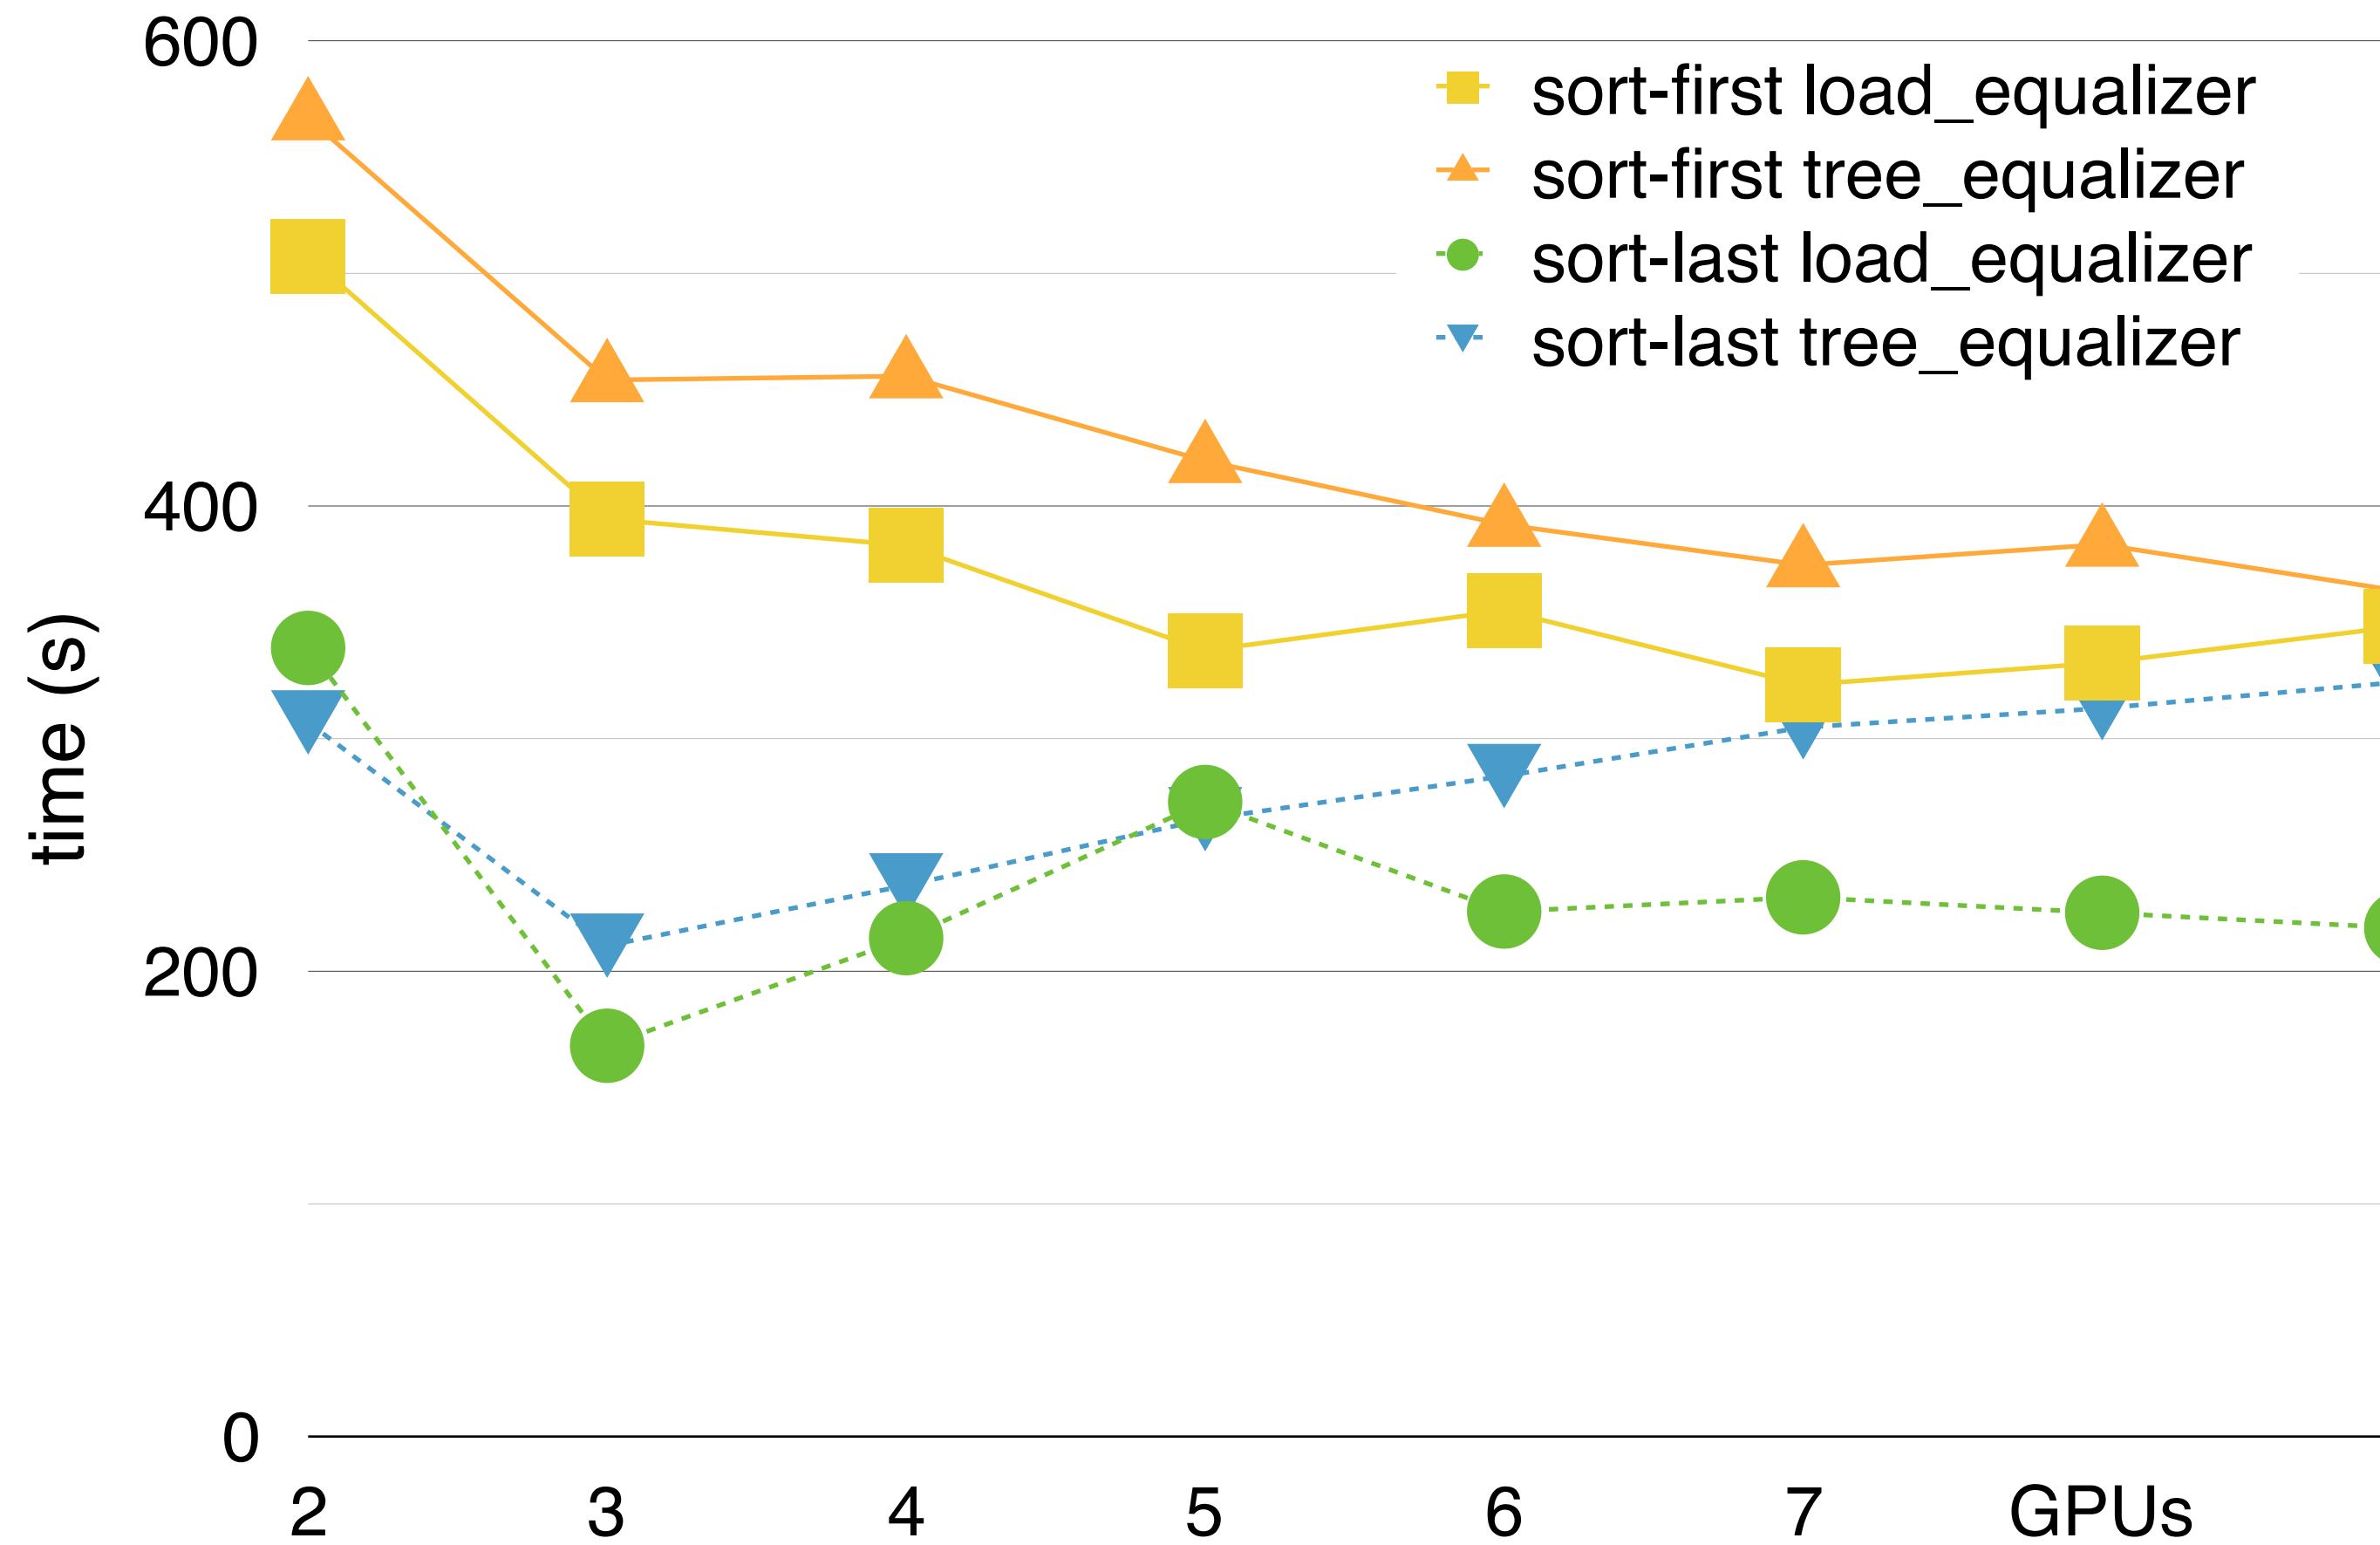
Single Display Load Balancing

- Static, Reactive, Predictive, Implicit
- Equalizers for runtime adaptations
- Binary split tree
 - ▶ Load Equalizer: Integrate over load grid
 - ▶ Tree Equalizer: Balance left/right render time

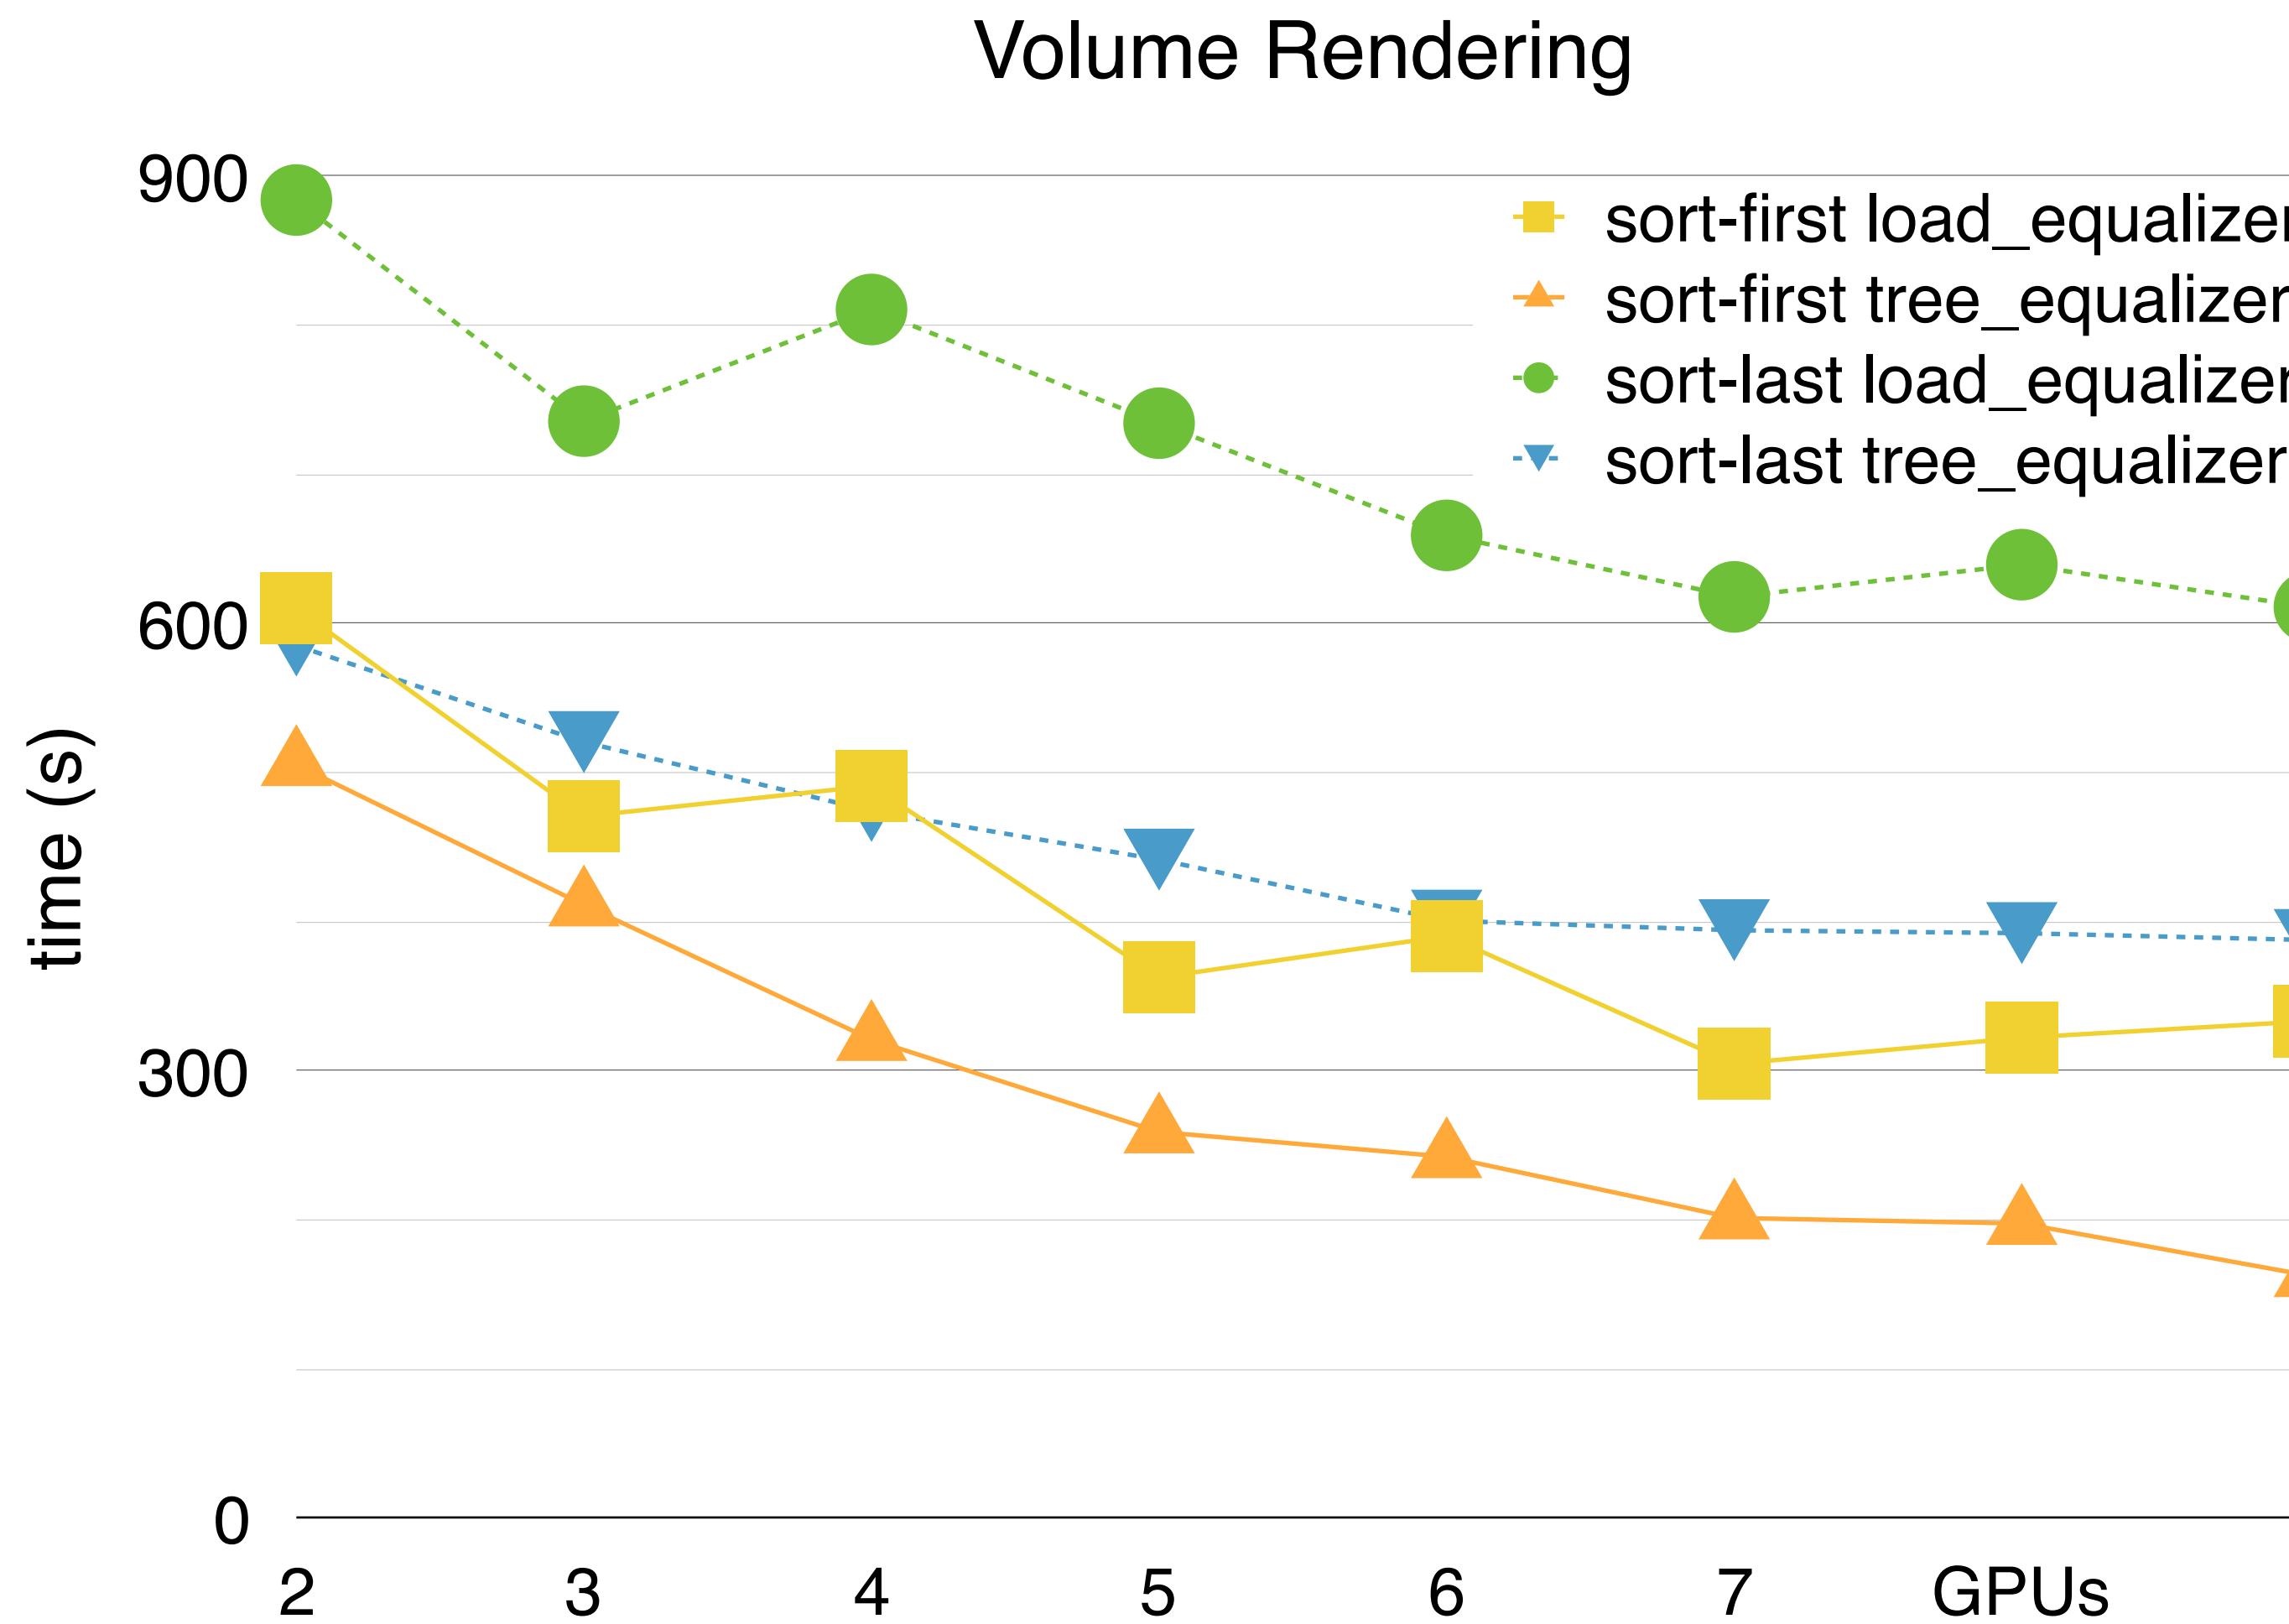


Single Display Load Balancing

Polygonal Rendering

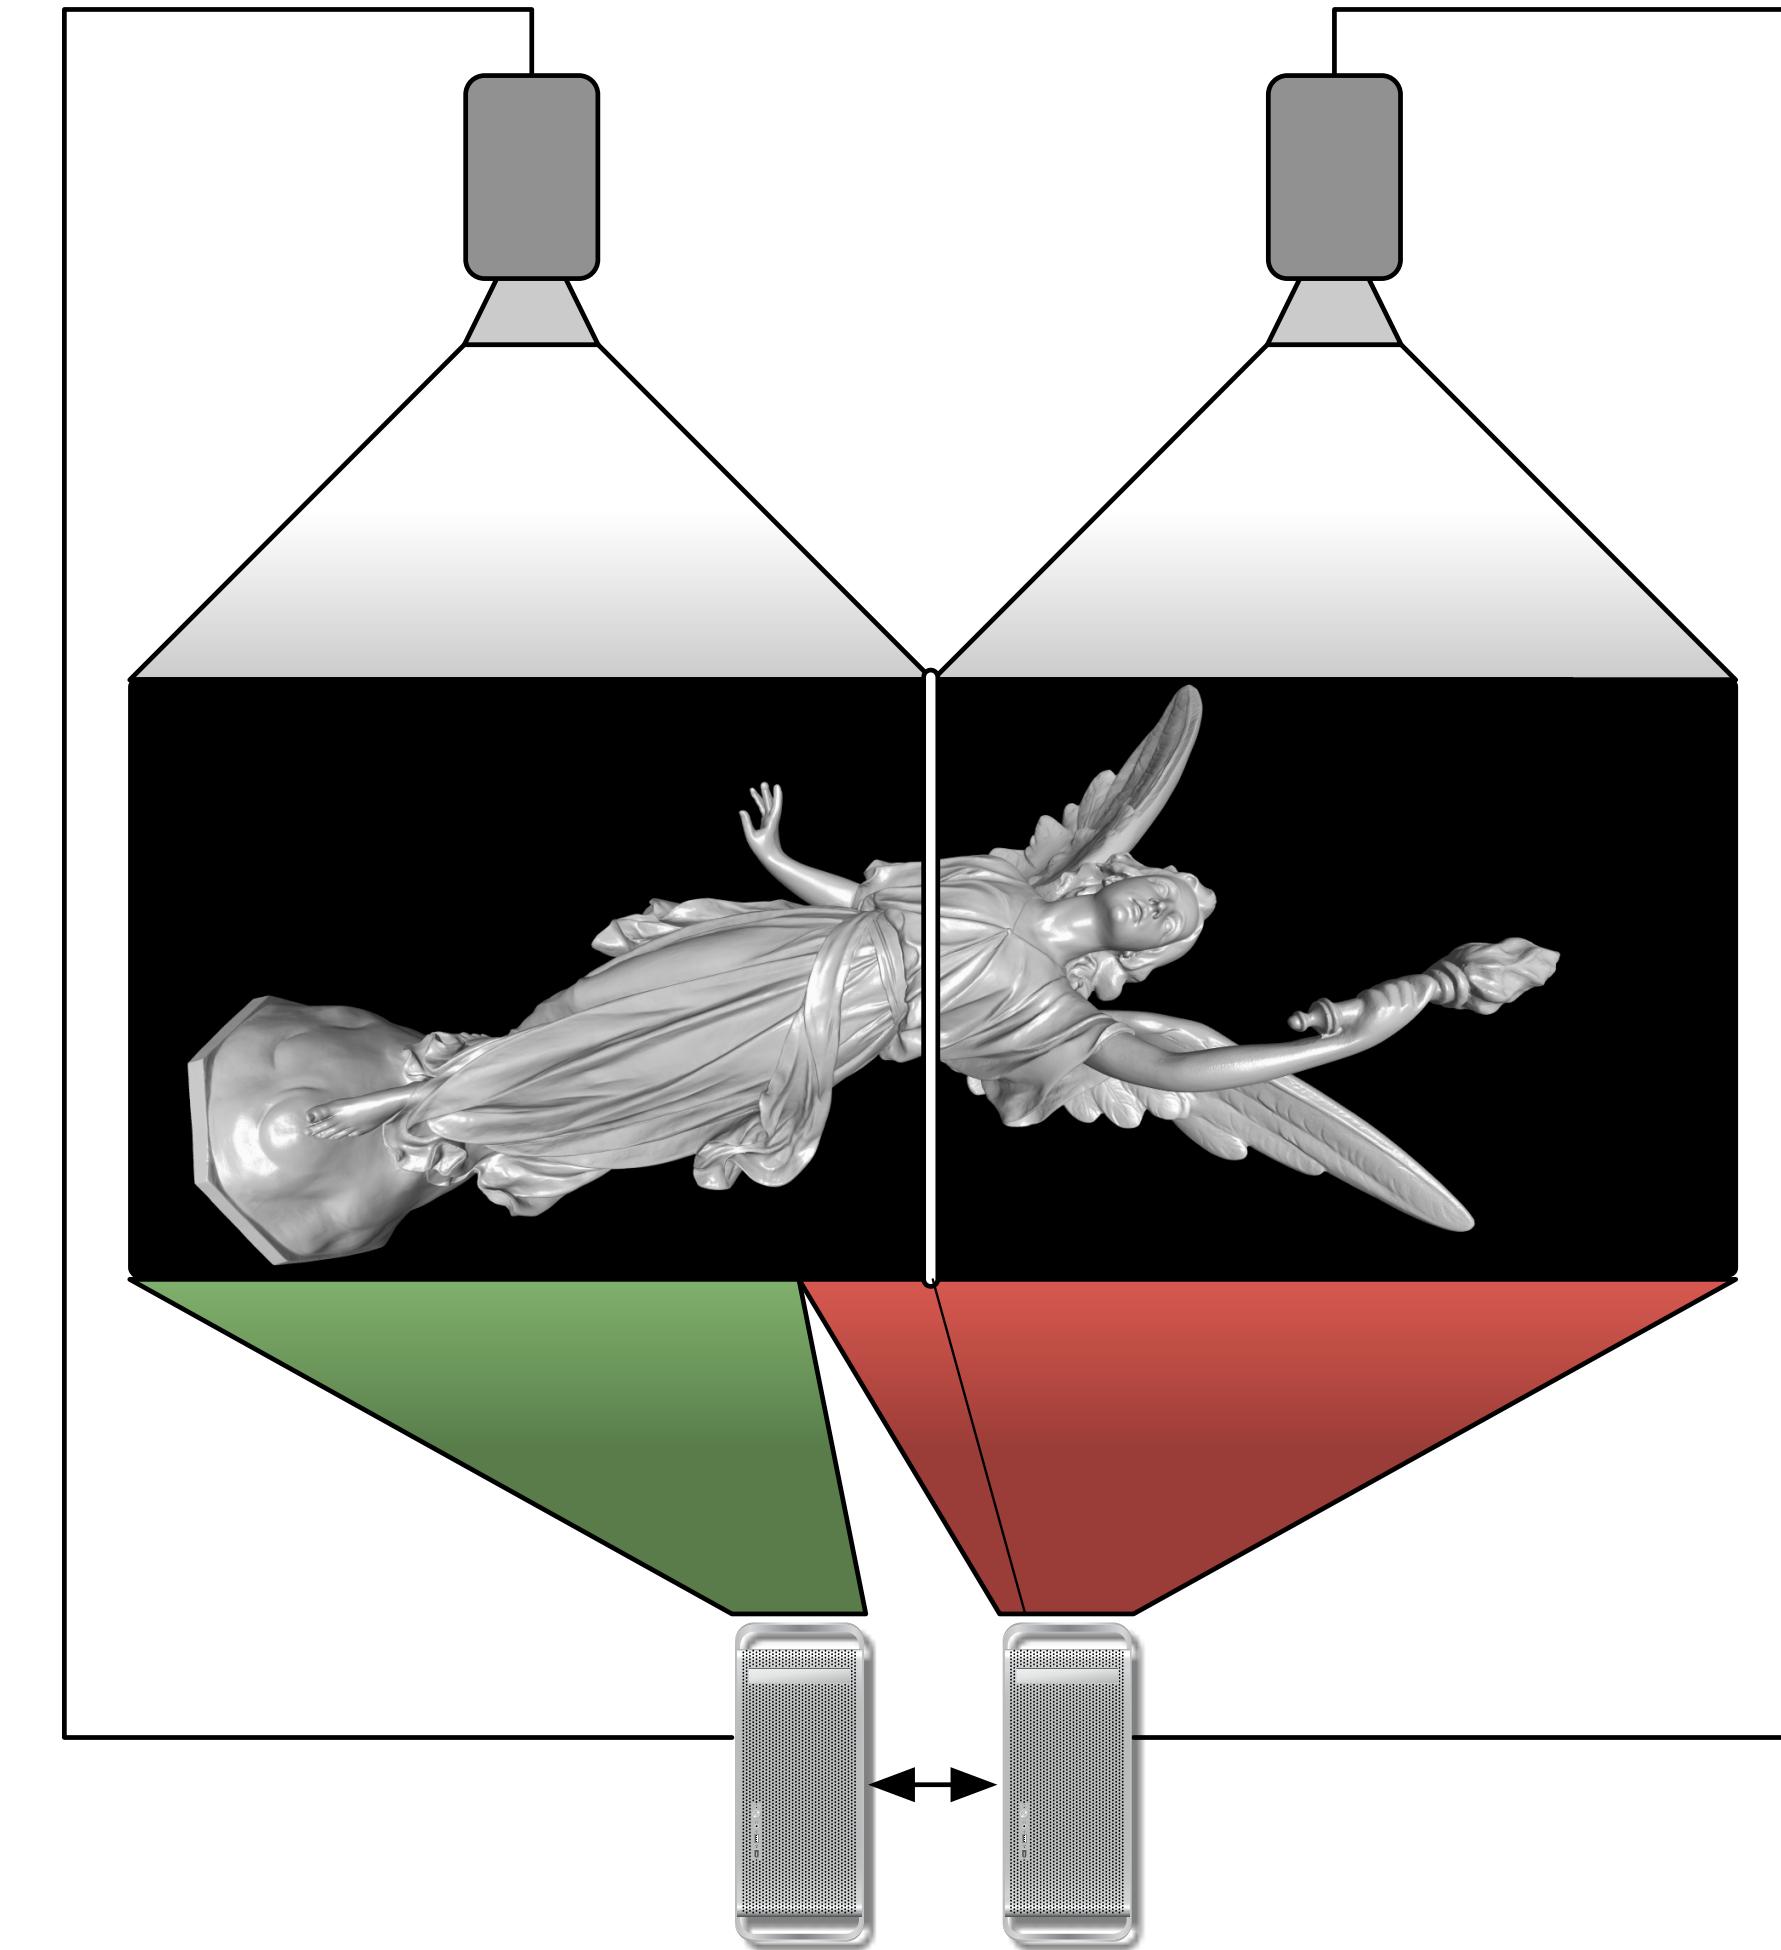


Single Display Load Balancing



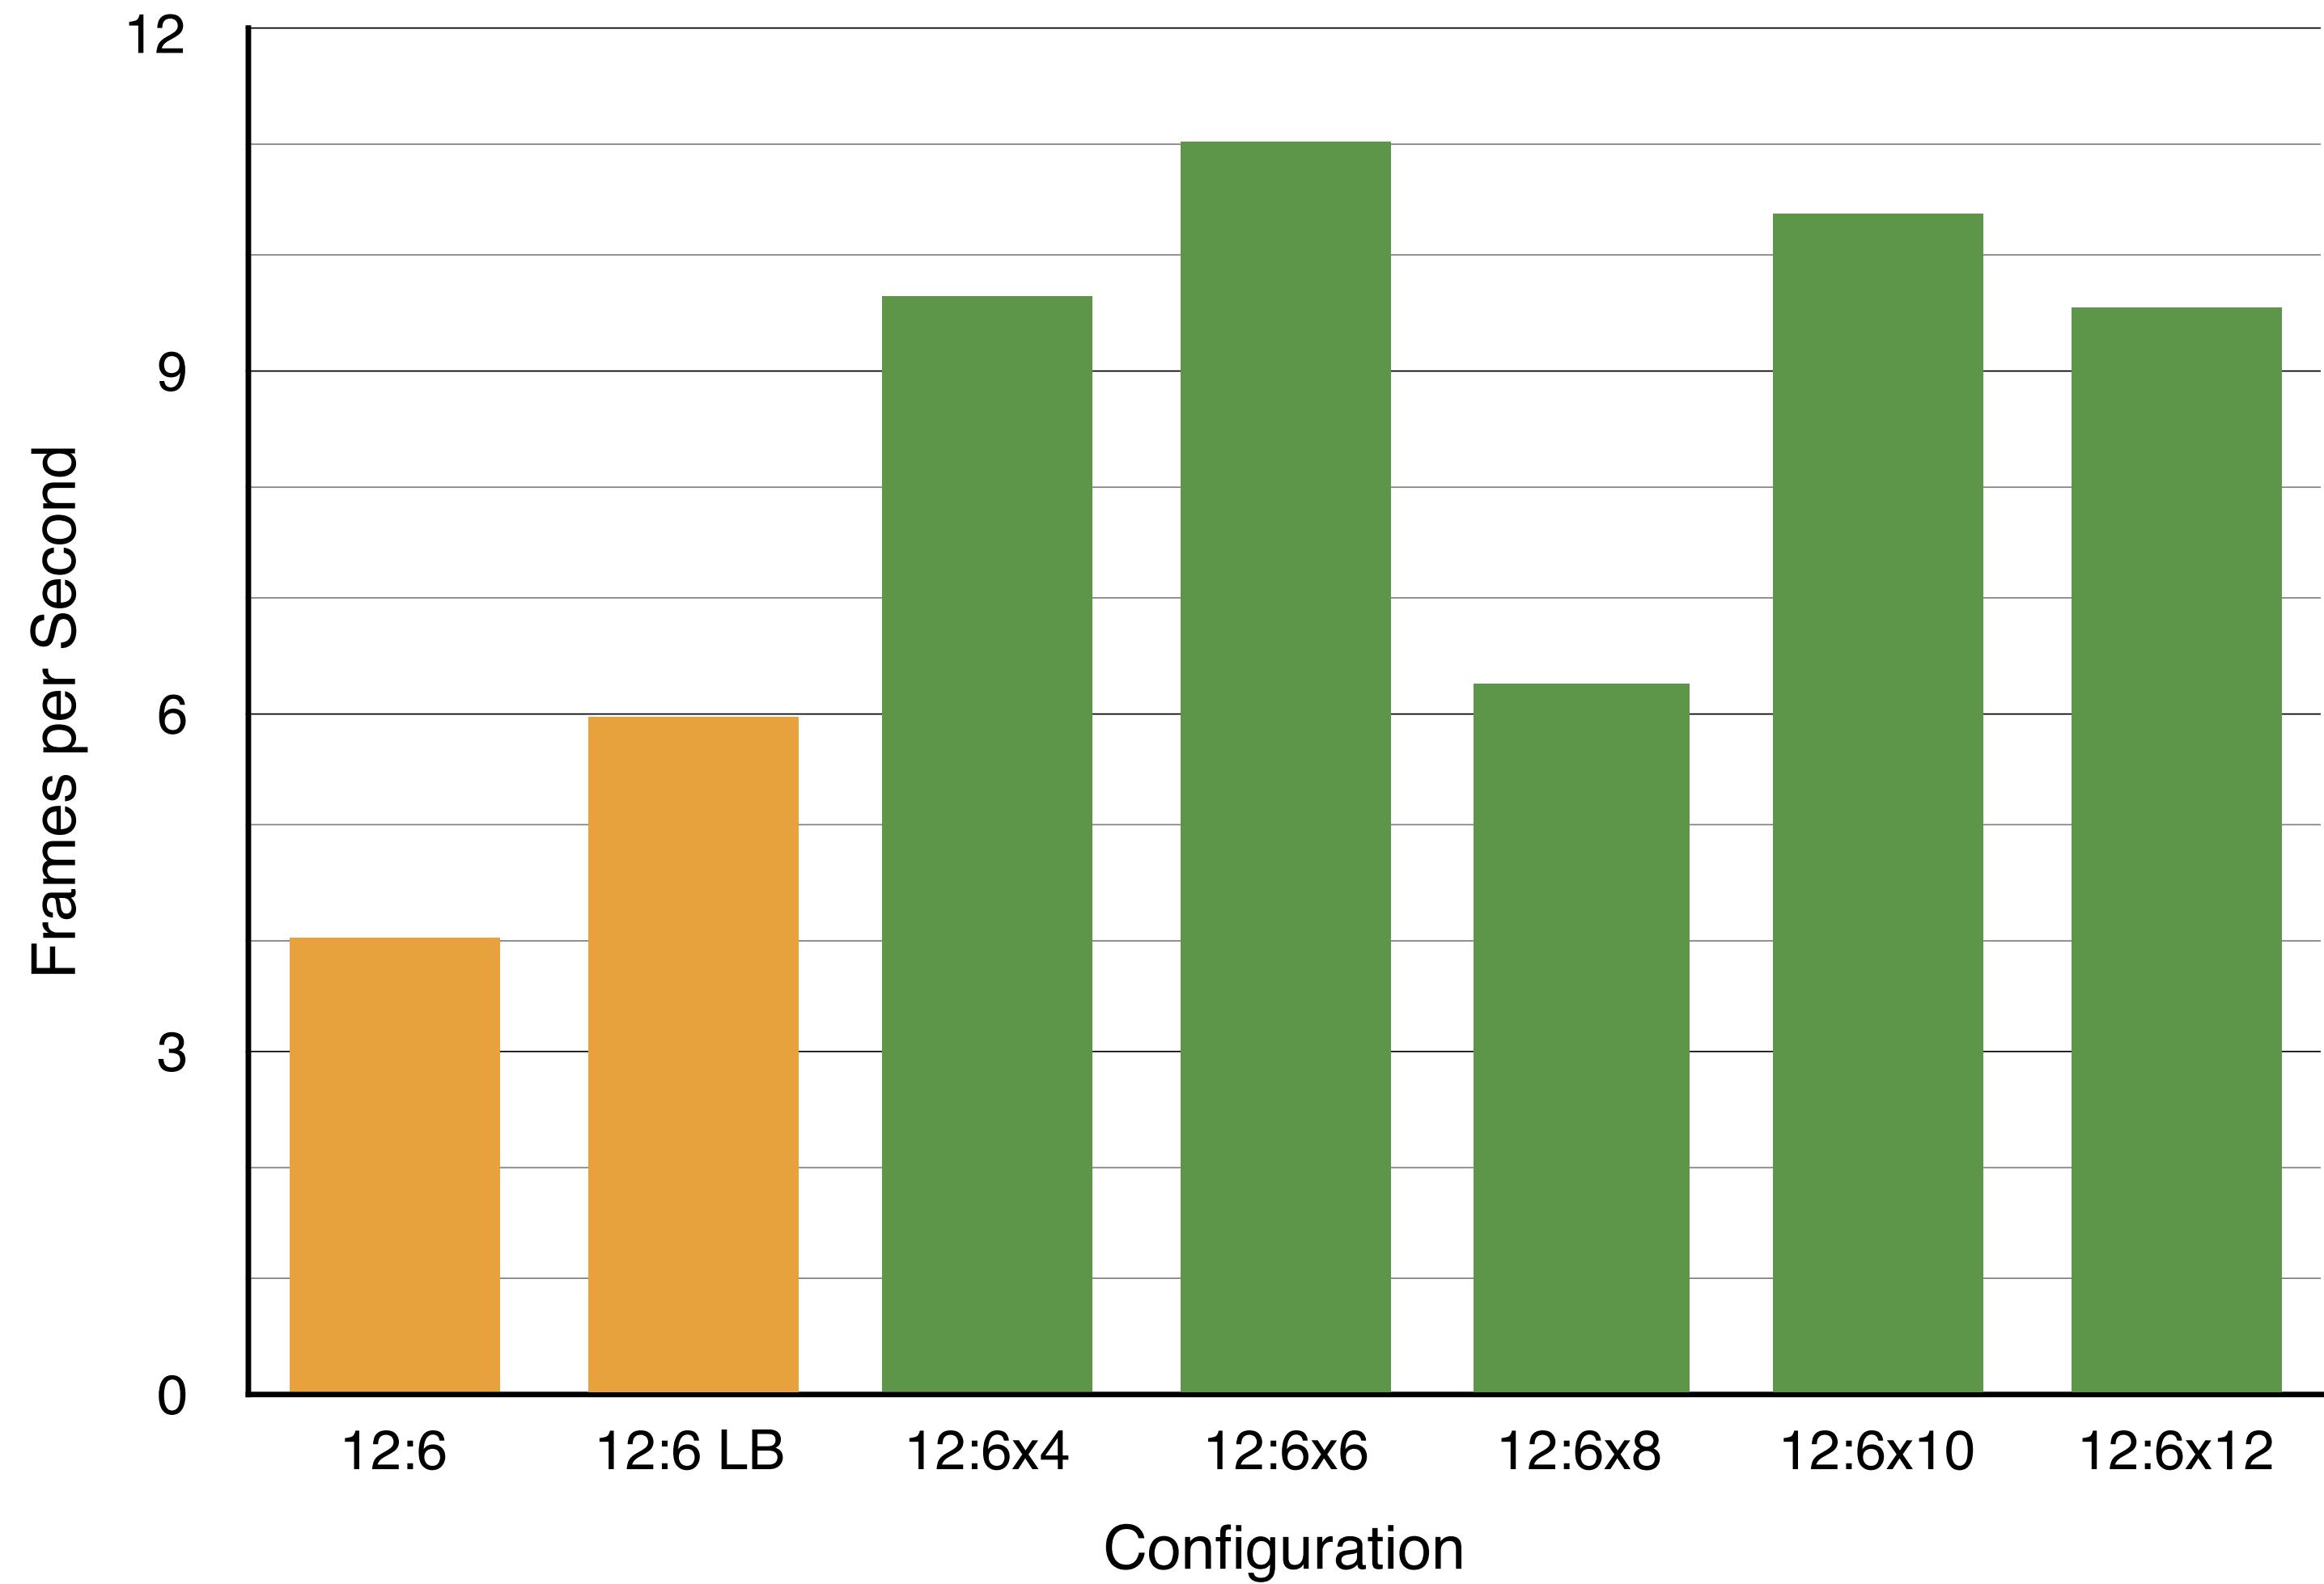
Cross-Display Load Balancing

- Assign m GPUs to n Displays ($m \geq n$)
- Per-display load balancer
 - ▶ Potential sources for each display
- Top-level view equalizer
 - ▶ Usage value per source
 - ▶ Source channels at most used twice
- Works for non-planar systems



Cross-Display Load Balancing

- 3x2 tiled display wall
- 12 GPU cluster
- Polygonal rendering
- Camera animation



Framework Architecture

How can we reduce end-to-end system latency?

How can we architect the parallel rendering framework to minimise synchronisation between threads?

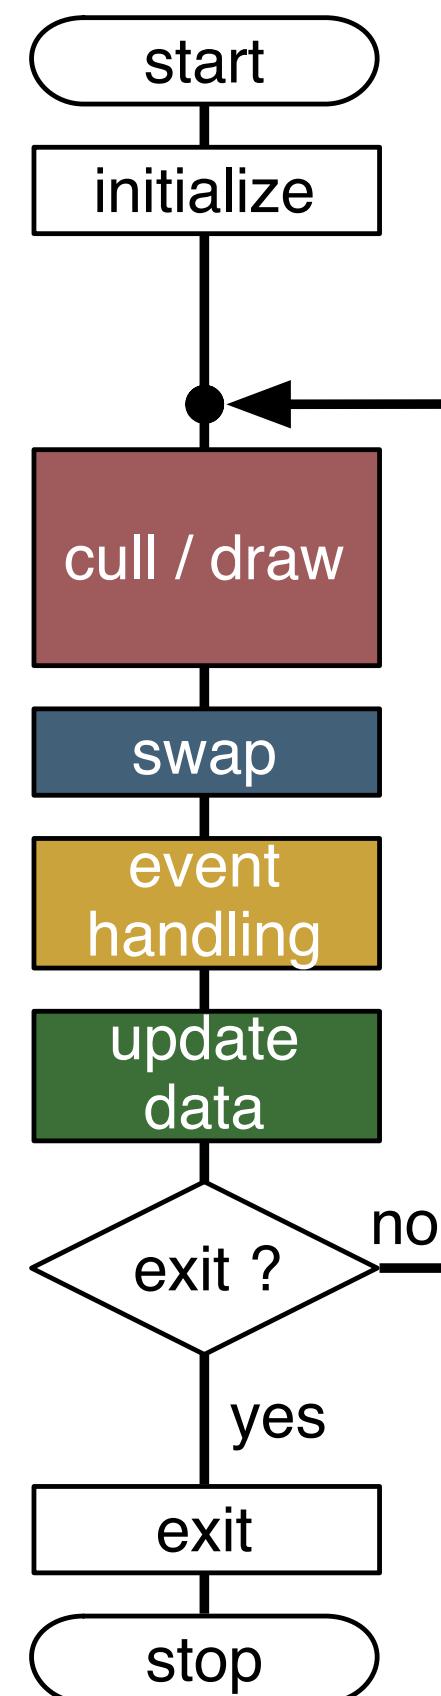
Which new algorithms will decrease the time needed to composite rendering results, in particular for sort-last rendering?

How can we maximise the impact of this research on large data scientists?

Minimally Invasive Architecture

Configuration vs Application

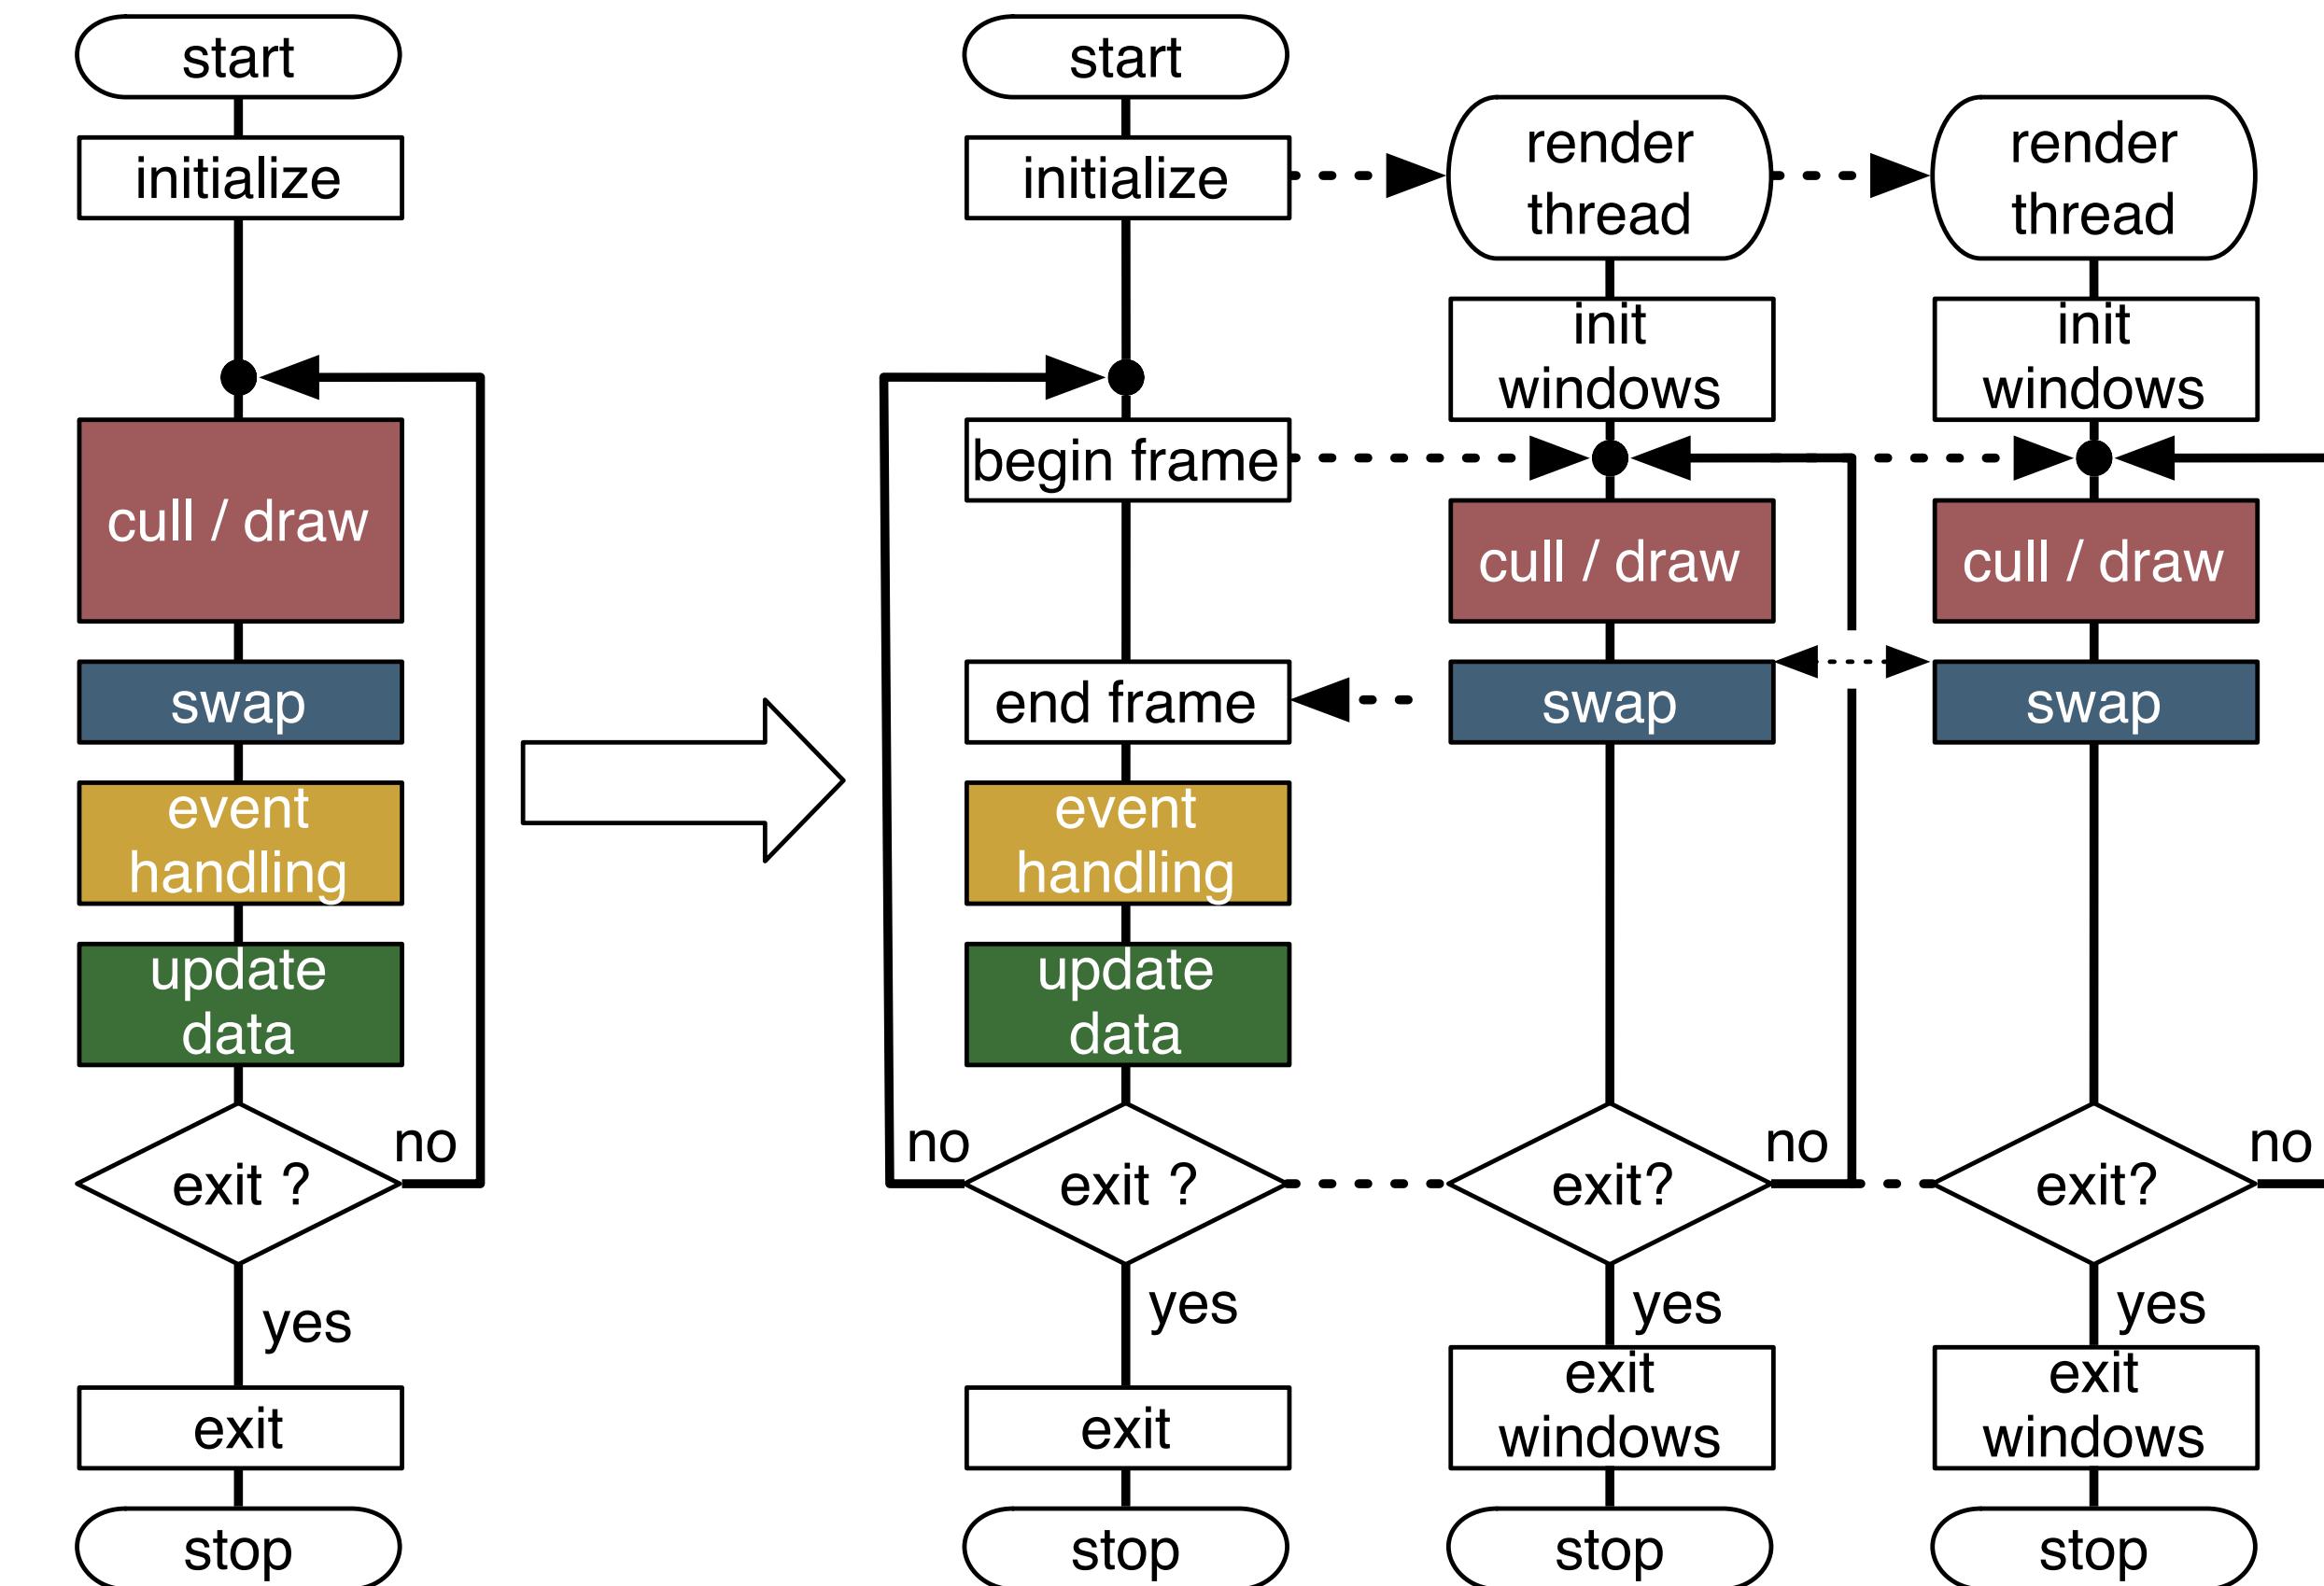
- Main Loop



Minimally Invasive Architecture

Configuration vs Application

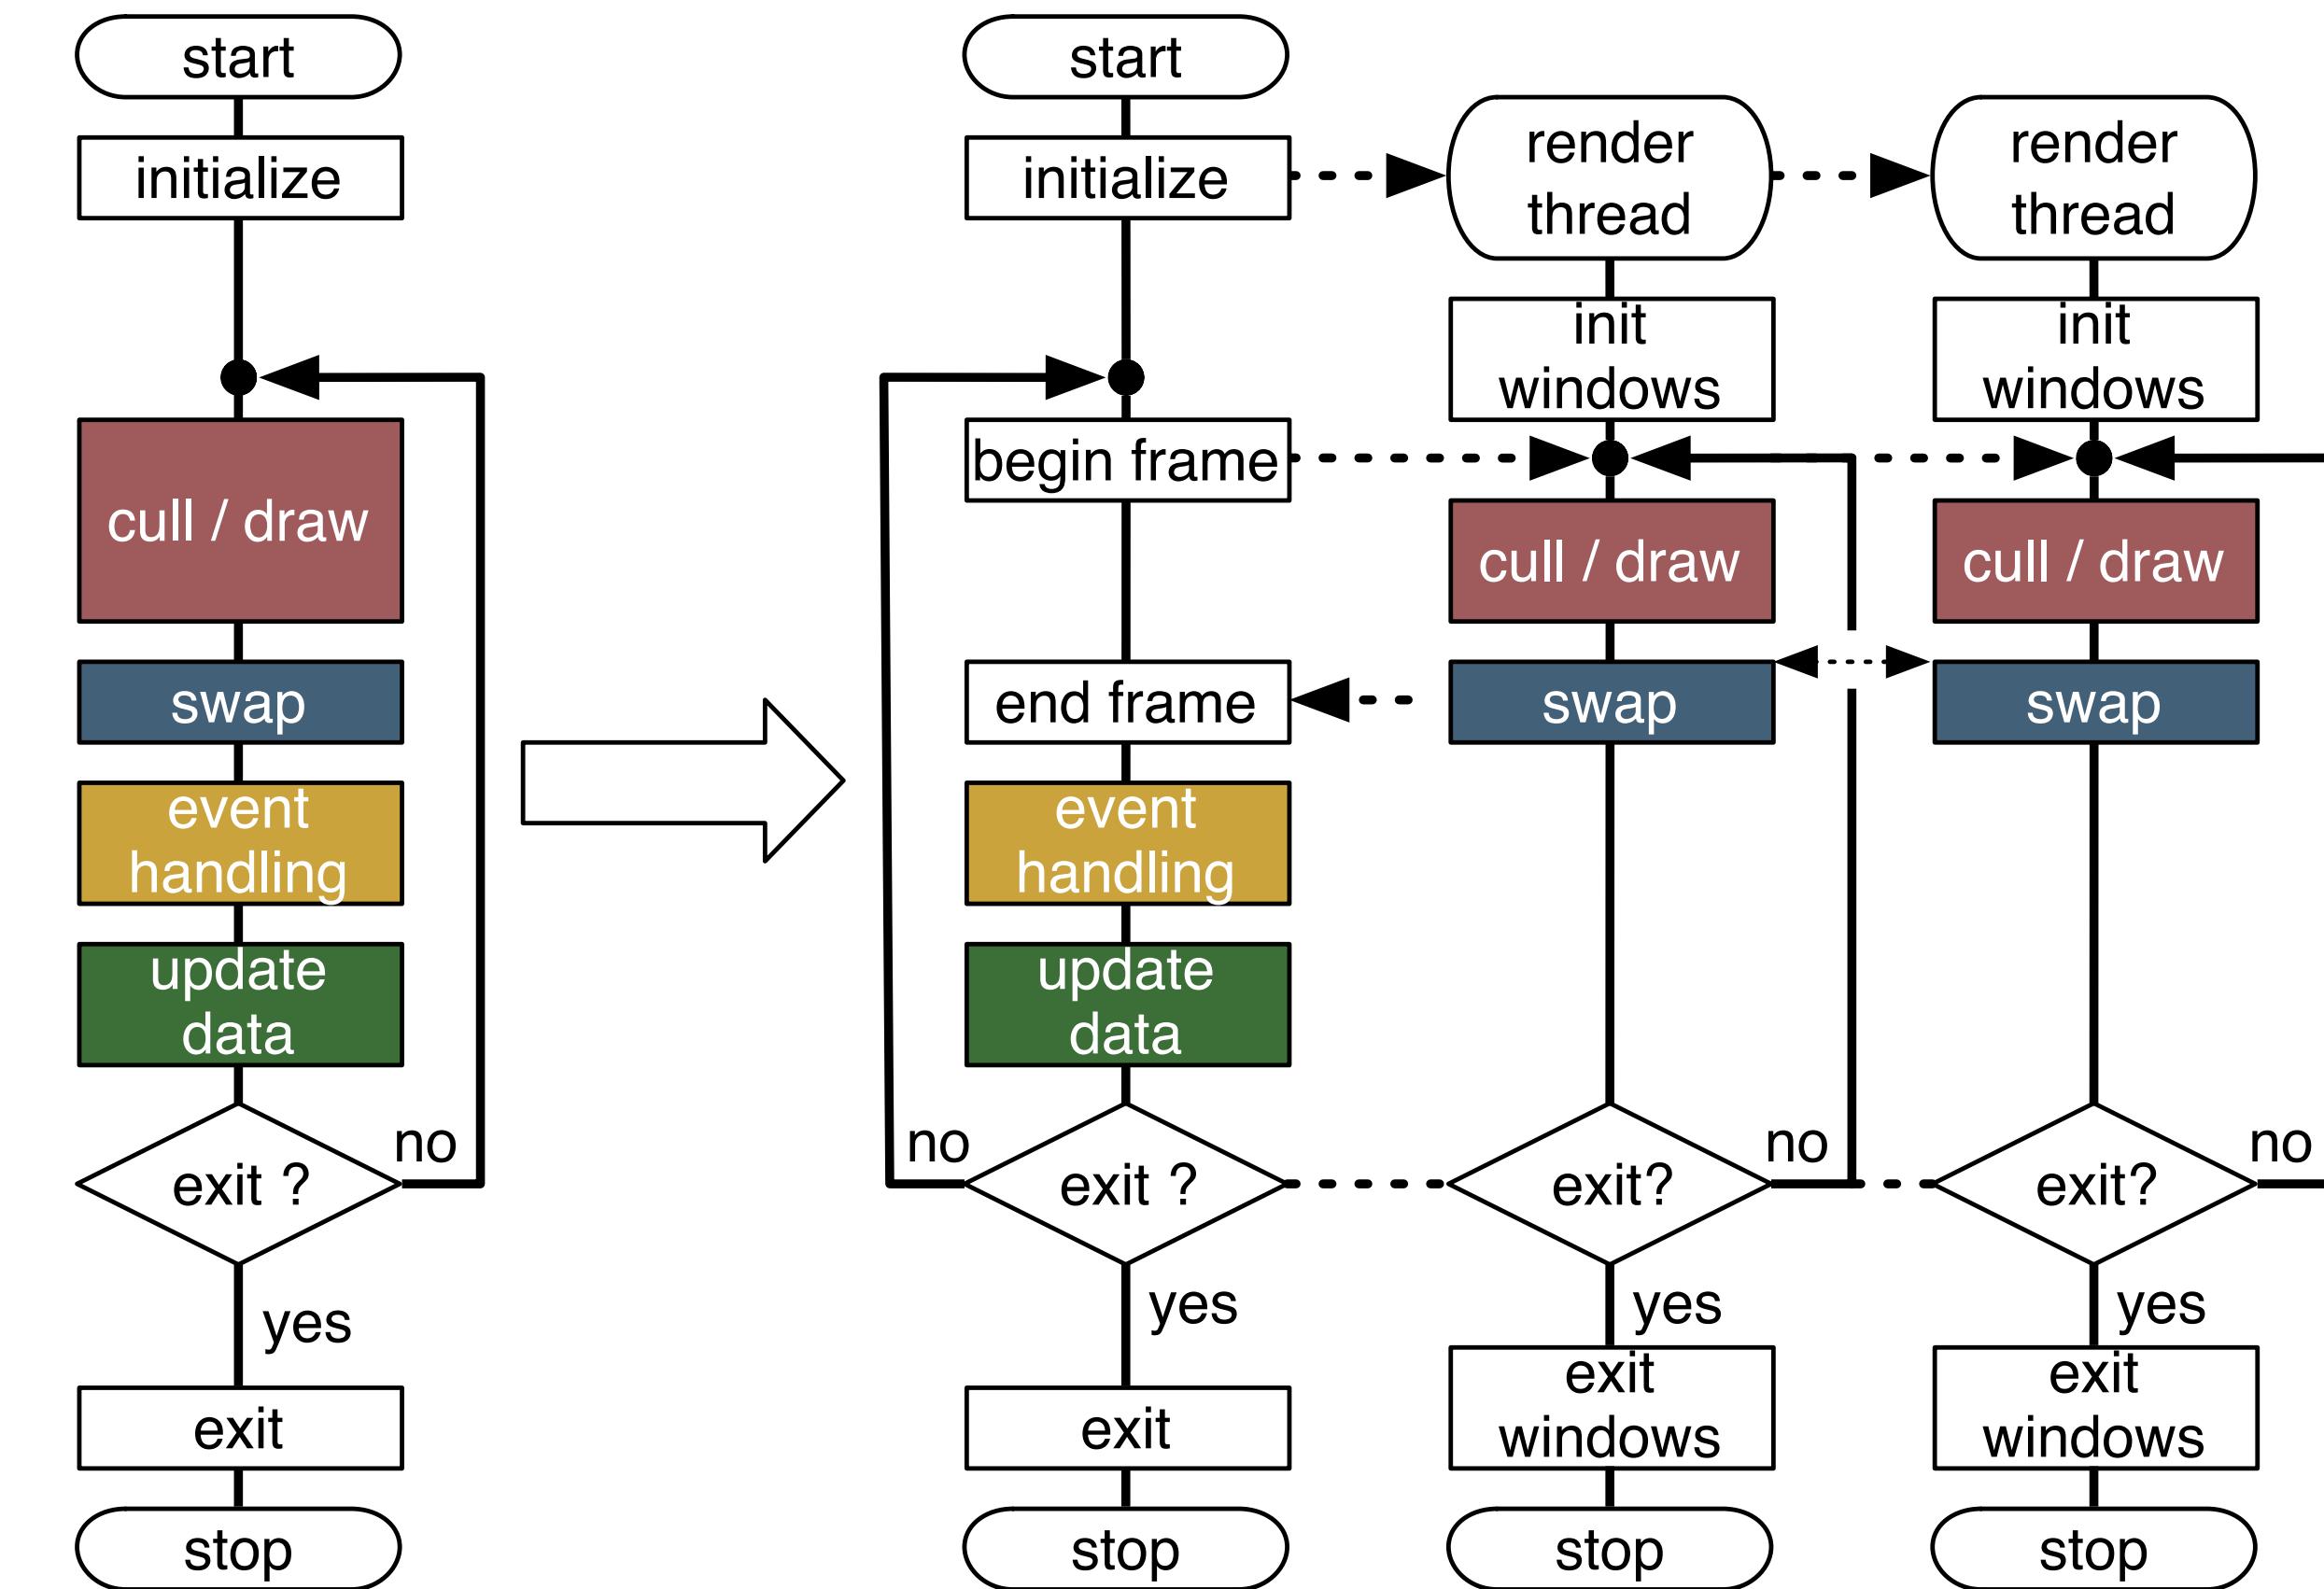
- Main Loop
- Render Client



Minimally Invasive Architecture

Configuration vs Application

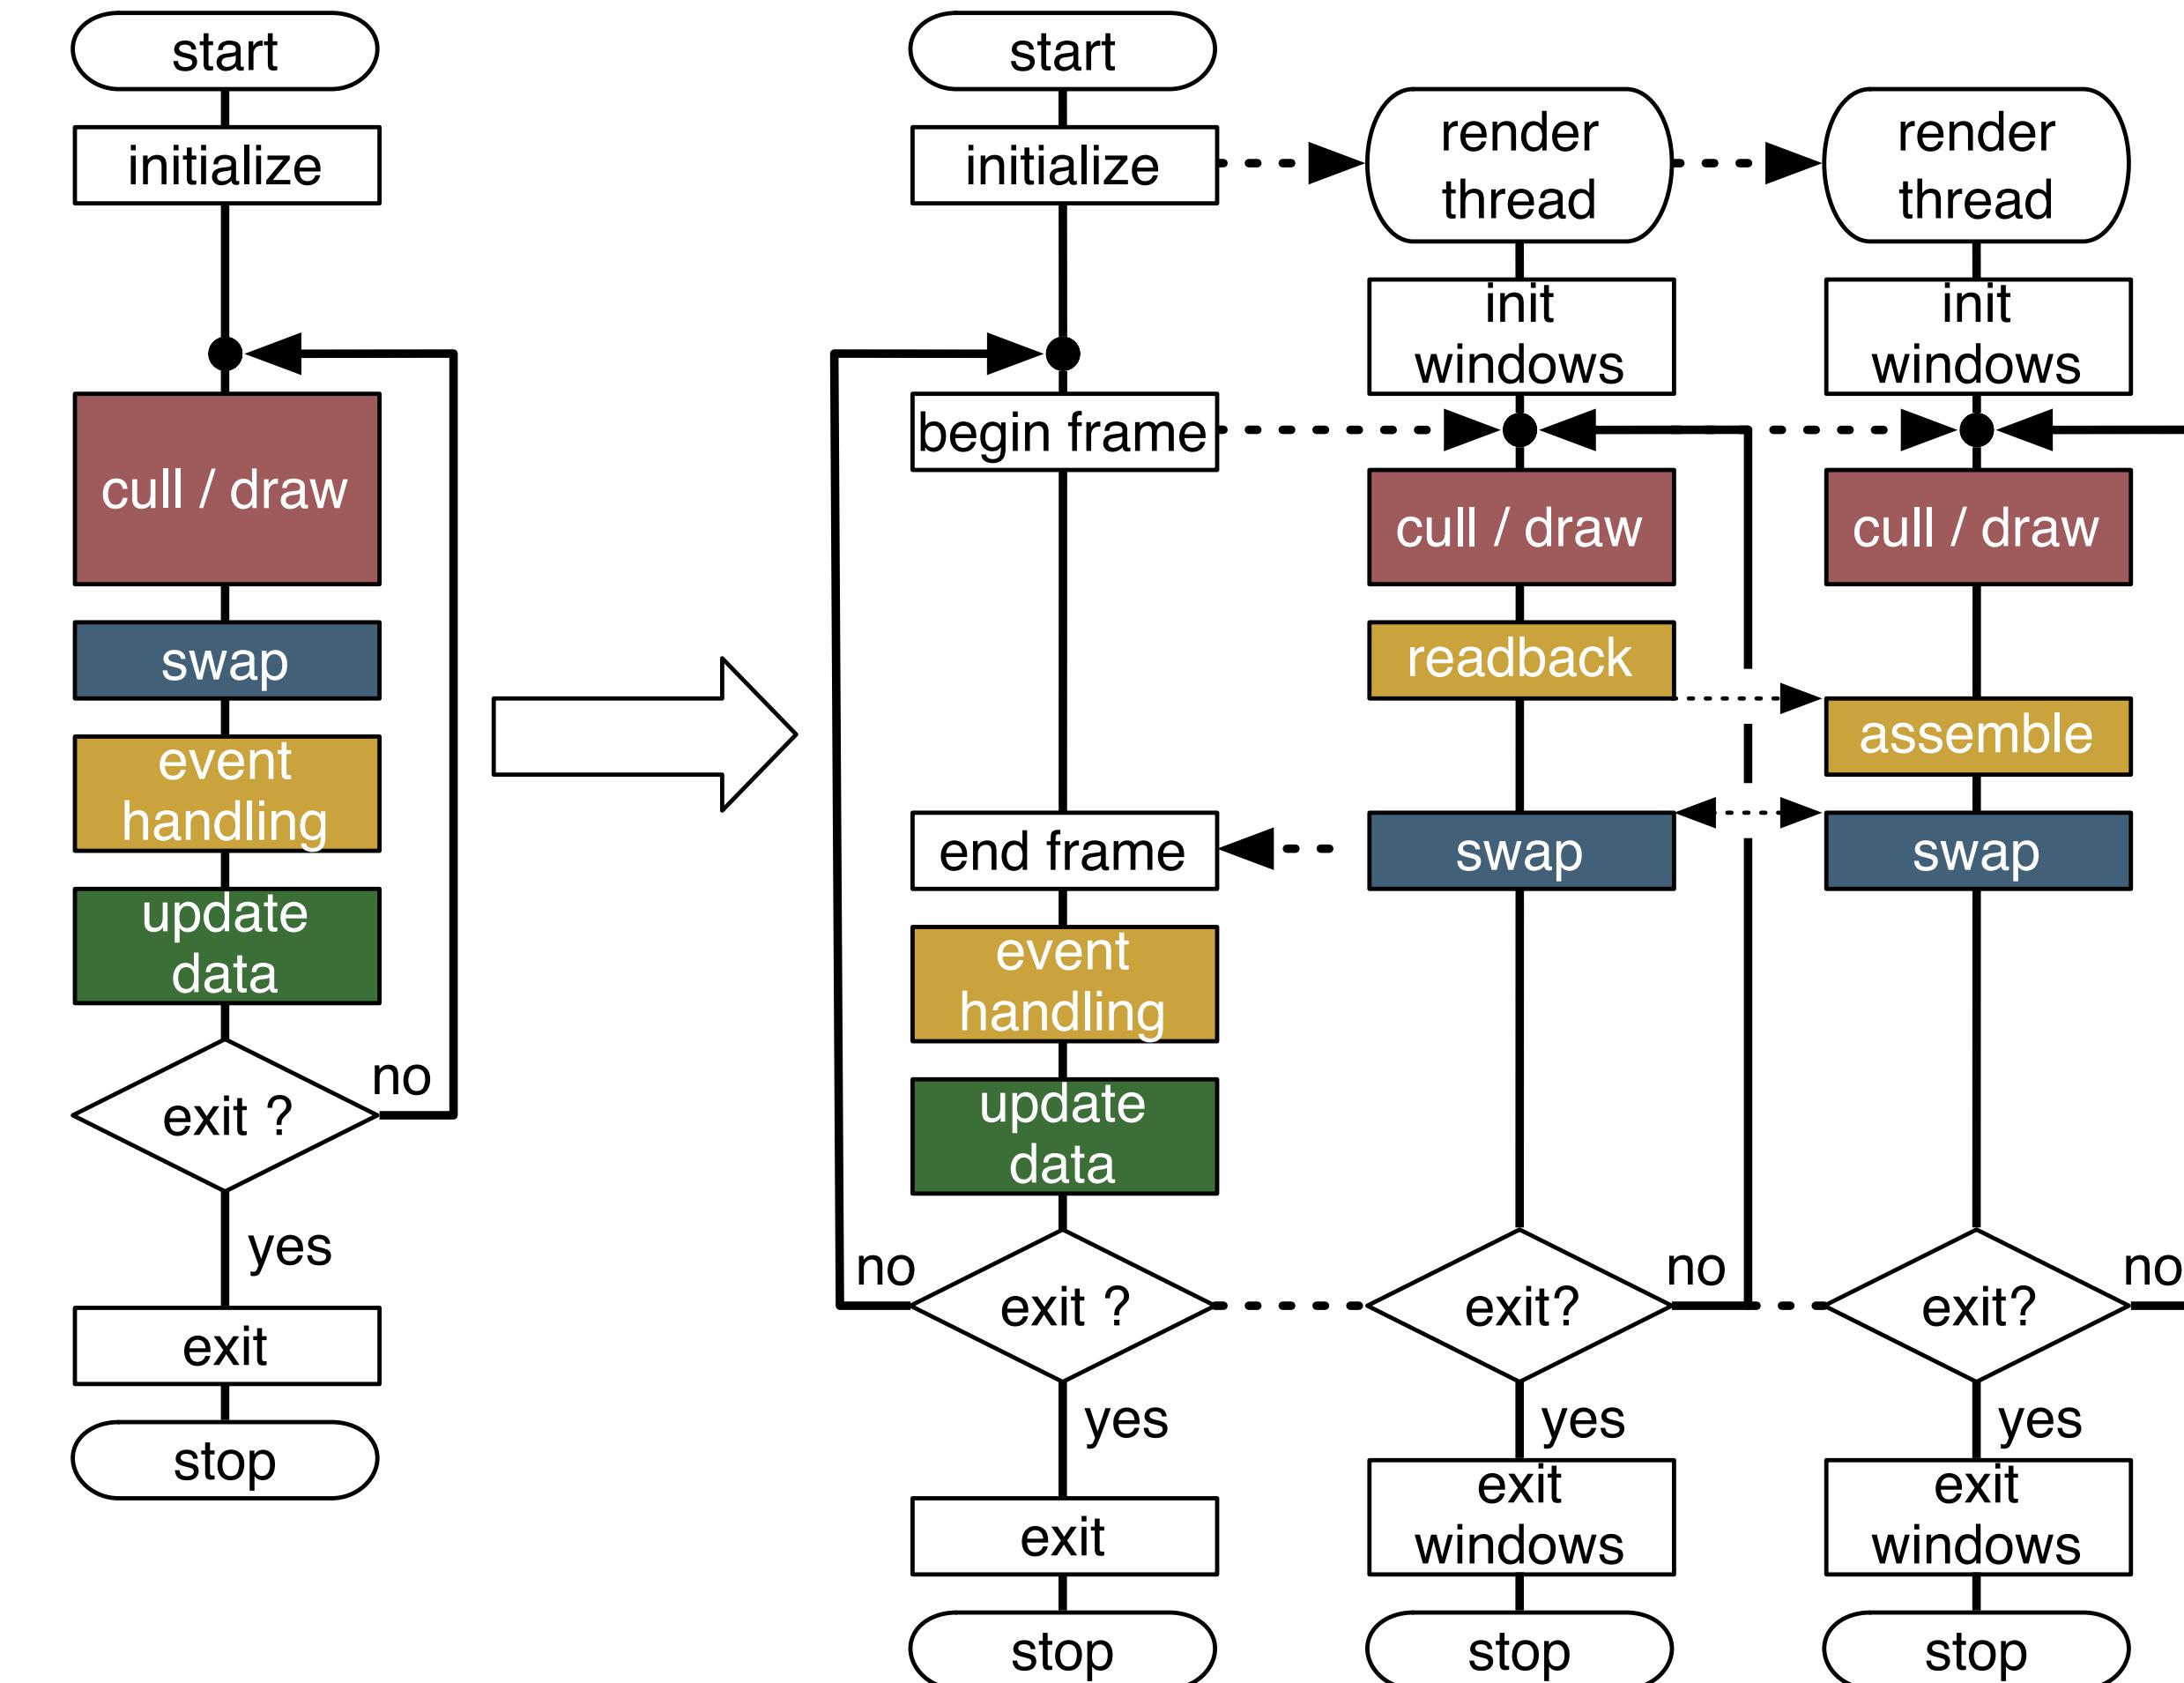
- Main Loop
- Render Client
- RenderContext
 - ▶ Viewport
 - ▶ Frustum
 - ▶ Head Transformation
 - ▶ Buffer
 - ▶ Range



Minimally Invasive Architecture

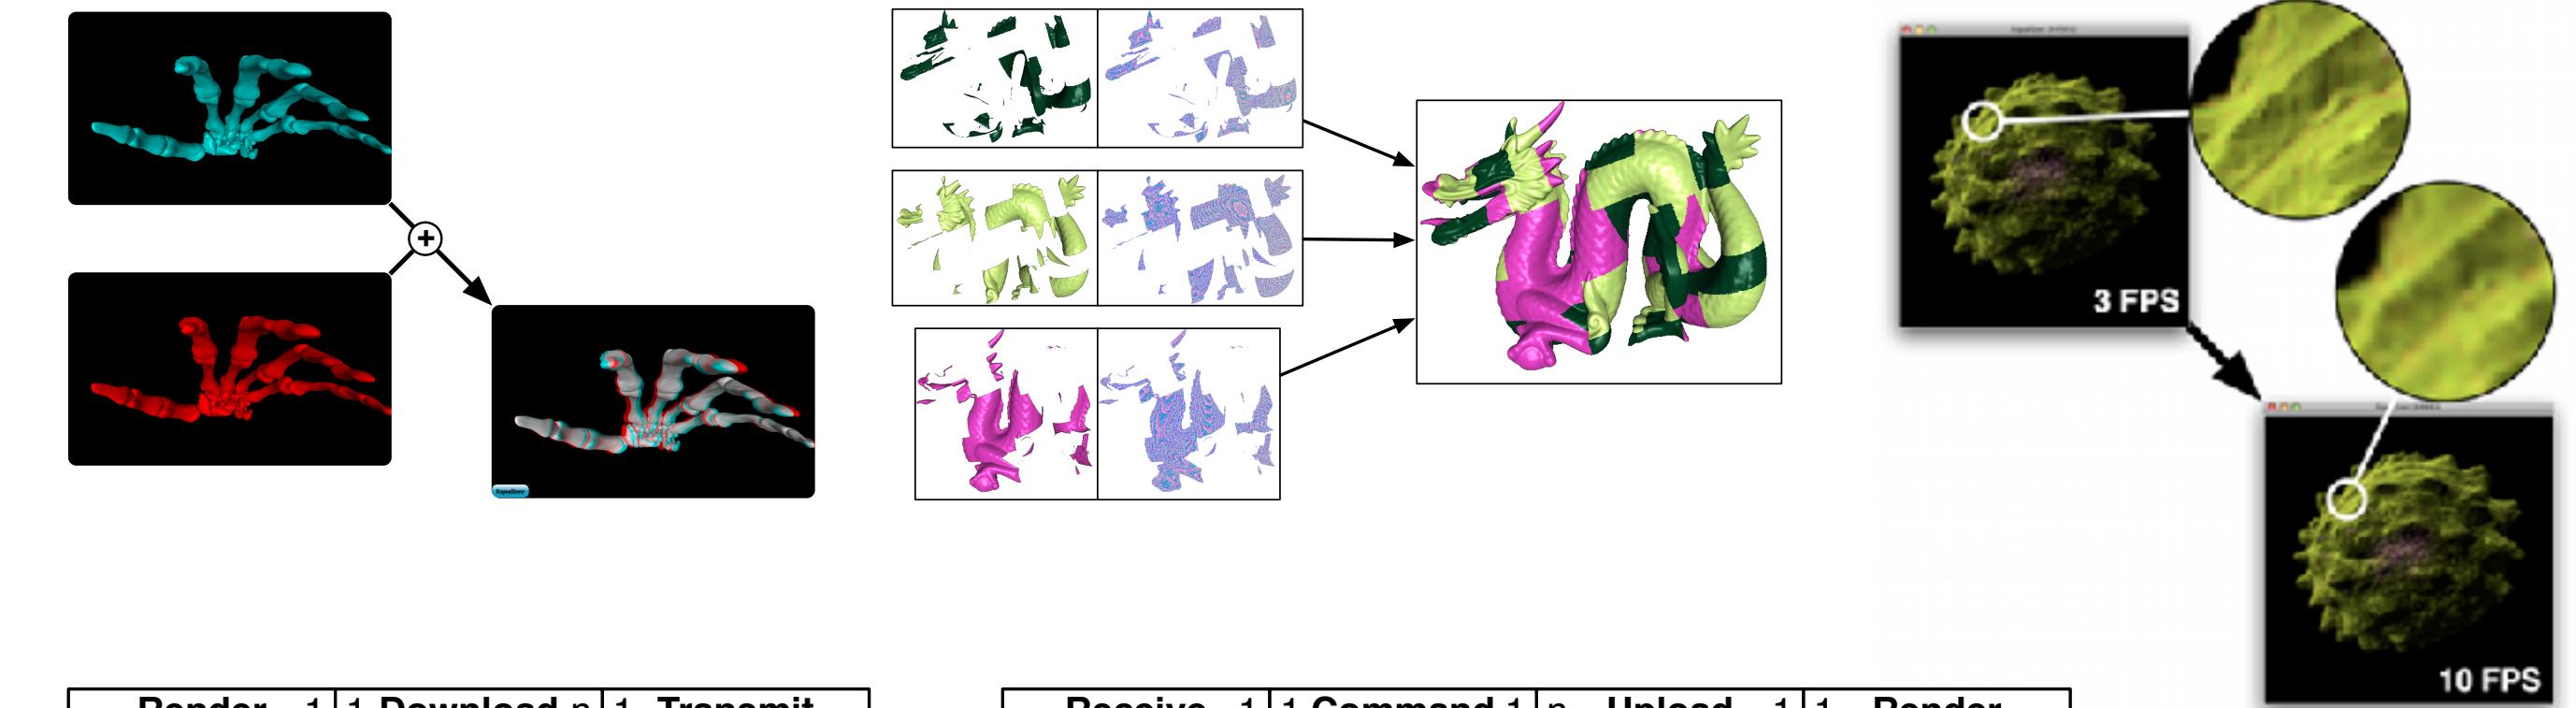
Configuration vs Application

- Main Loop
- Render Client
- RenderContext
 - ▶ Viewport
 - ▶ Frustum
 - ▶ Head Transformation
 - ▶ Buffer
 - ▶ Range

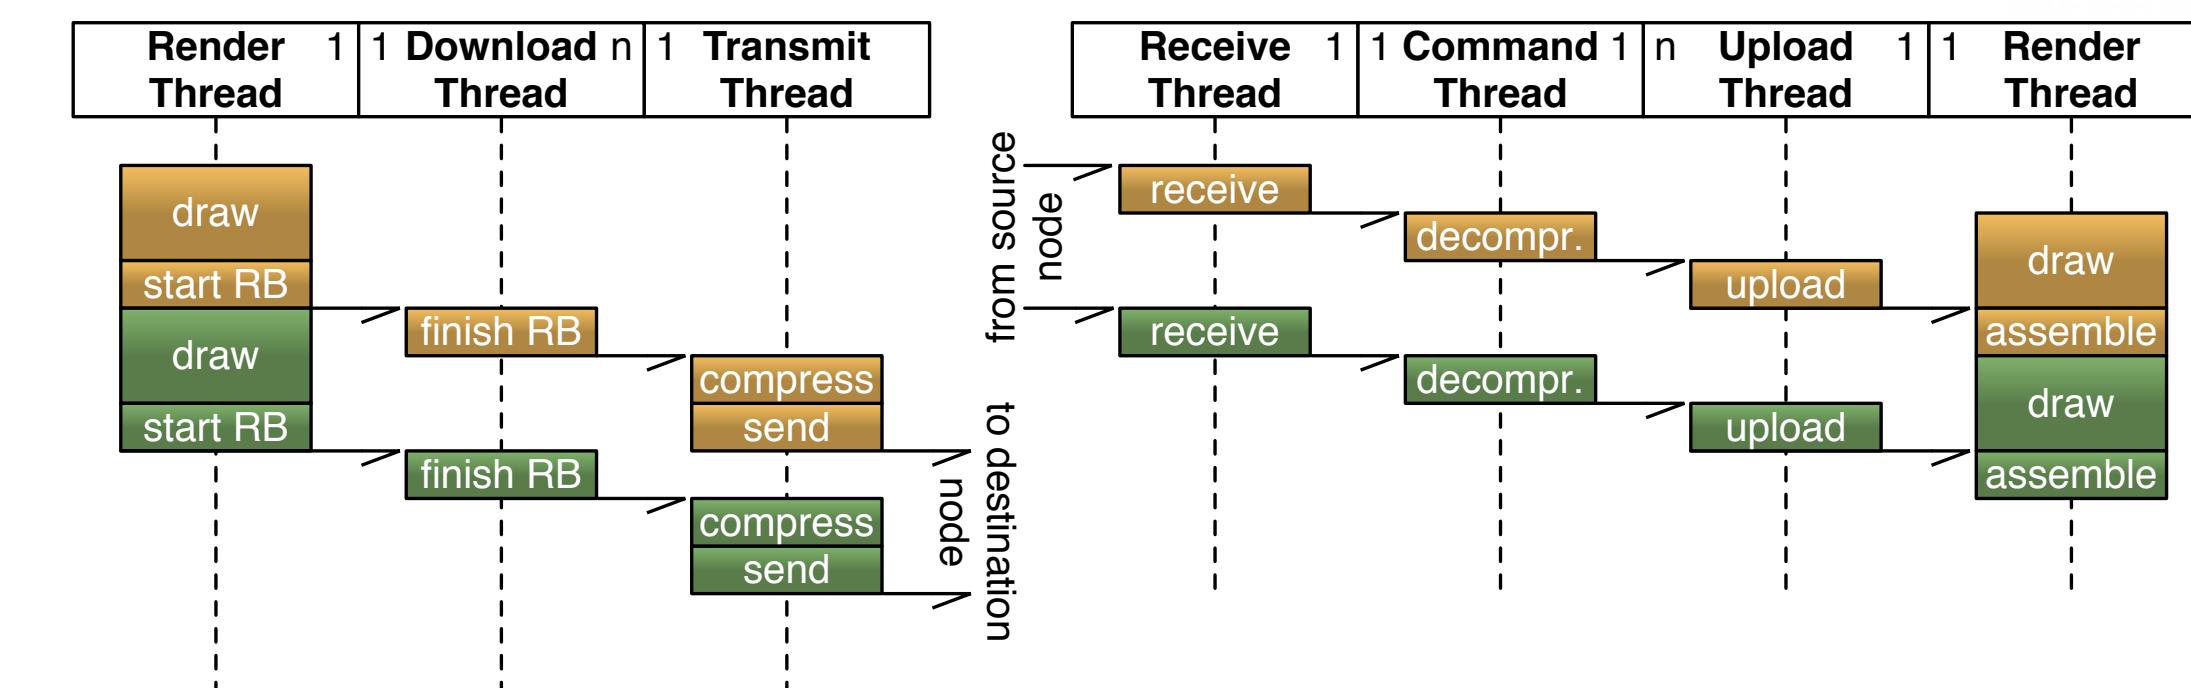


Features

Eight Decomposition Modes



Six Equalizers



Whole System Optimization

Full Application Support

Eilemann, S., Steiner, D., and Pajarola, R. (2018). Equalizer 2.0 – Convergence of a Parallel Rendering Framework. *IEEE Transactions on Visualization and Computer Graphics*

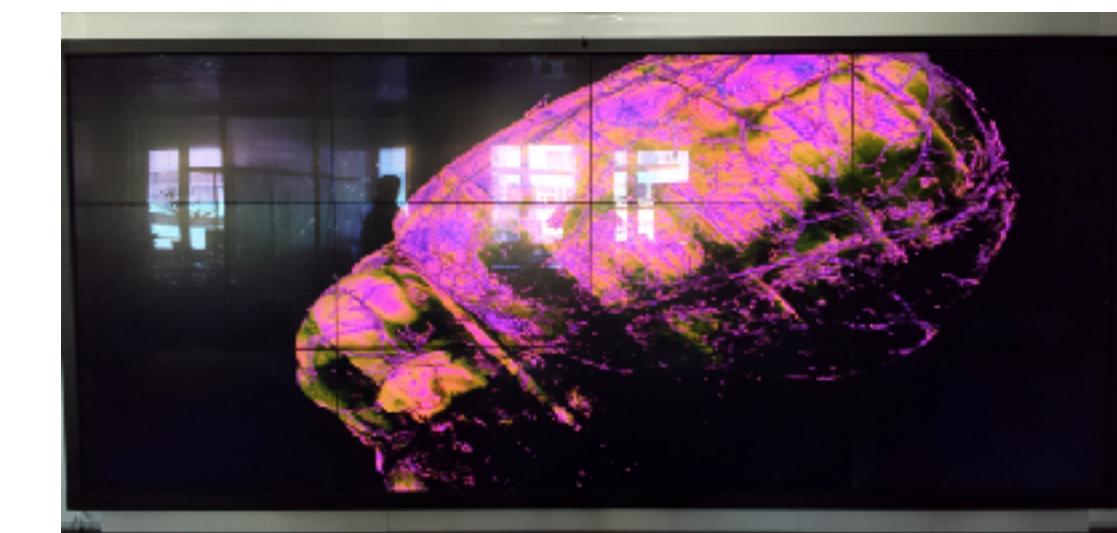
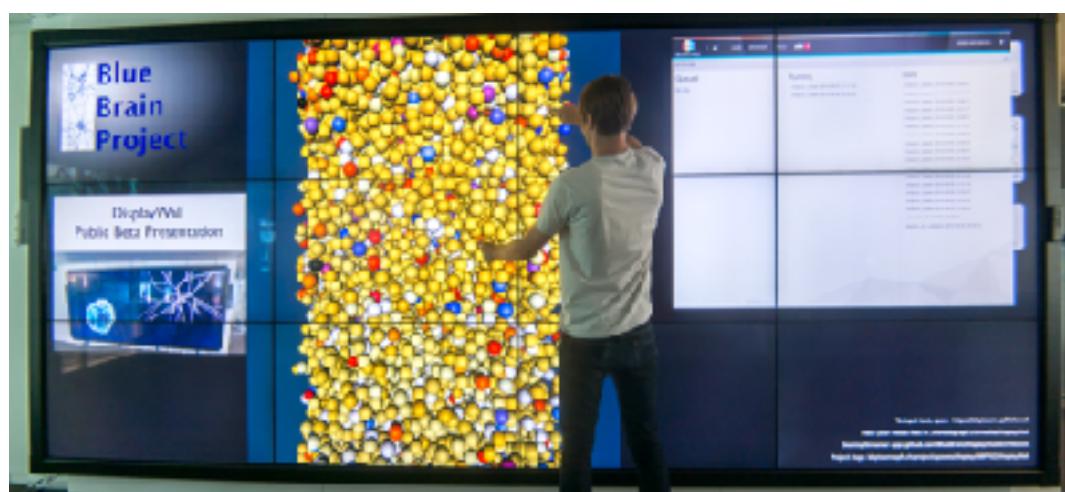
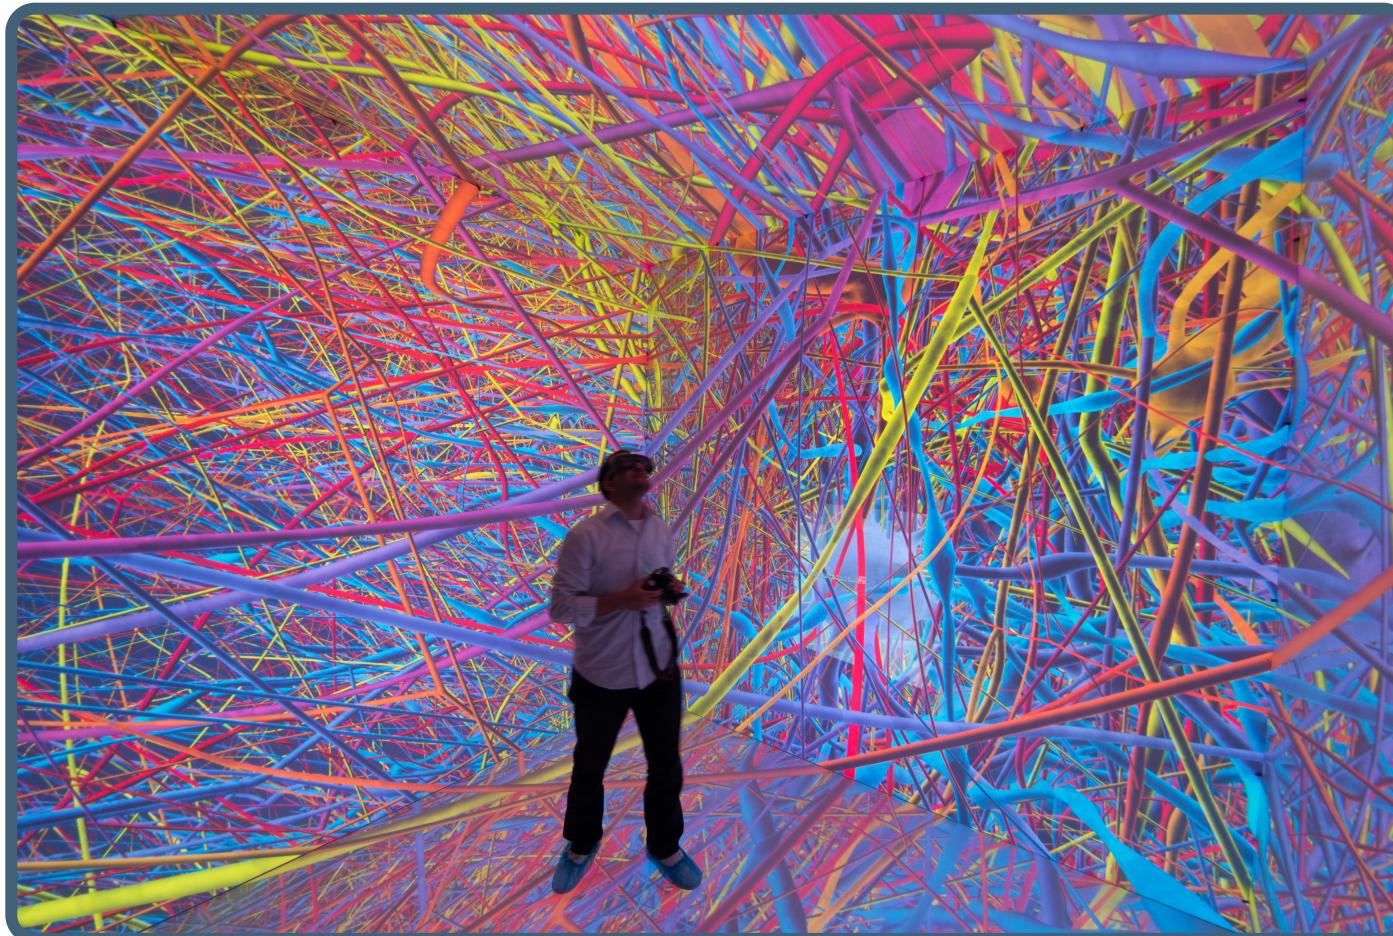
Eilemann, S., Bilgili, A., Abdellah, M., Hernando, J., Makhinya, M., Pajarola, R., and Schürmann, F. (2012). Parallel Rendering on Hybrid Multi-GPU Clusters. *Eurographics Symposium on Parallel Graphics and Visualization*

Eilemann, S. and Pajarola, R. (2007). Direct send compositing for parallel sort-last rendering. *Eurographics Symposium on Parallel Graphics and Visualization*.

Summary

- Minimally invasive, flexible framework architecture
- Full-fledged application support
- Whole system optimization
- Reliable multicast protocol
- Compression techniques for interactive use
- Novel load balancing algorithms for single and multi display systems

Applications



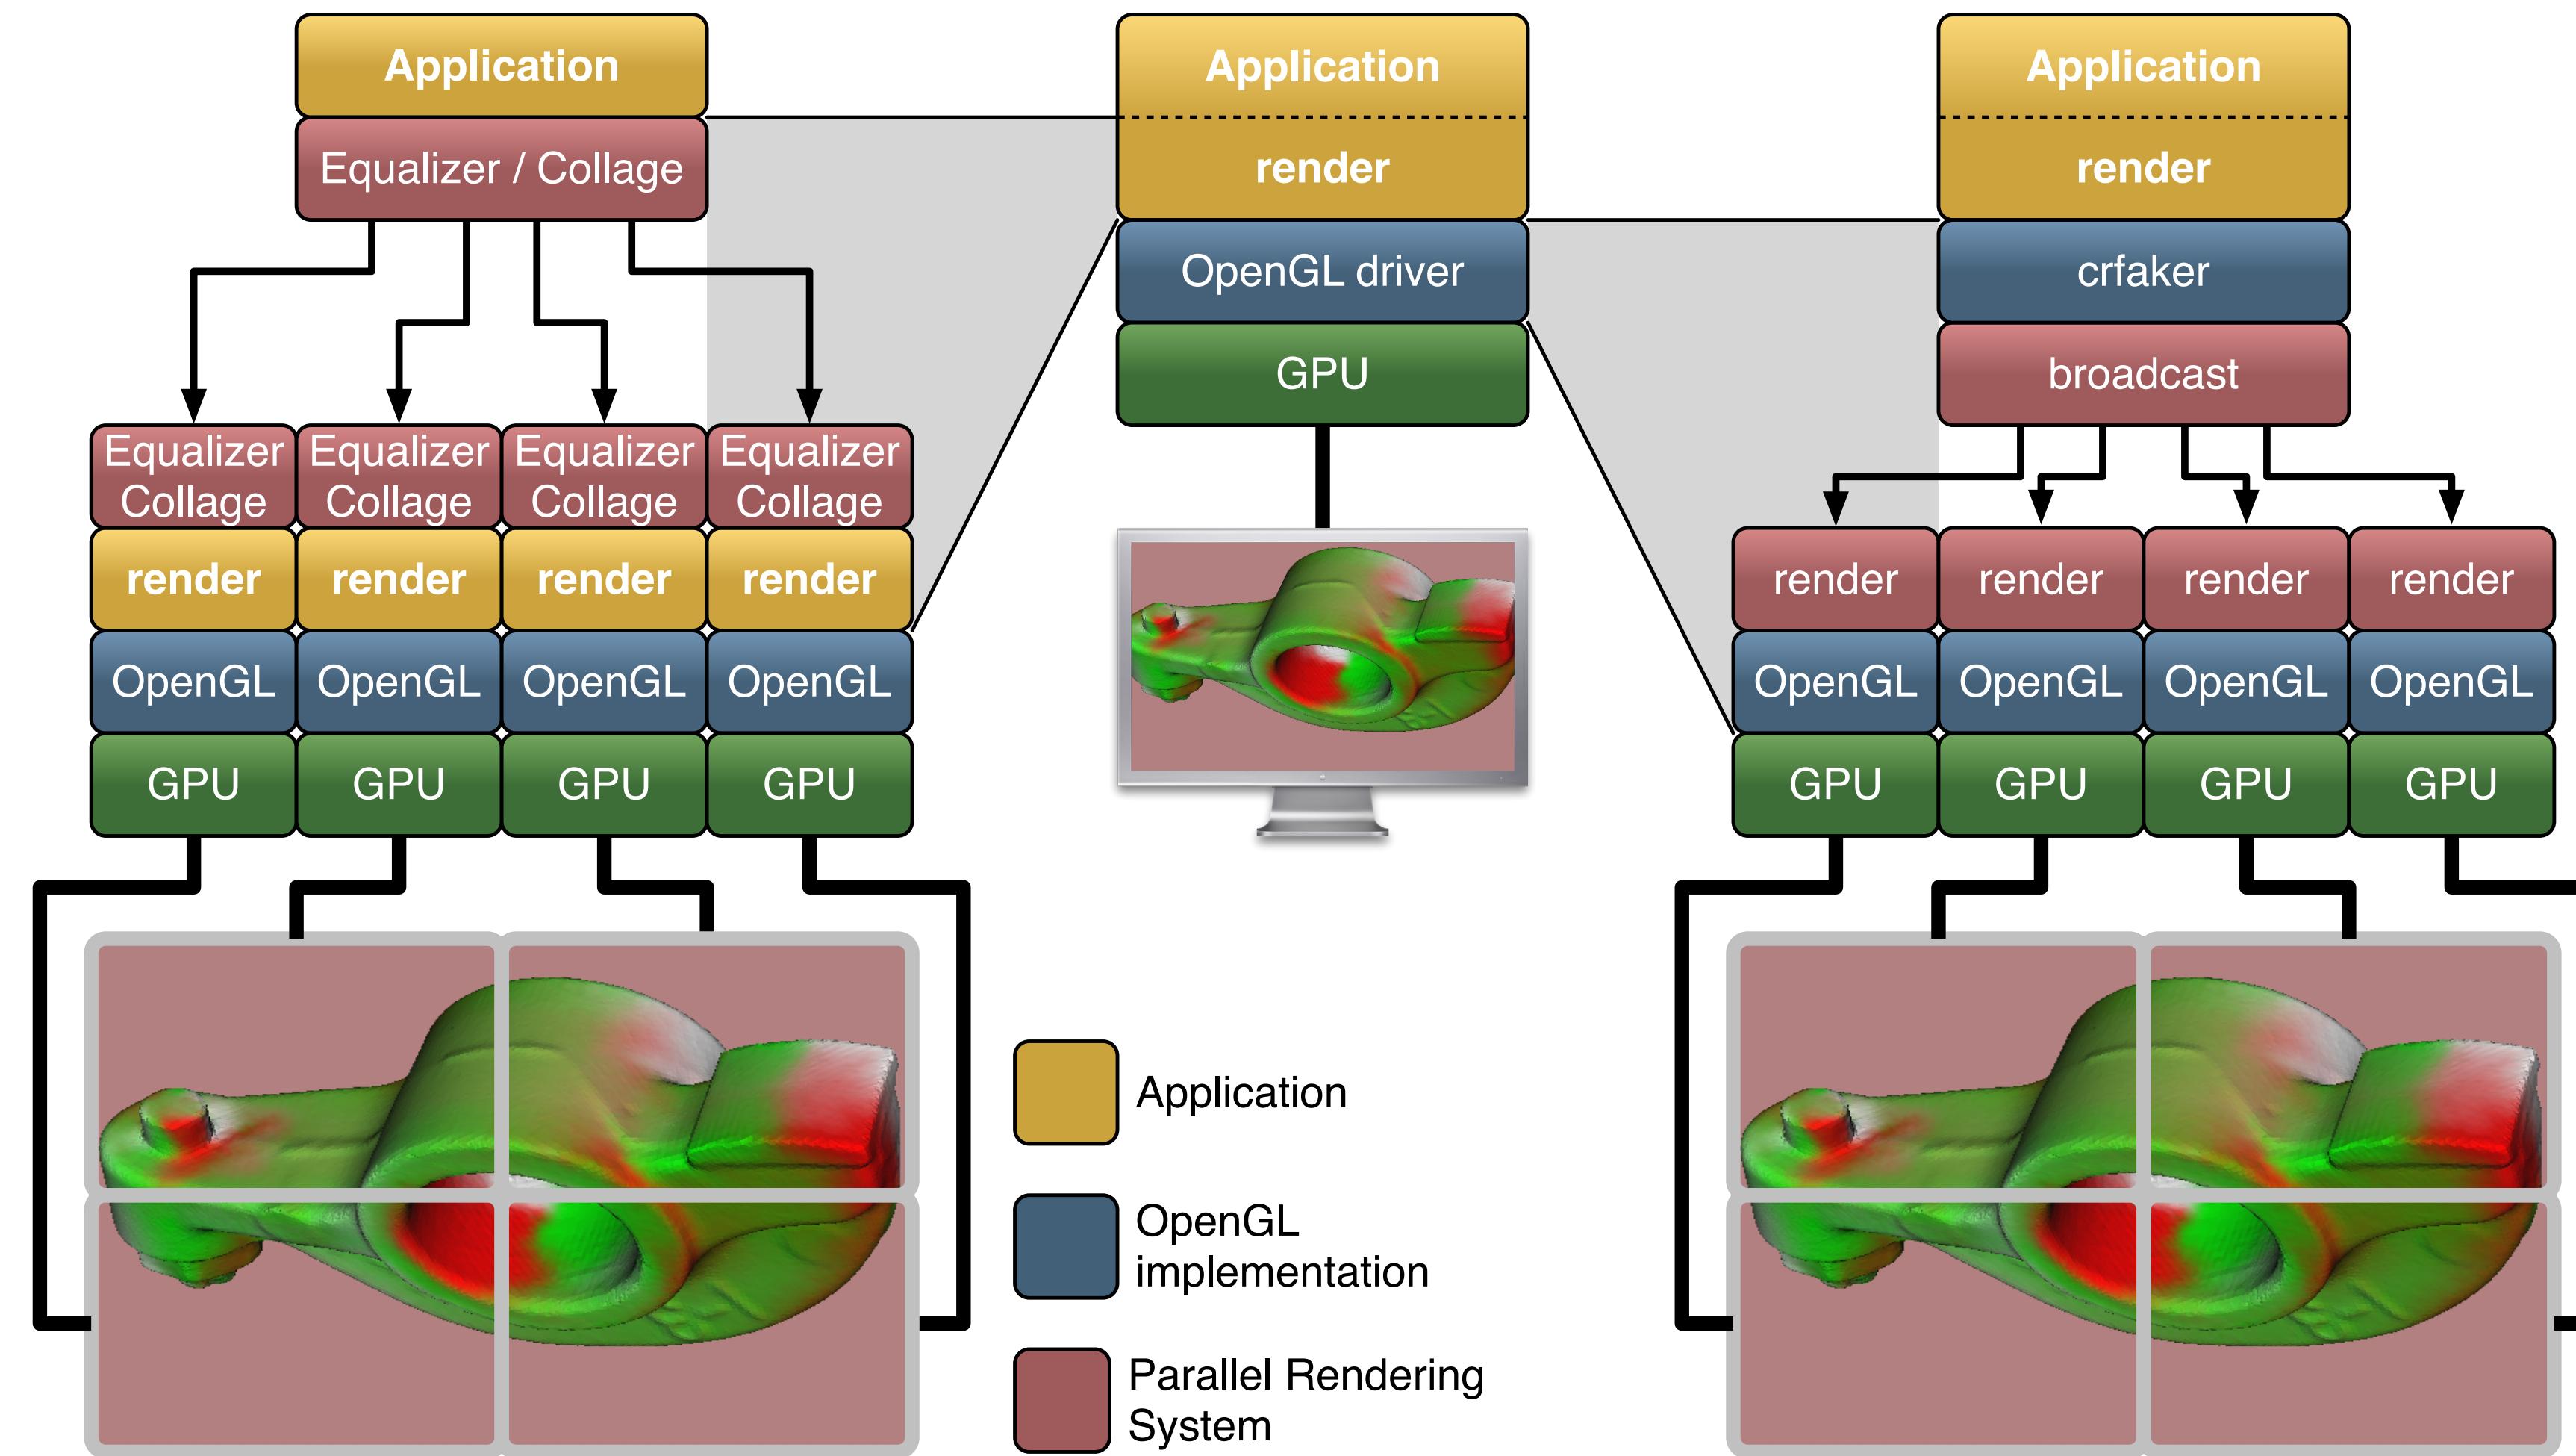
Thank You

Questions?

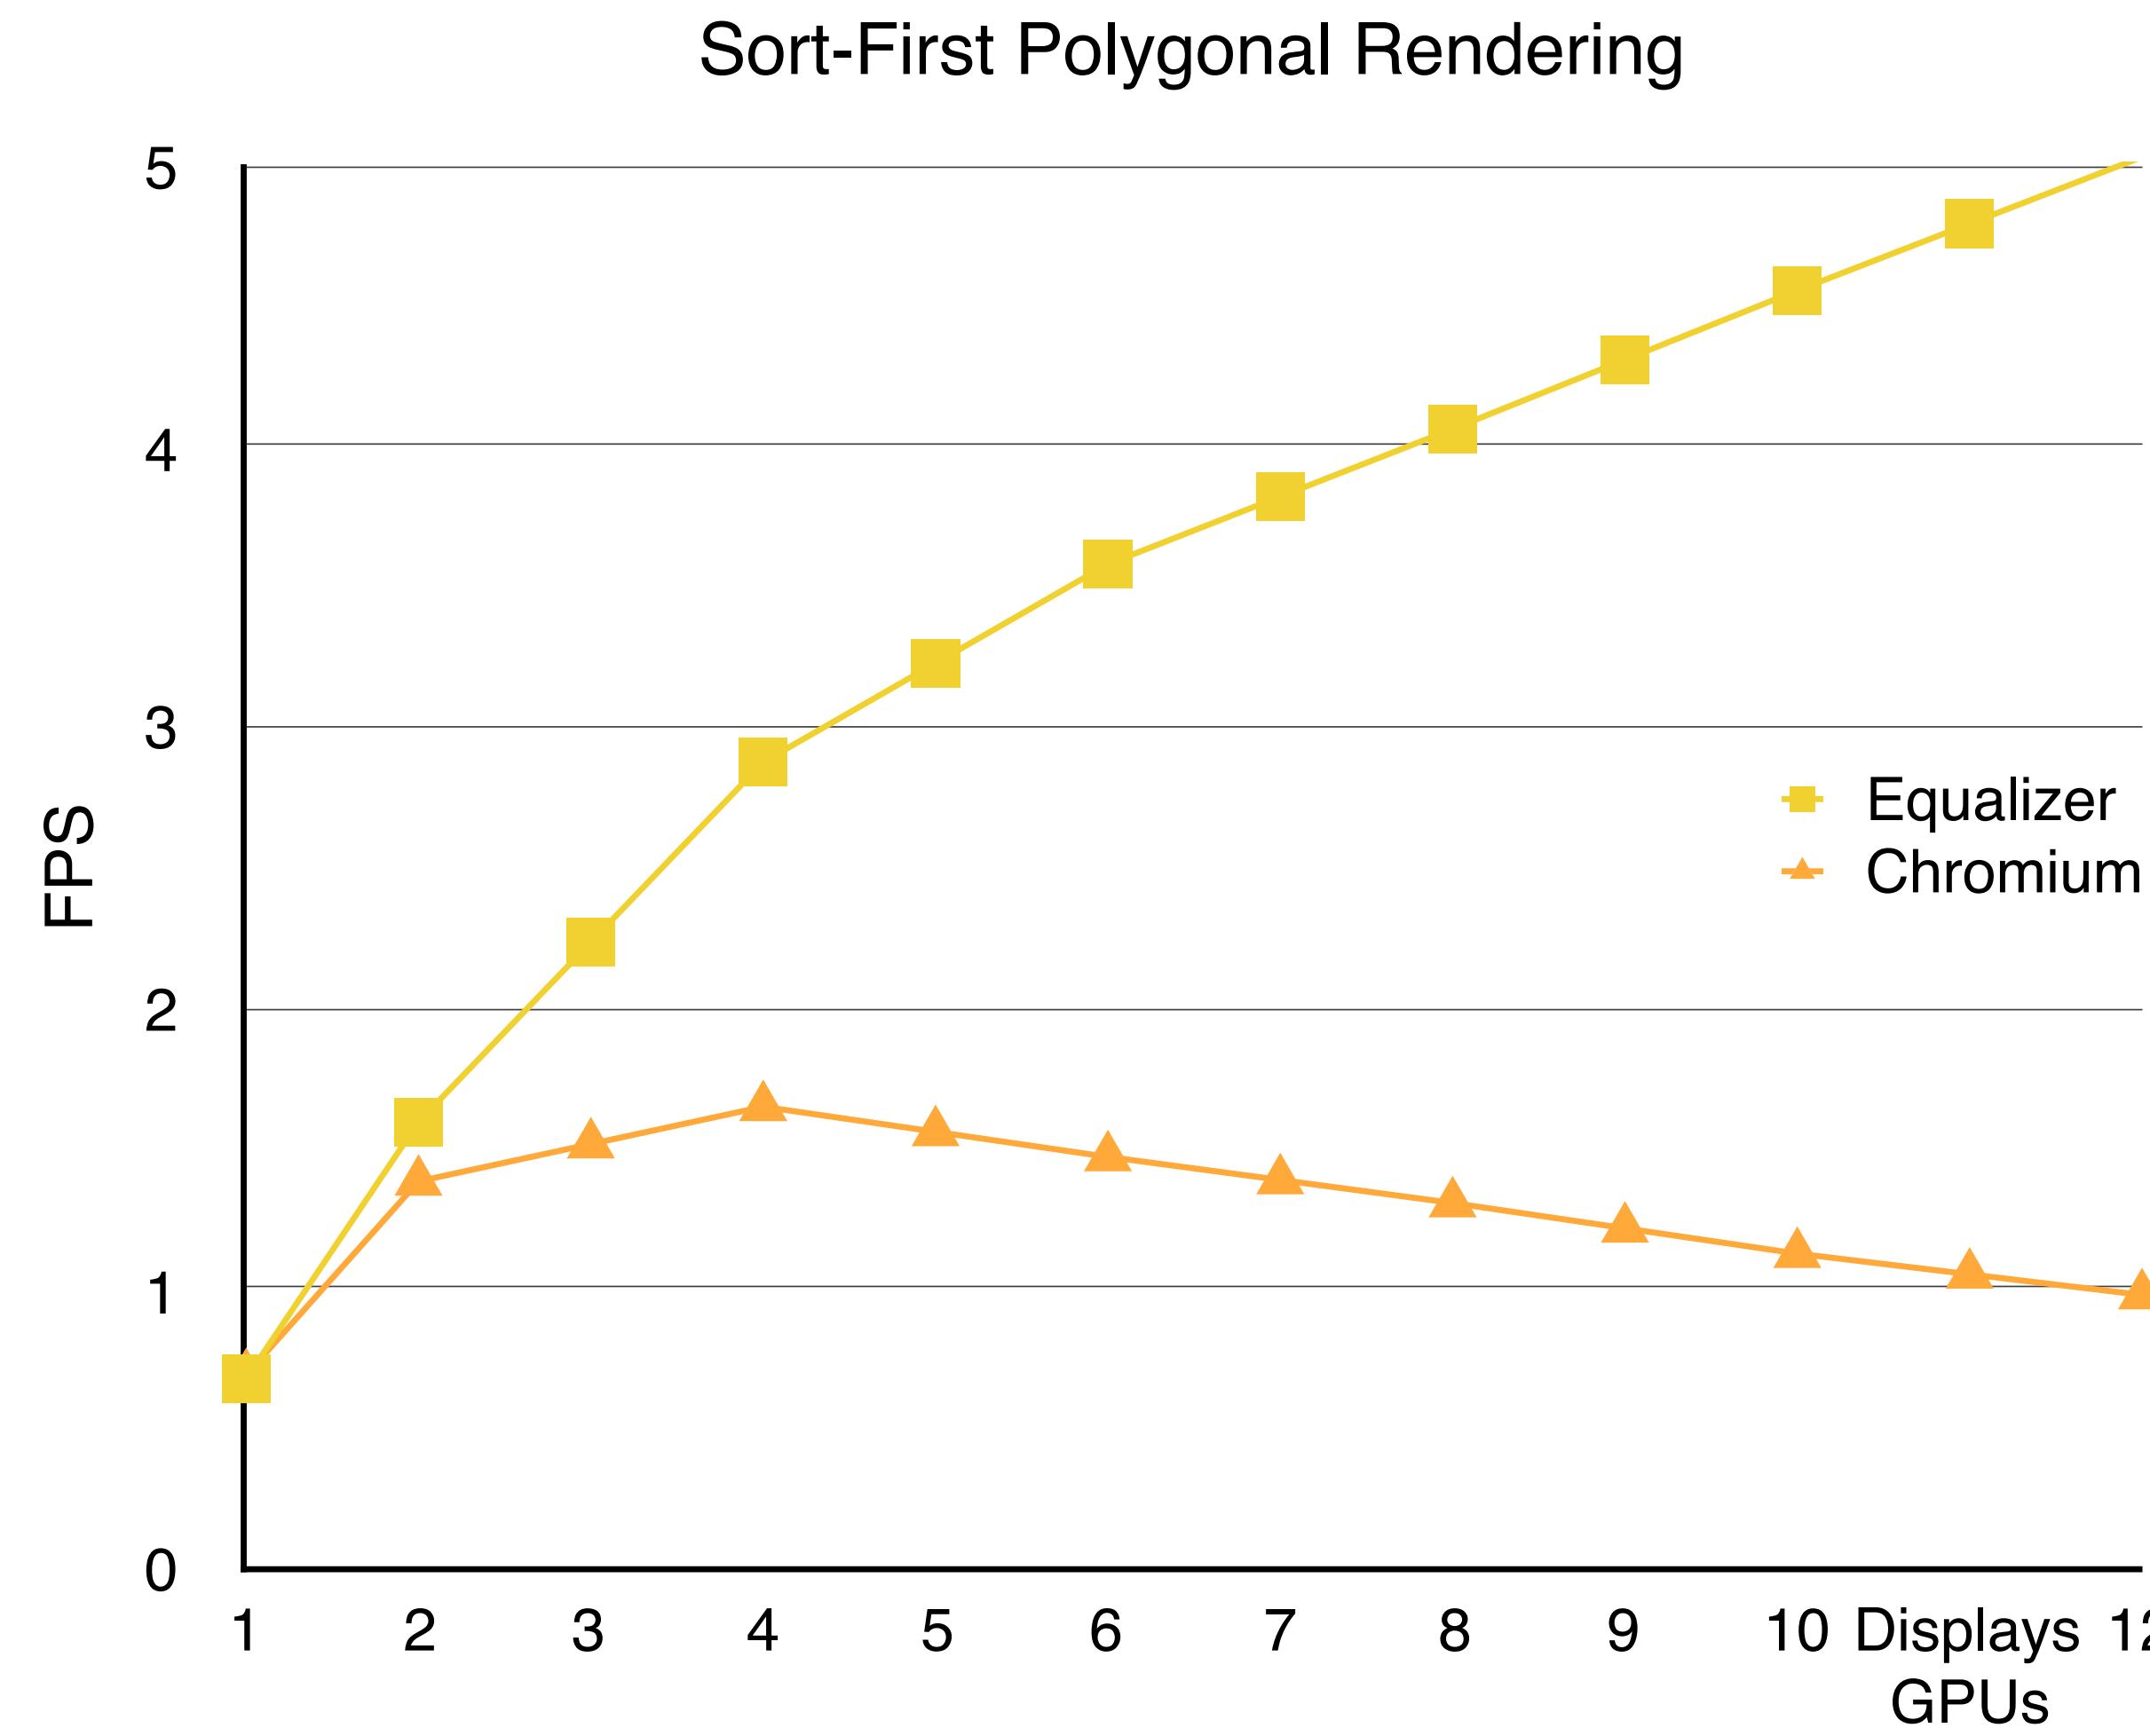
Future Work

- Scene graph and rendering toolkit integration
- Auto-selection of scalable rendering algorithms and resources
- Load balancing sort-last compositing
- Predictive load balancing
- Sort-last scalable raytracing

Chromium vs Equalizer



Chromium vs Equalizer

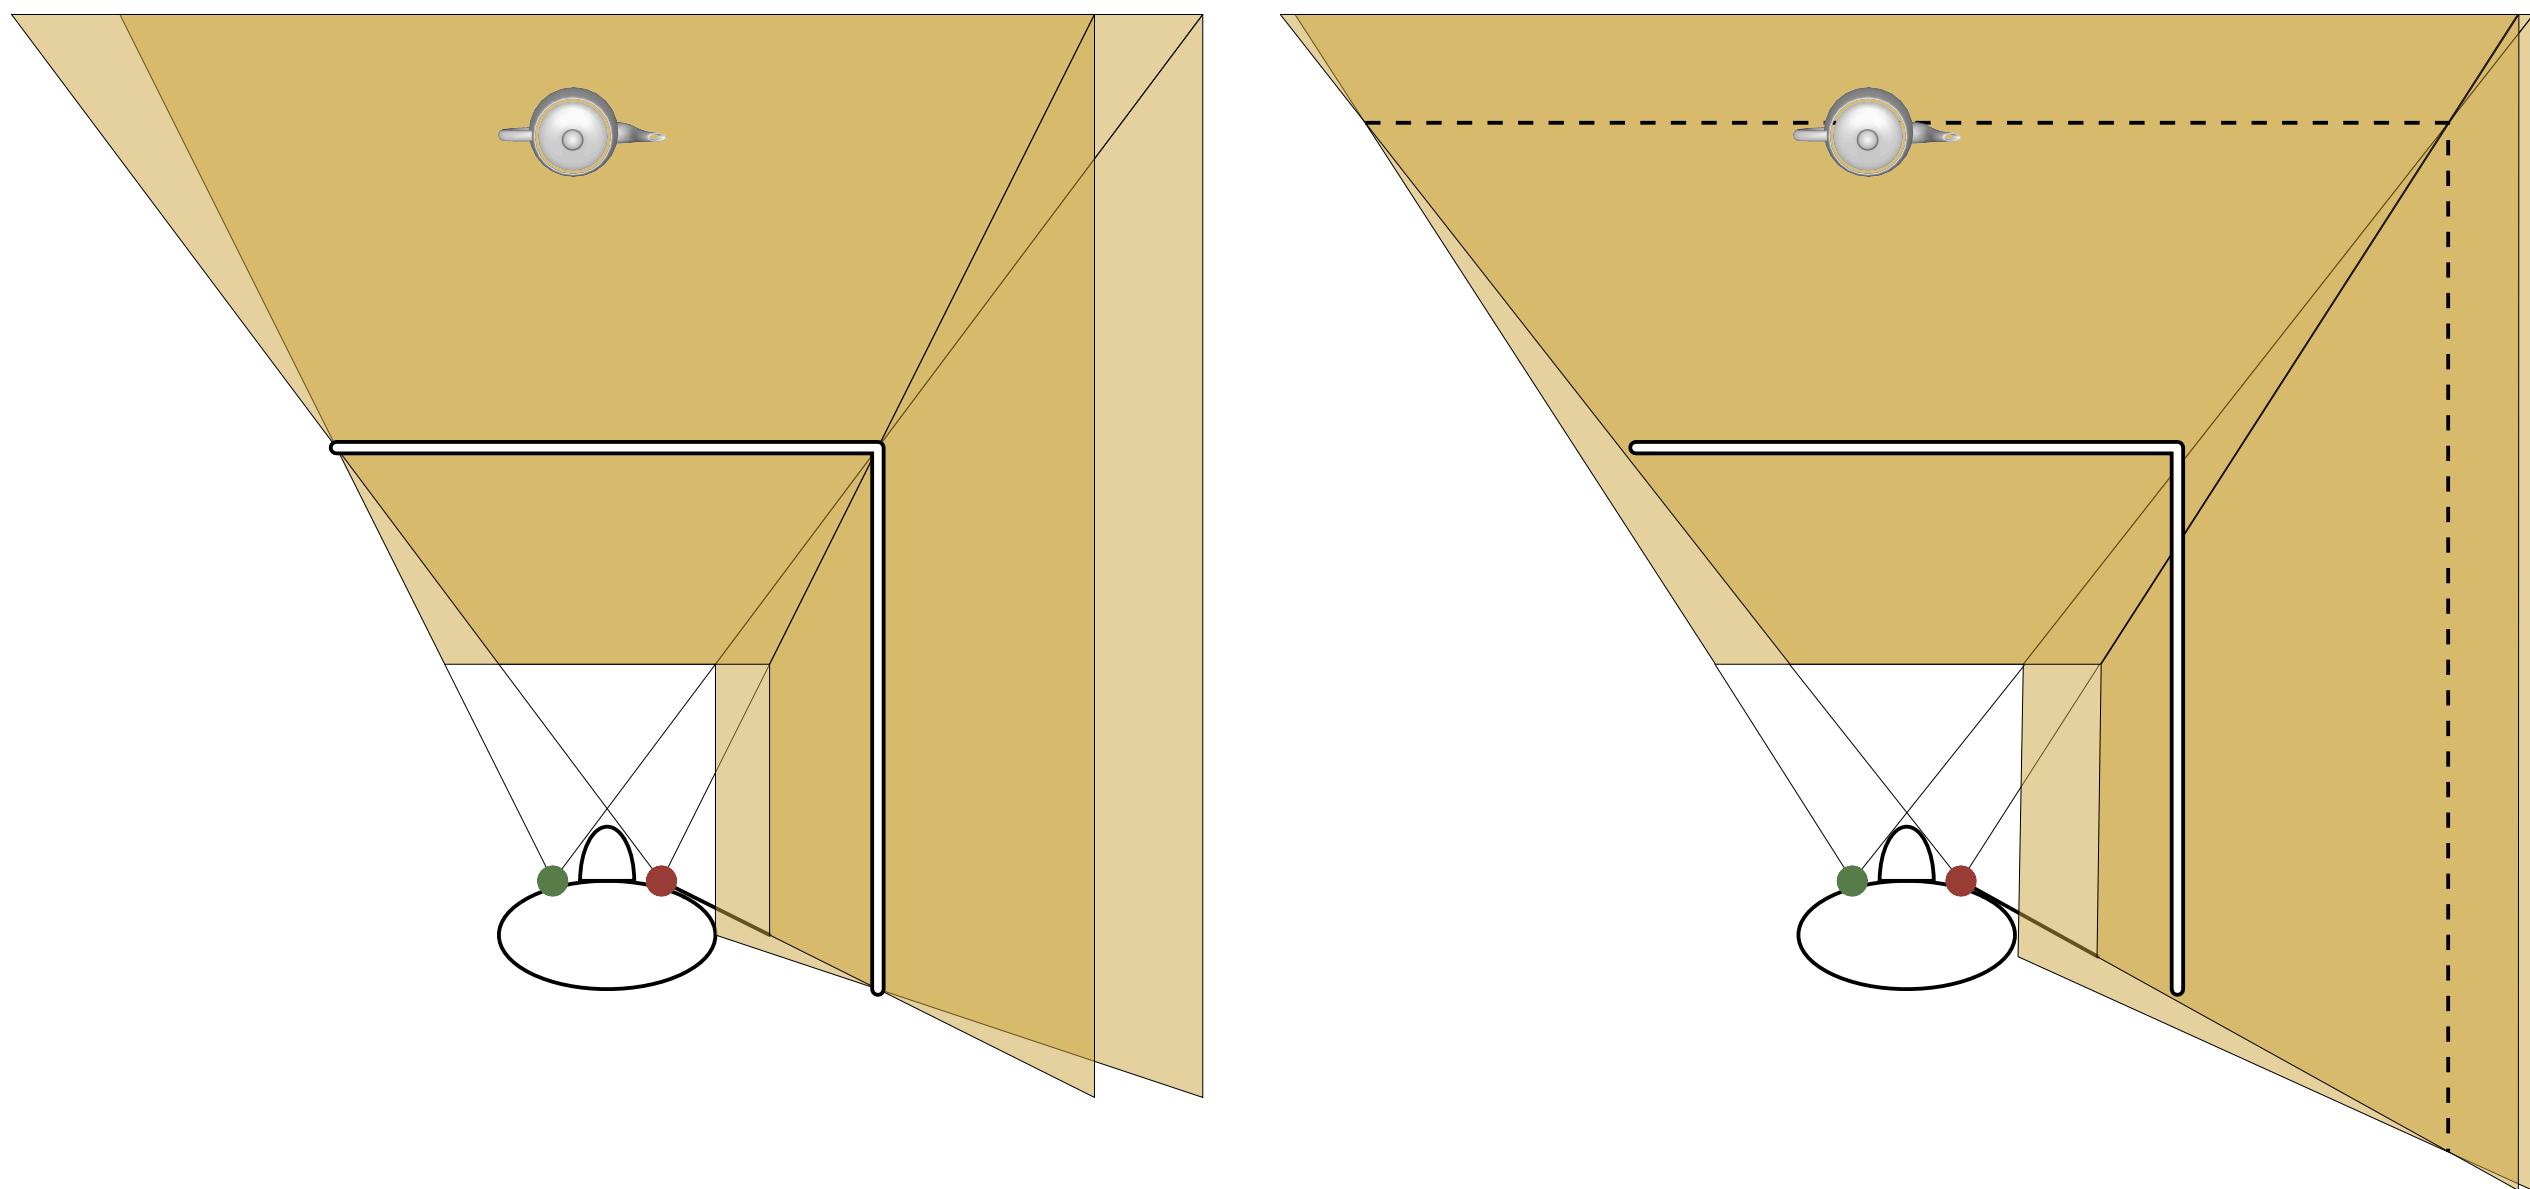


Display Abstraction



Virtual Reality

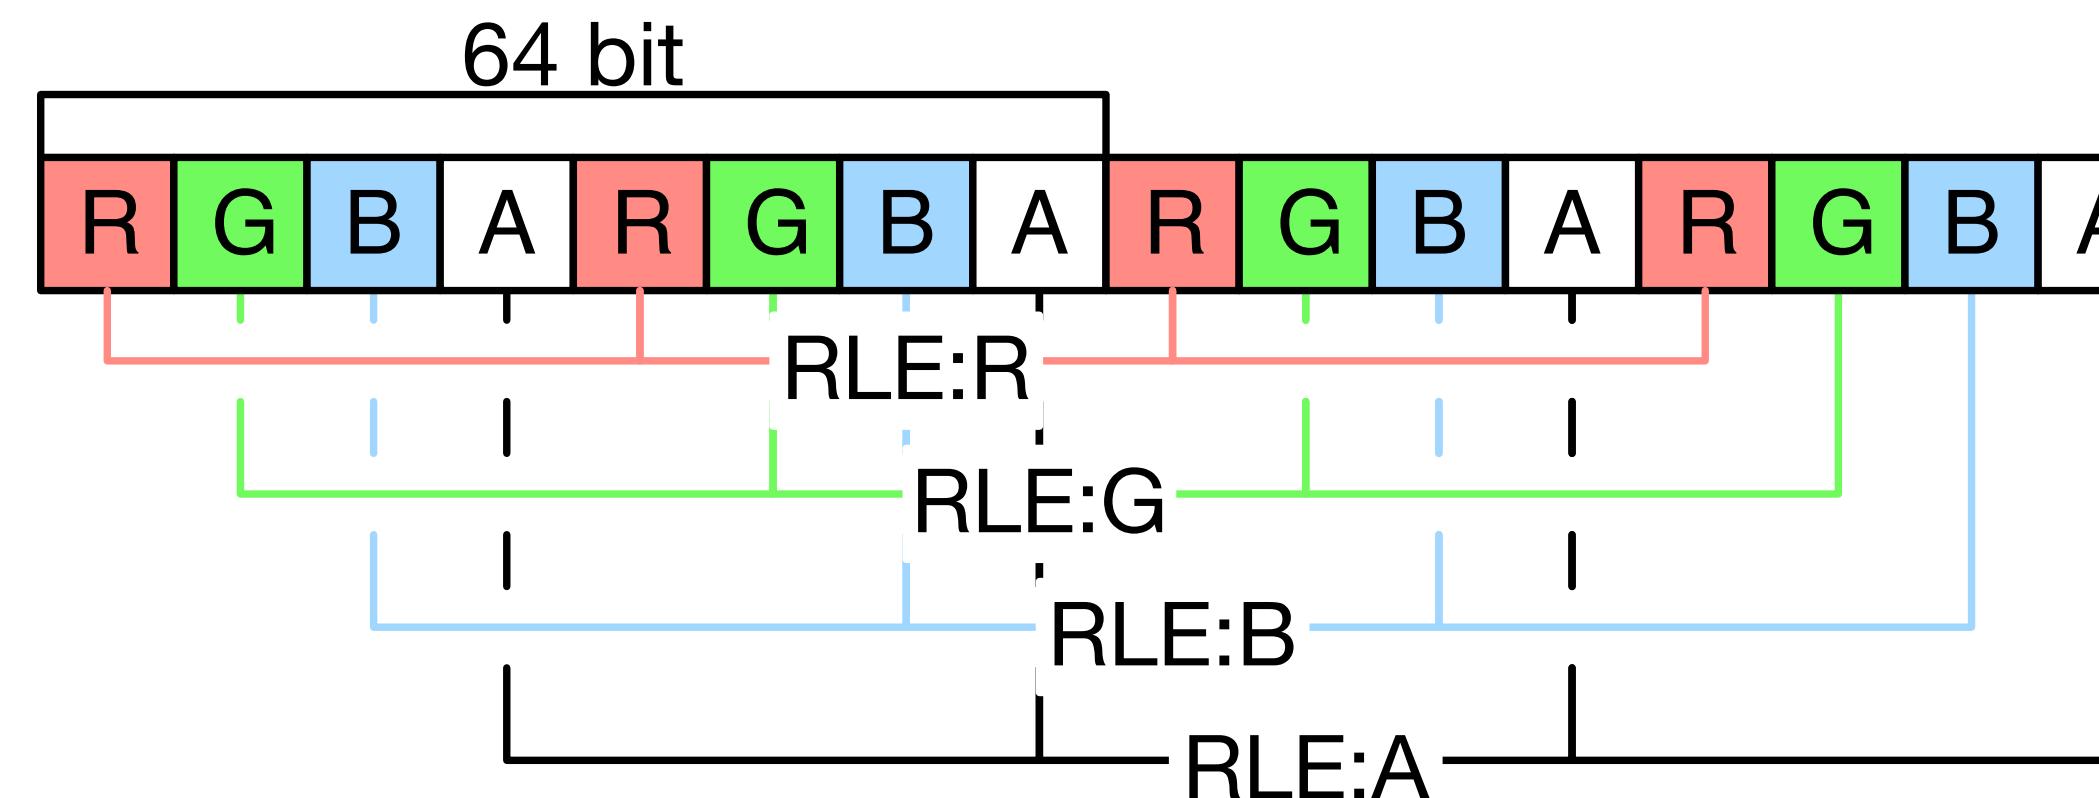
- Head Tracking
 - ▶ Cave
 - ▶ HMD
- Dynamic Focus Distance
- Asymmetric Eye Position
- Model Unit
- Runtime Stereo Switch



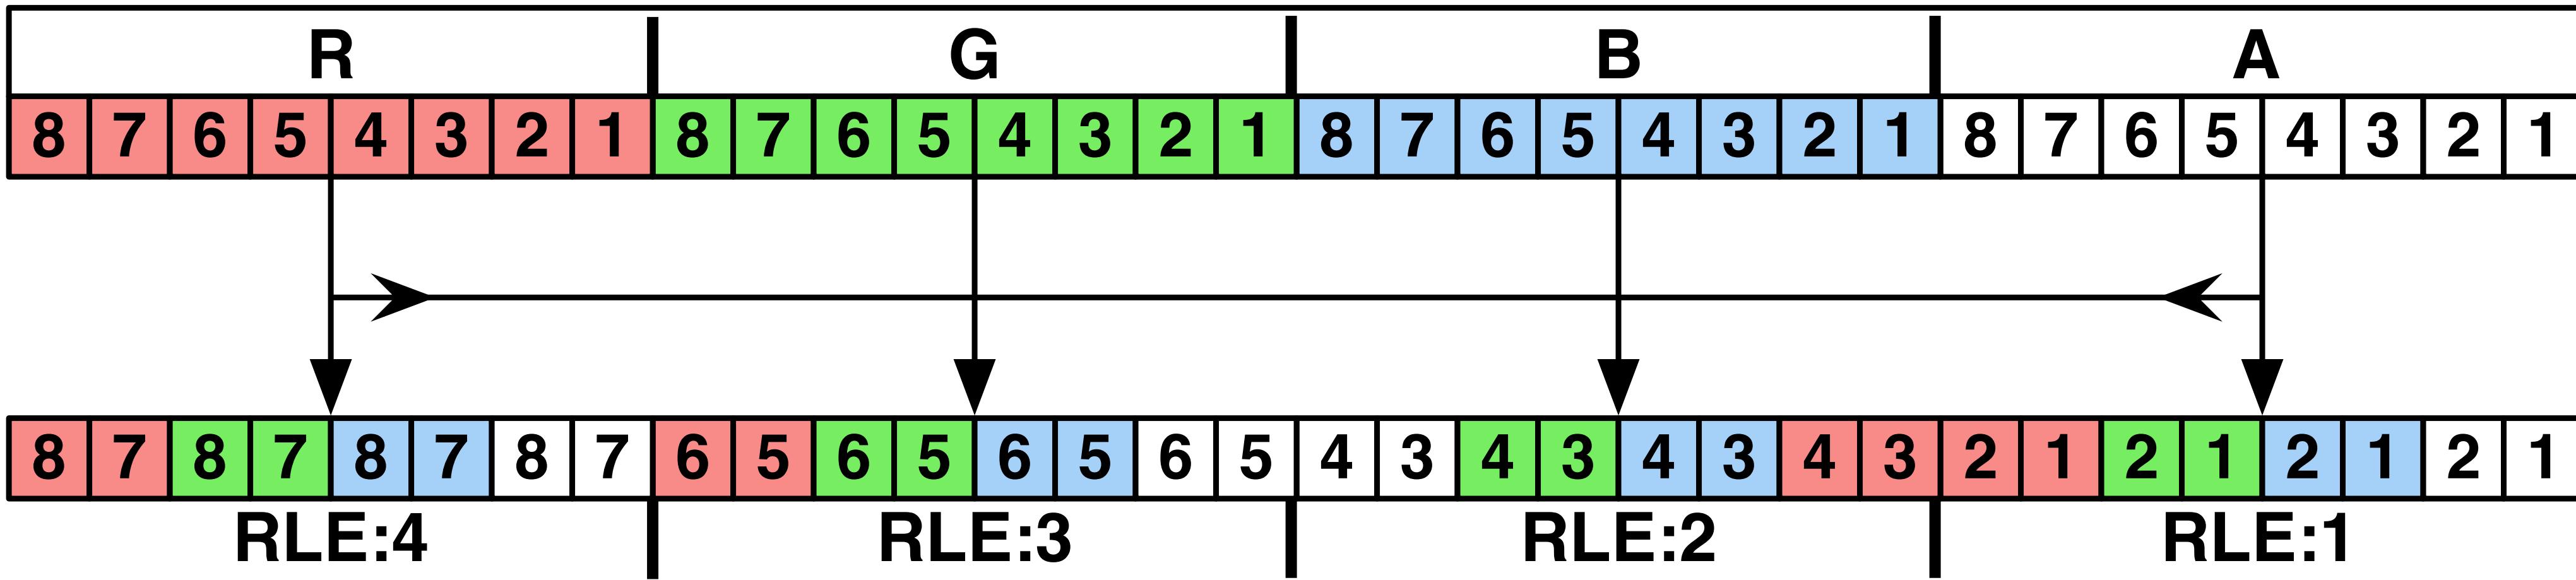
Collage

- Peer-to-peer networking library
 - ▶ Connection - Data/OStream - Node/LocalNode - Objects
- Versioned, distributed objects
 - ▶ Compression
 - ▶ Chunking, Buffering, Caching
 - ▶ Multicast
- Reliable Stream Protocol
 - ▶ TCP semantics over UDP Multicast

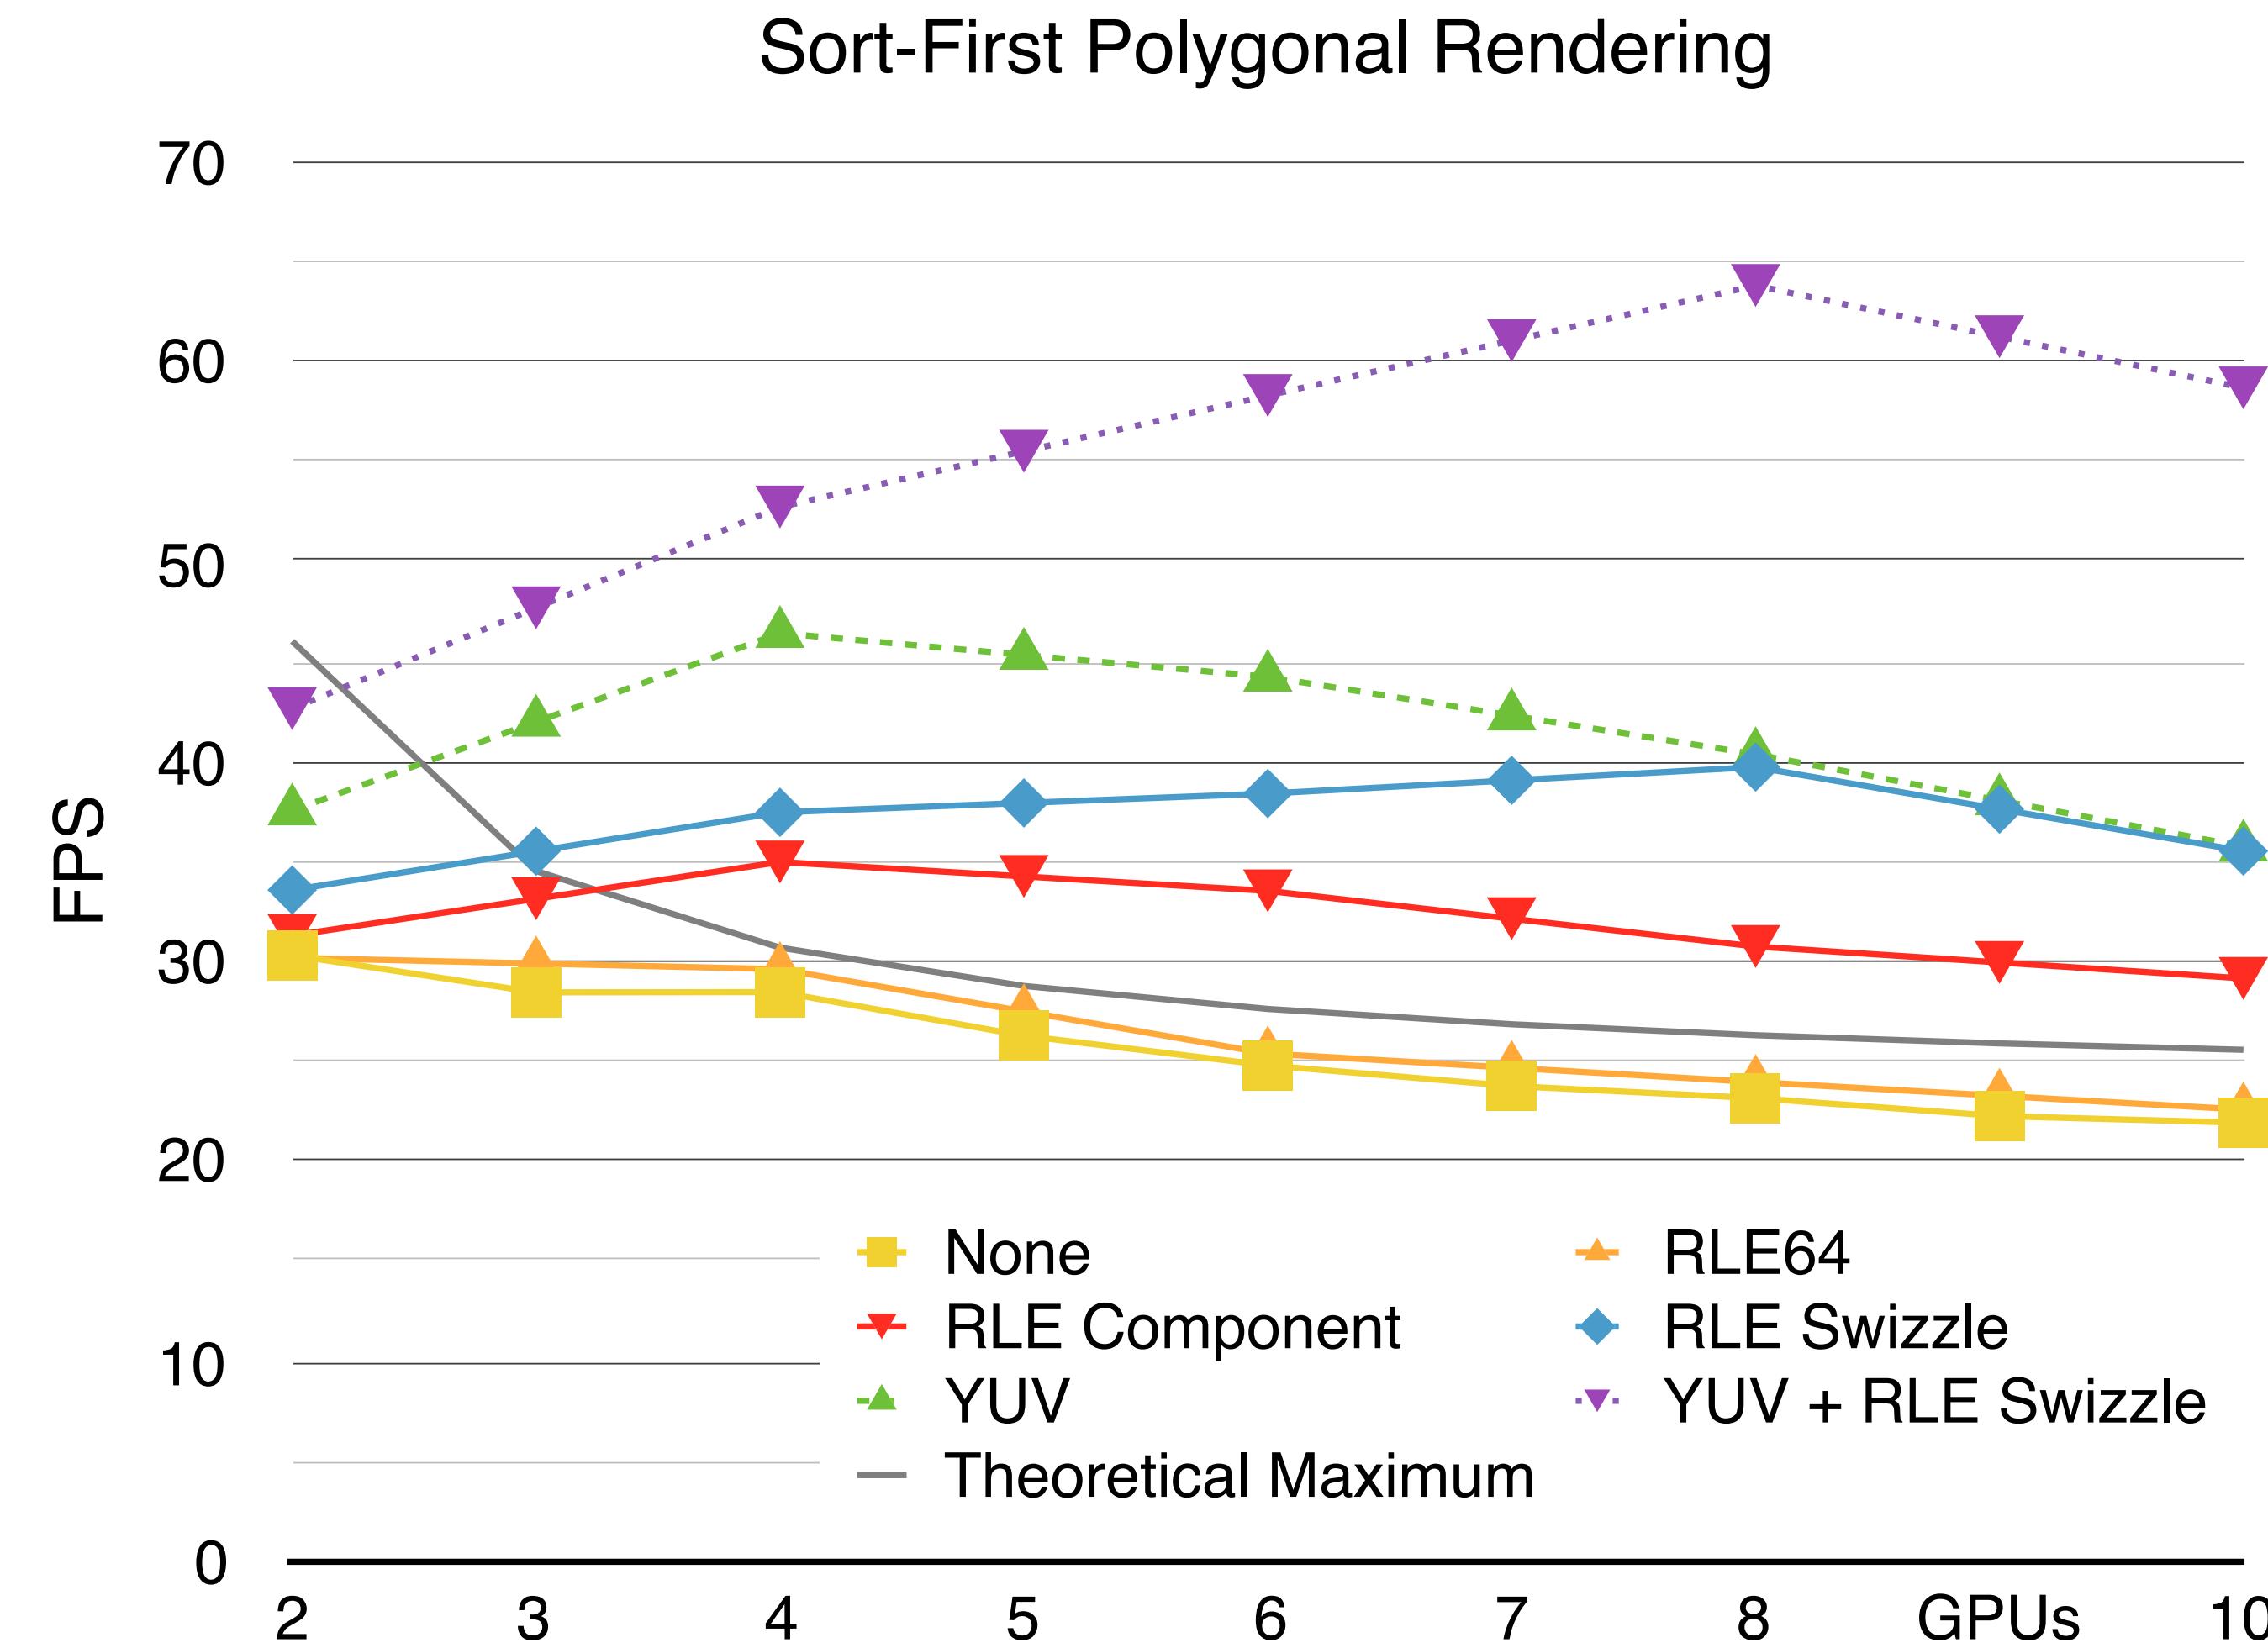
RLE Compression



32 bit

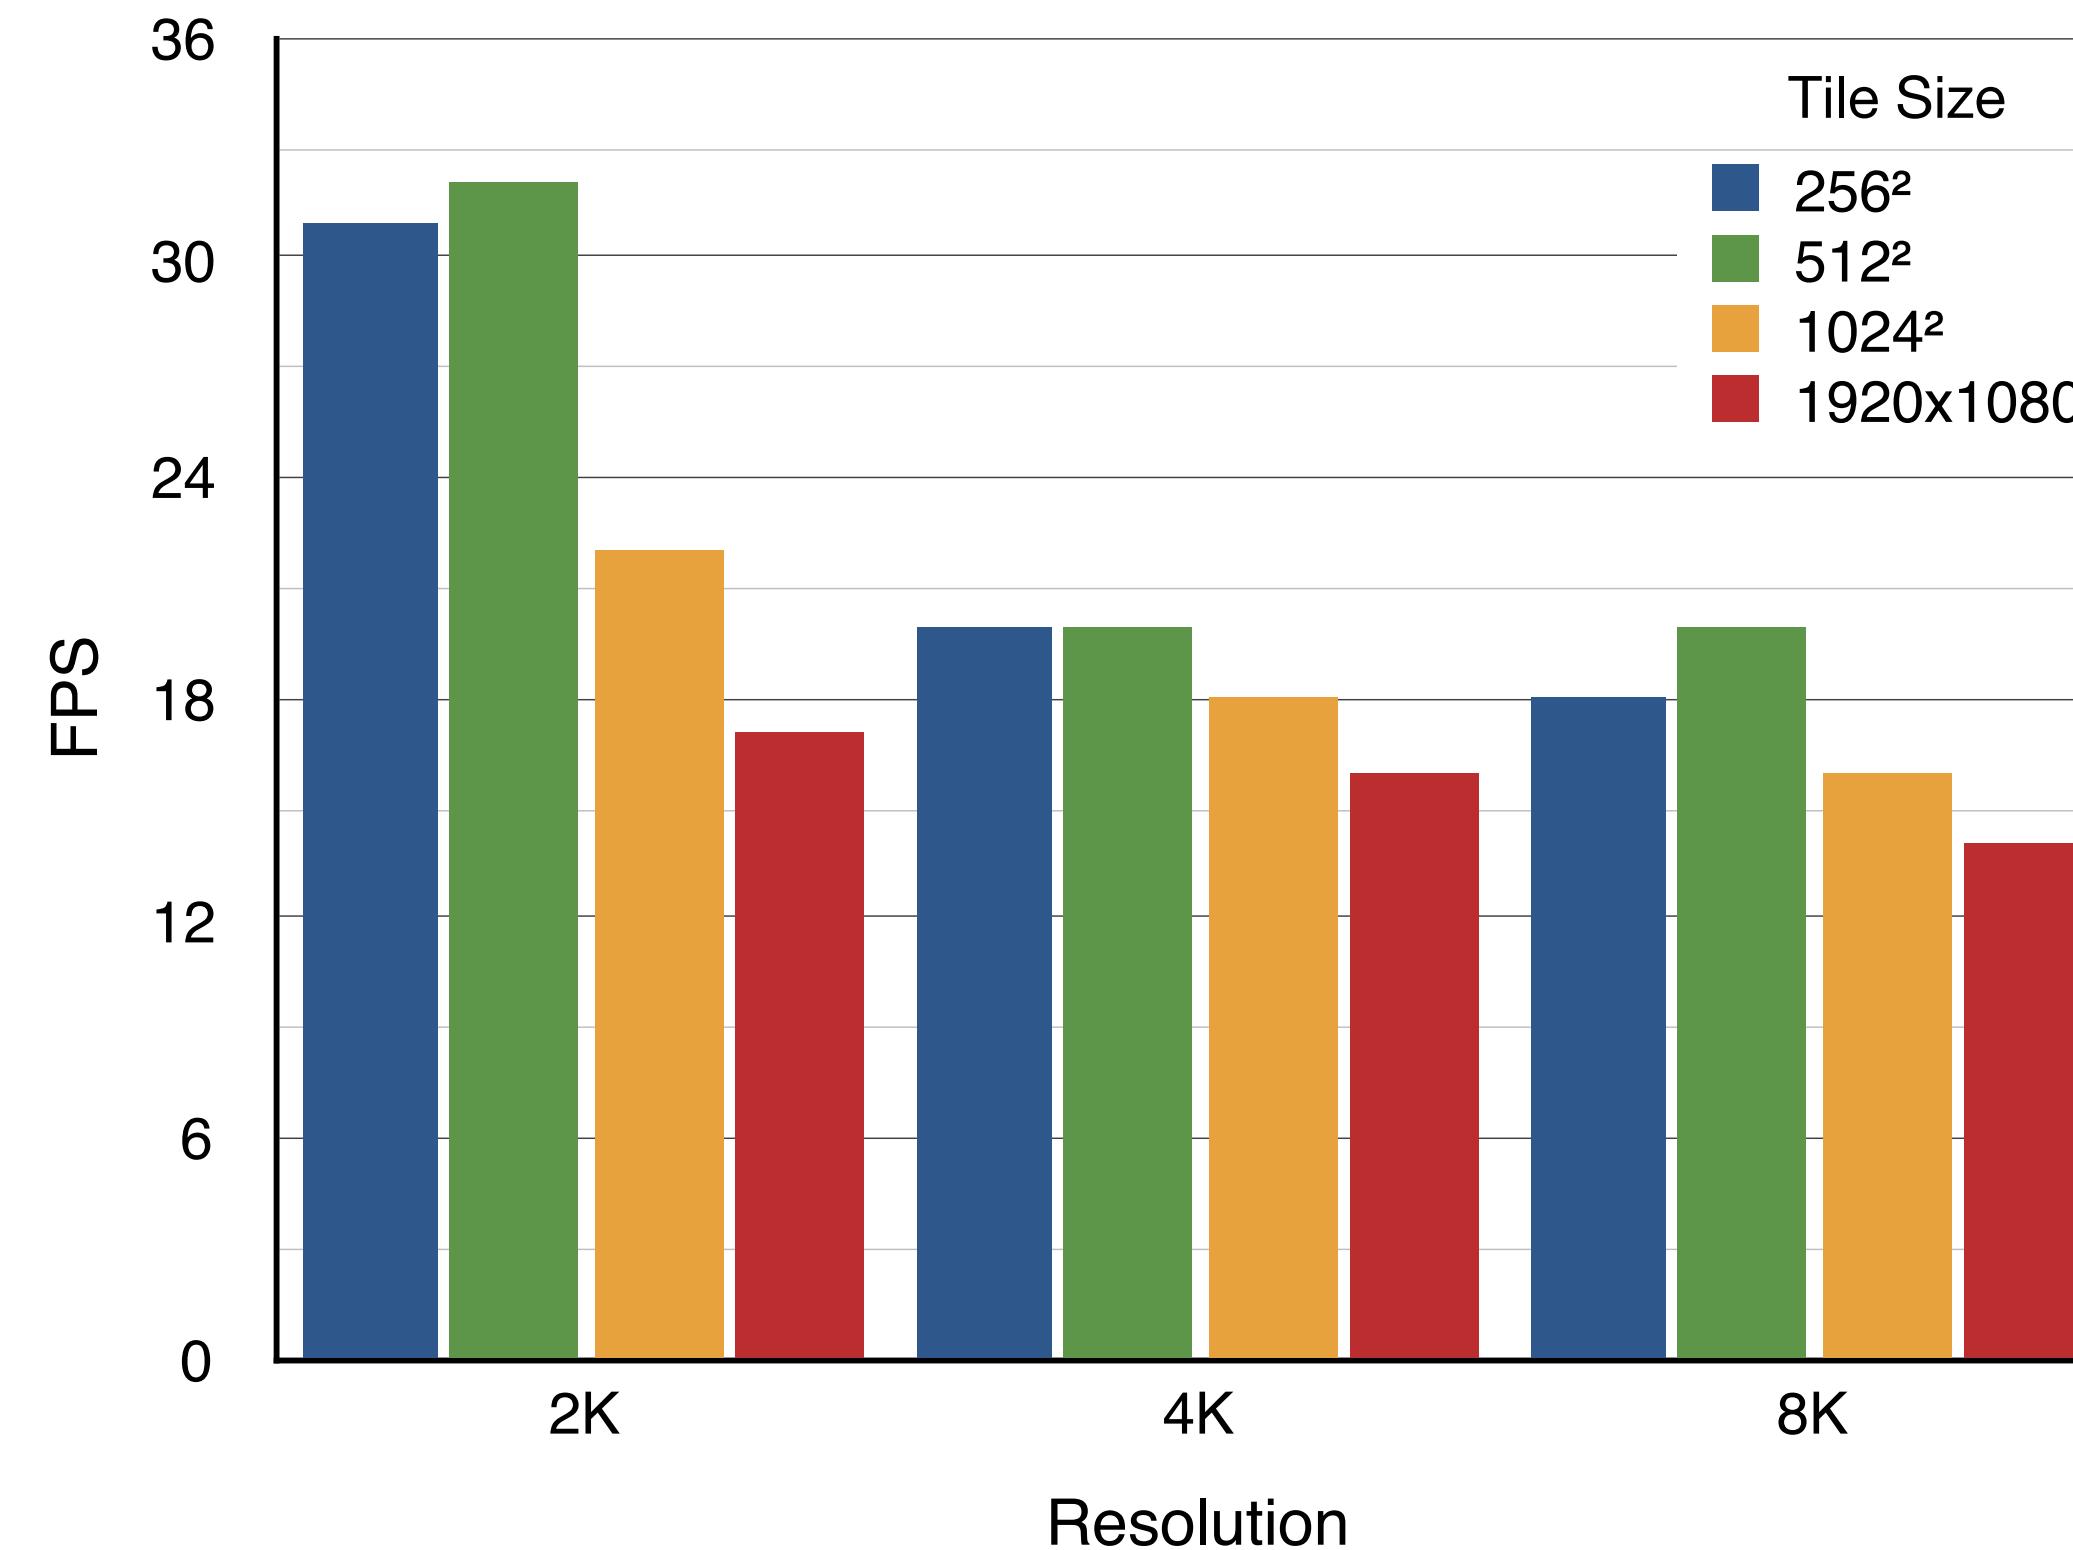


RLE Compression

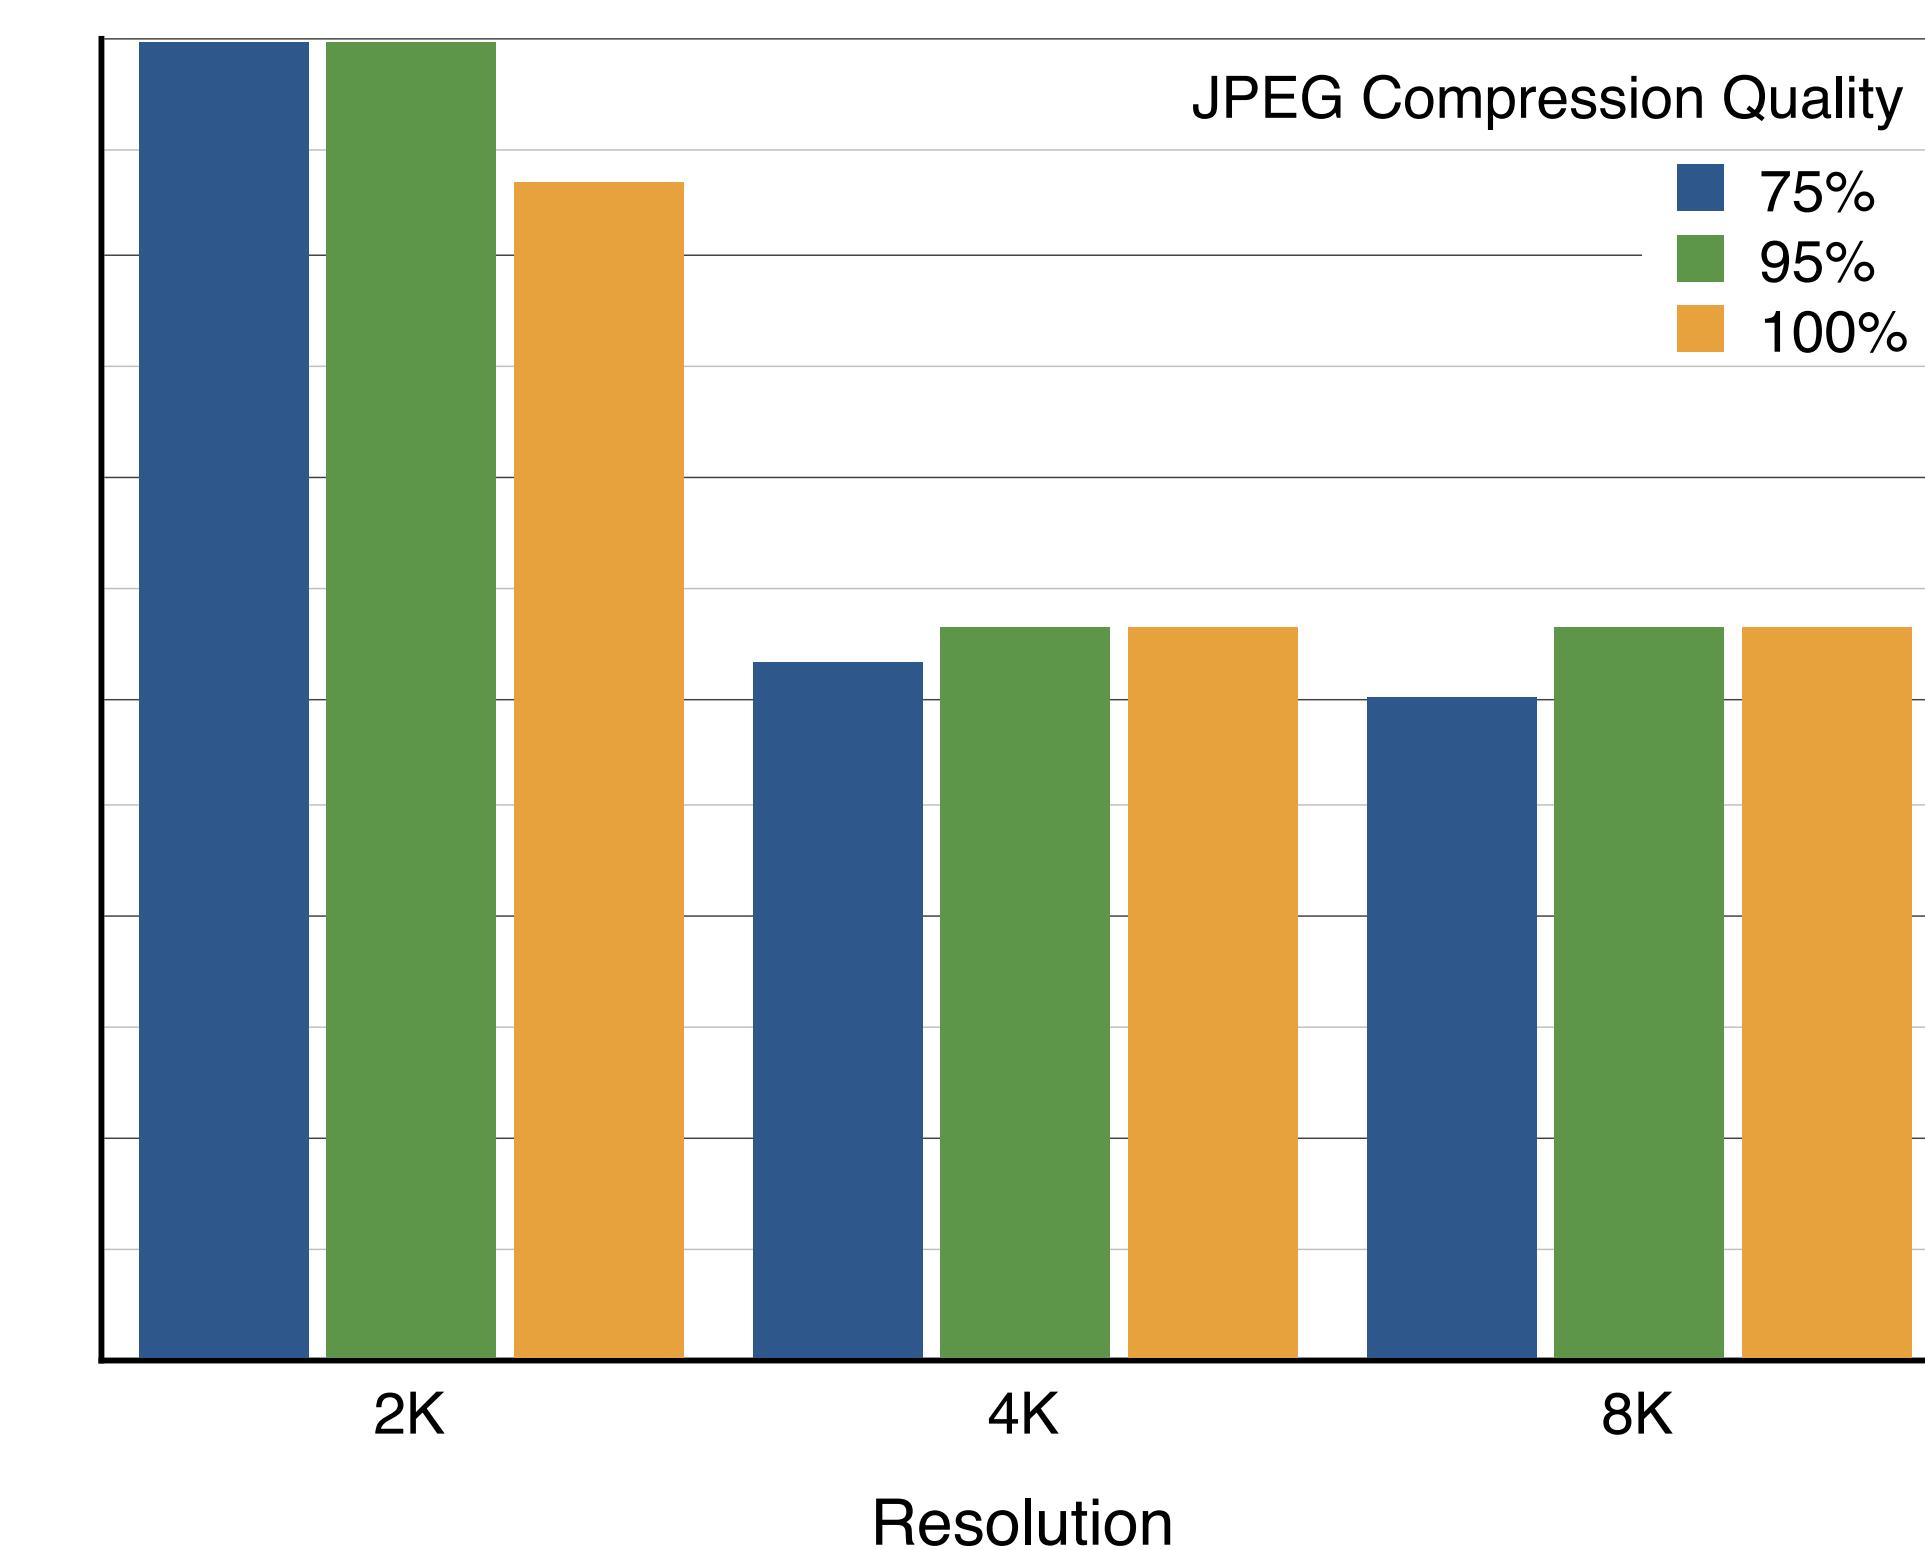


Tide Streaming

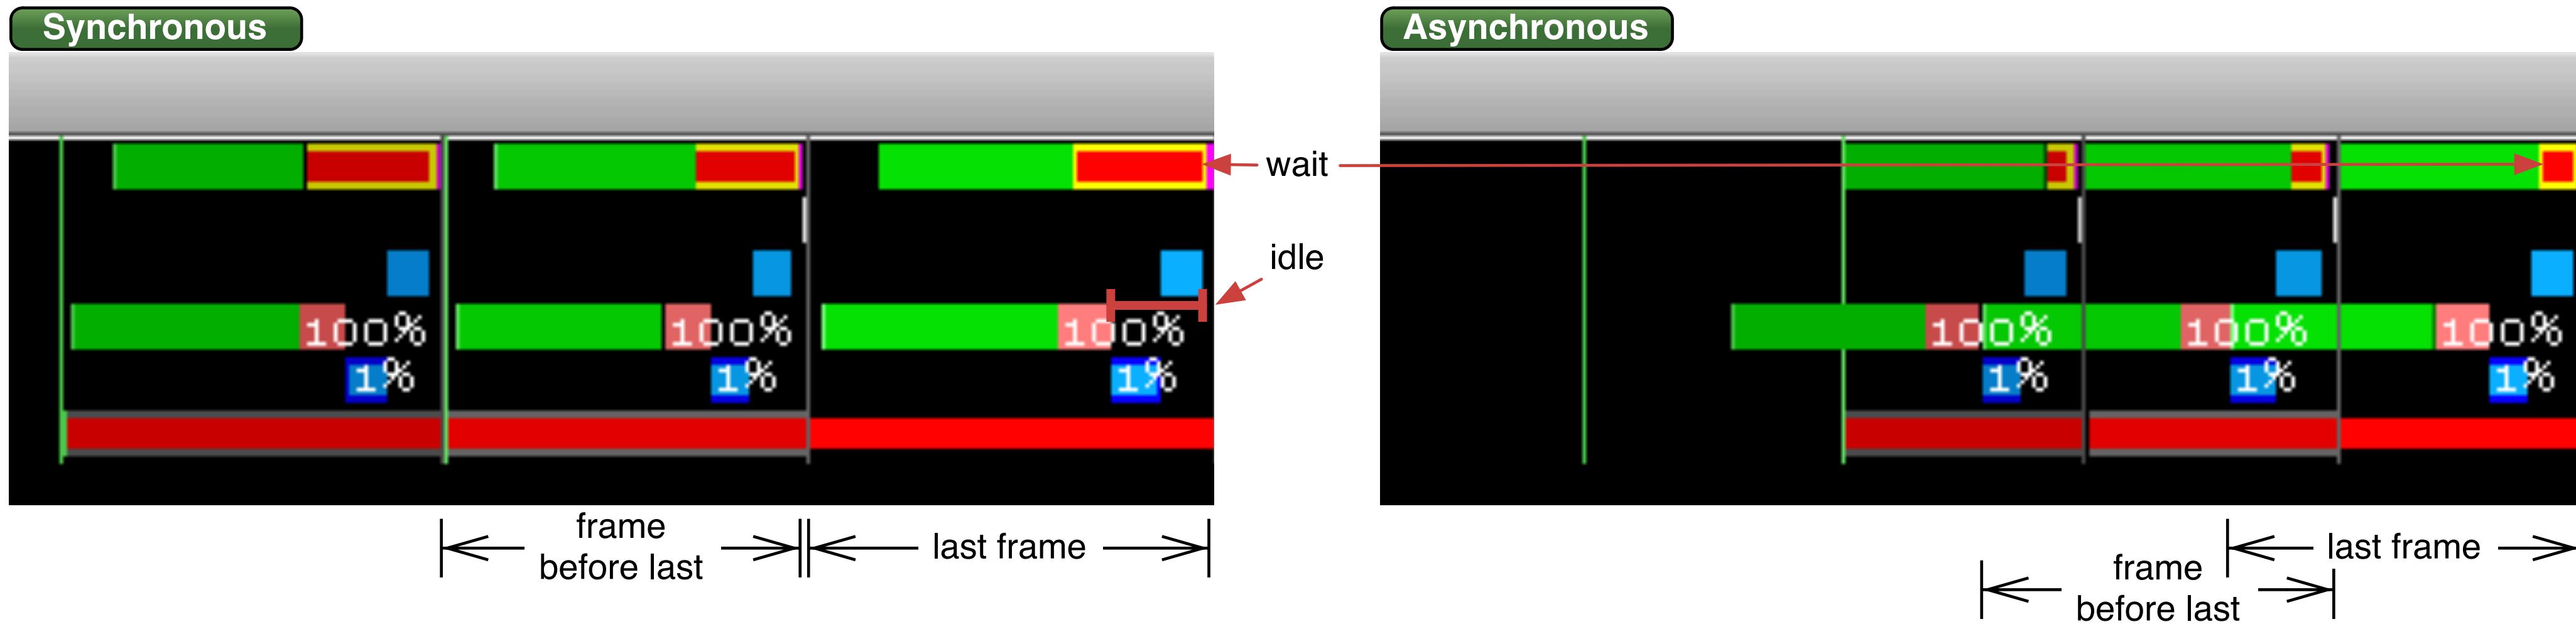
RTNeuron - Tide Remote Parallel Rendering



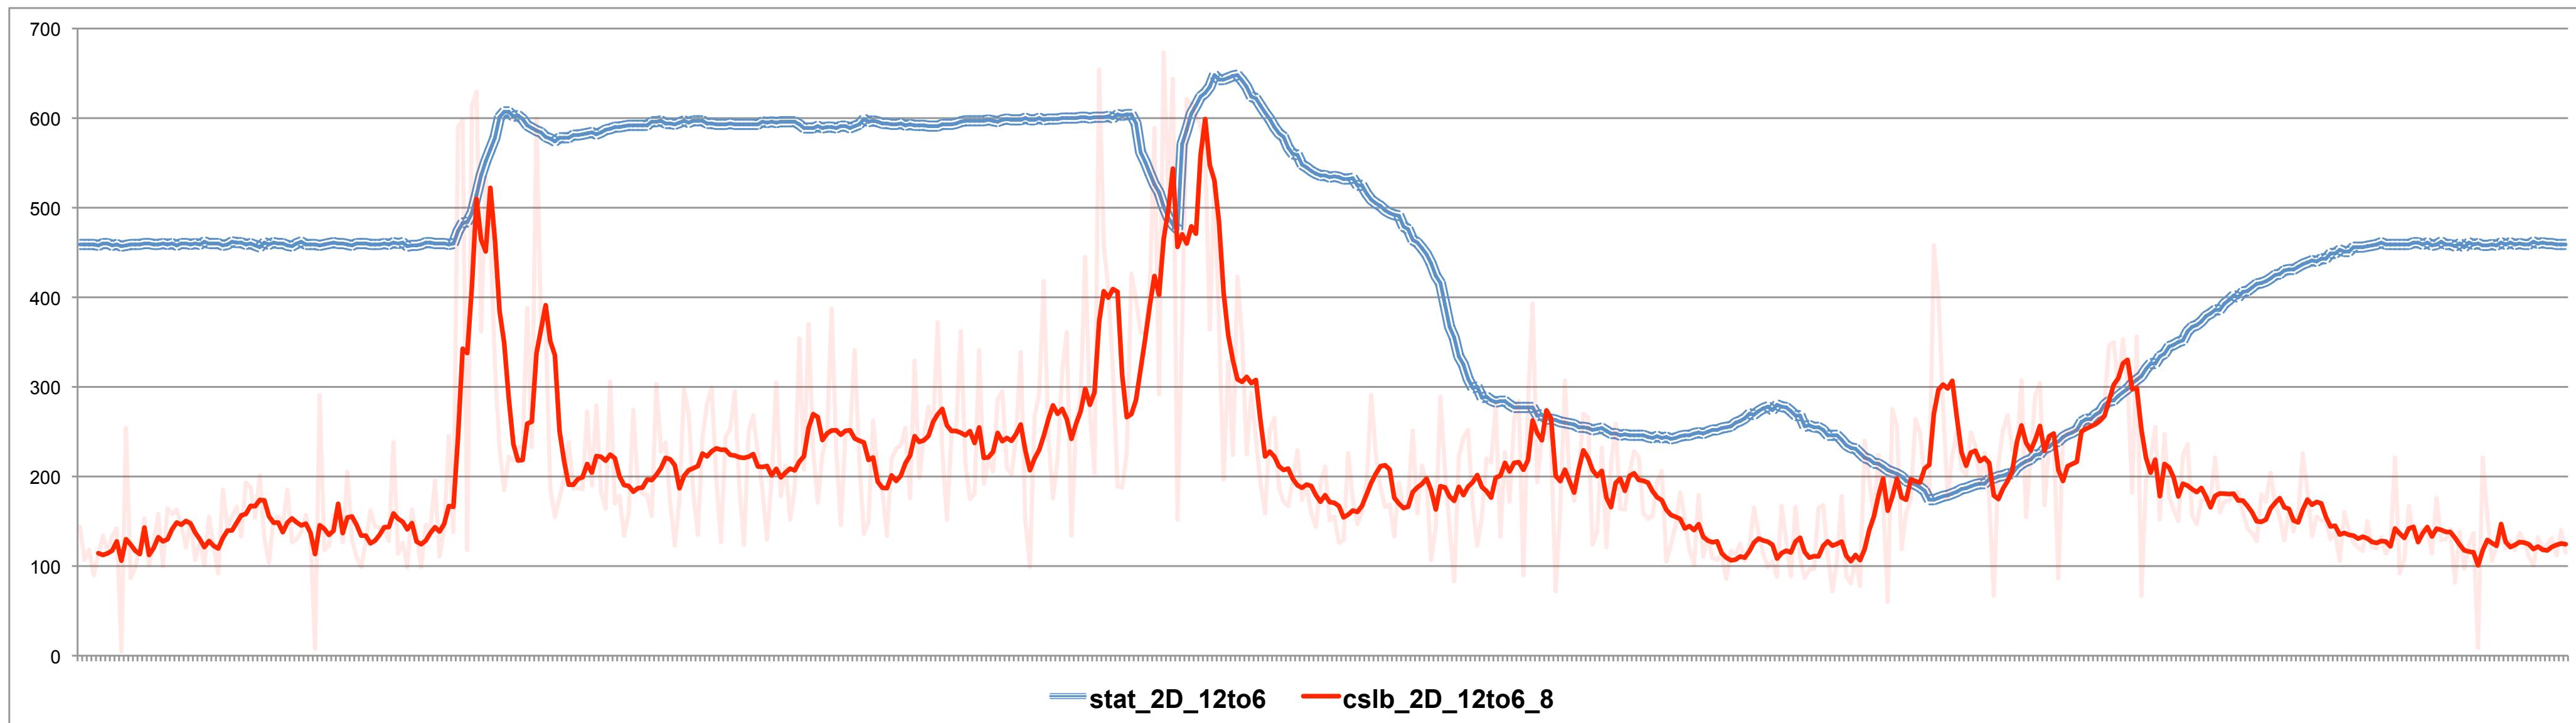
512 2 Tile Size



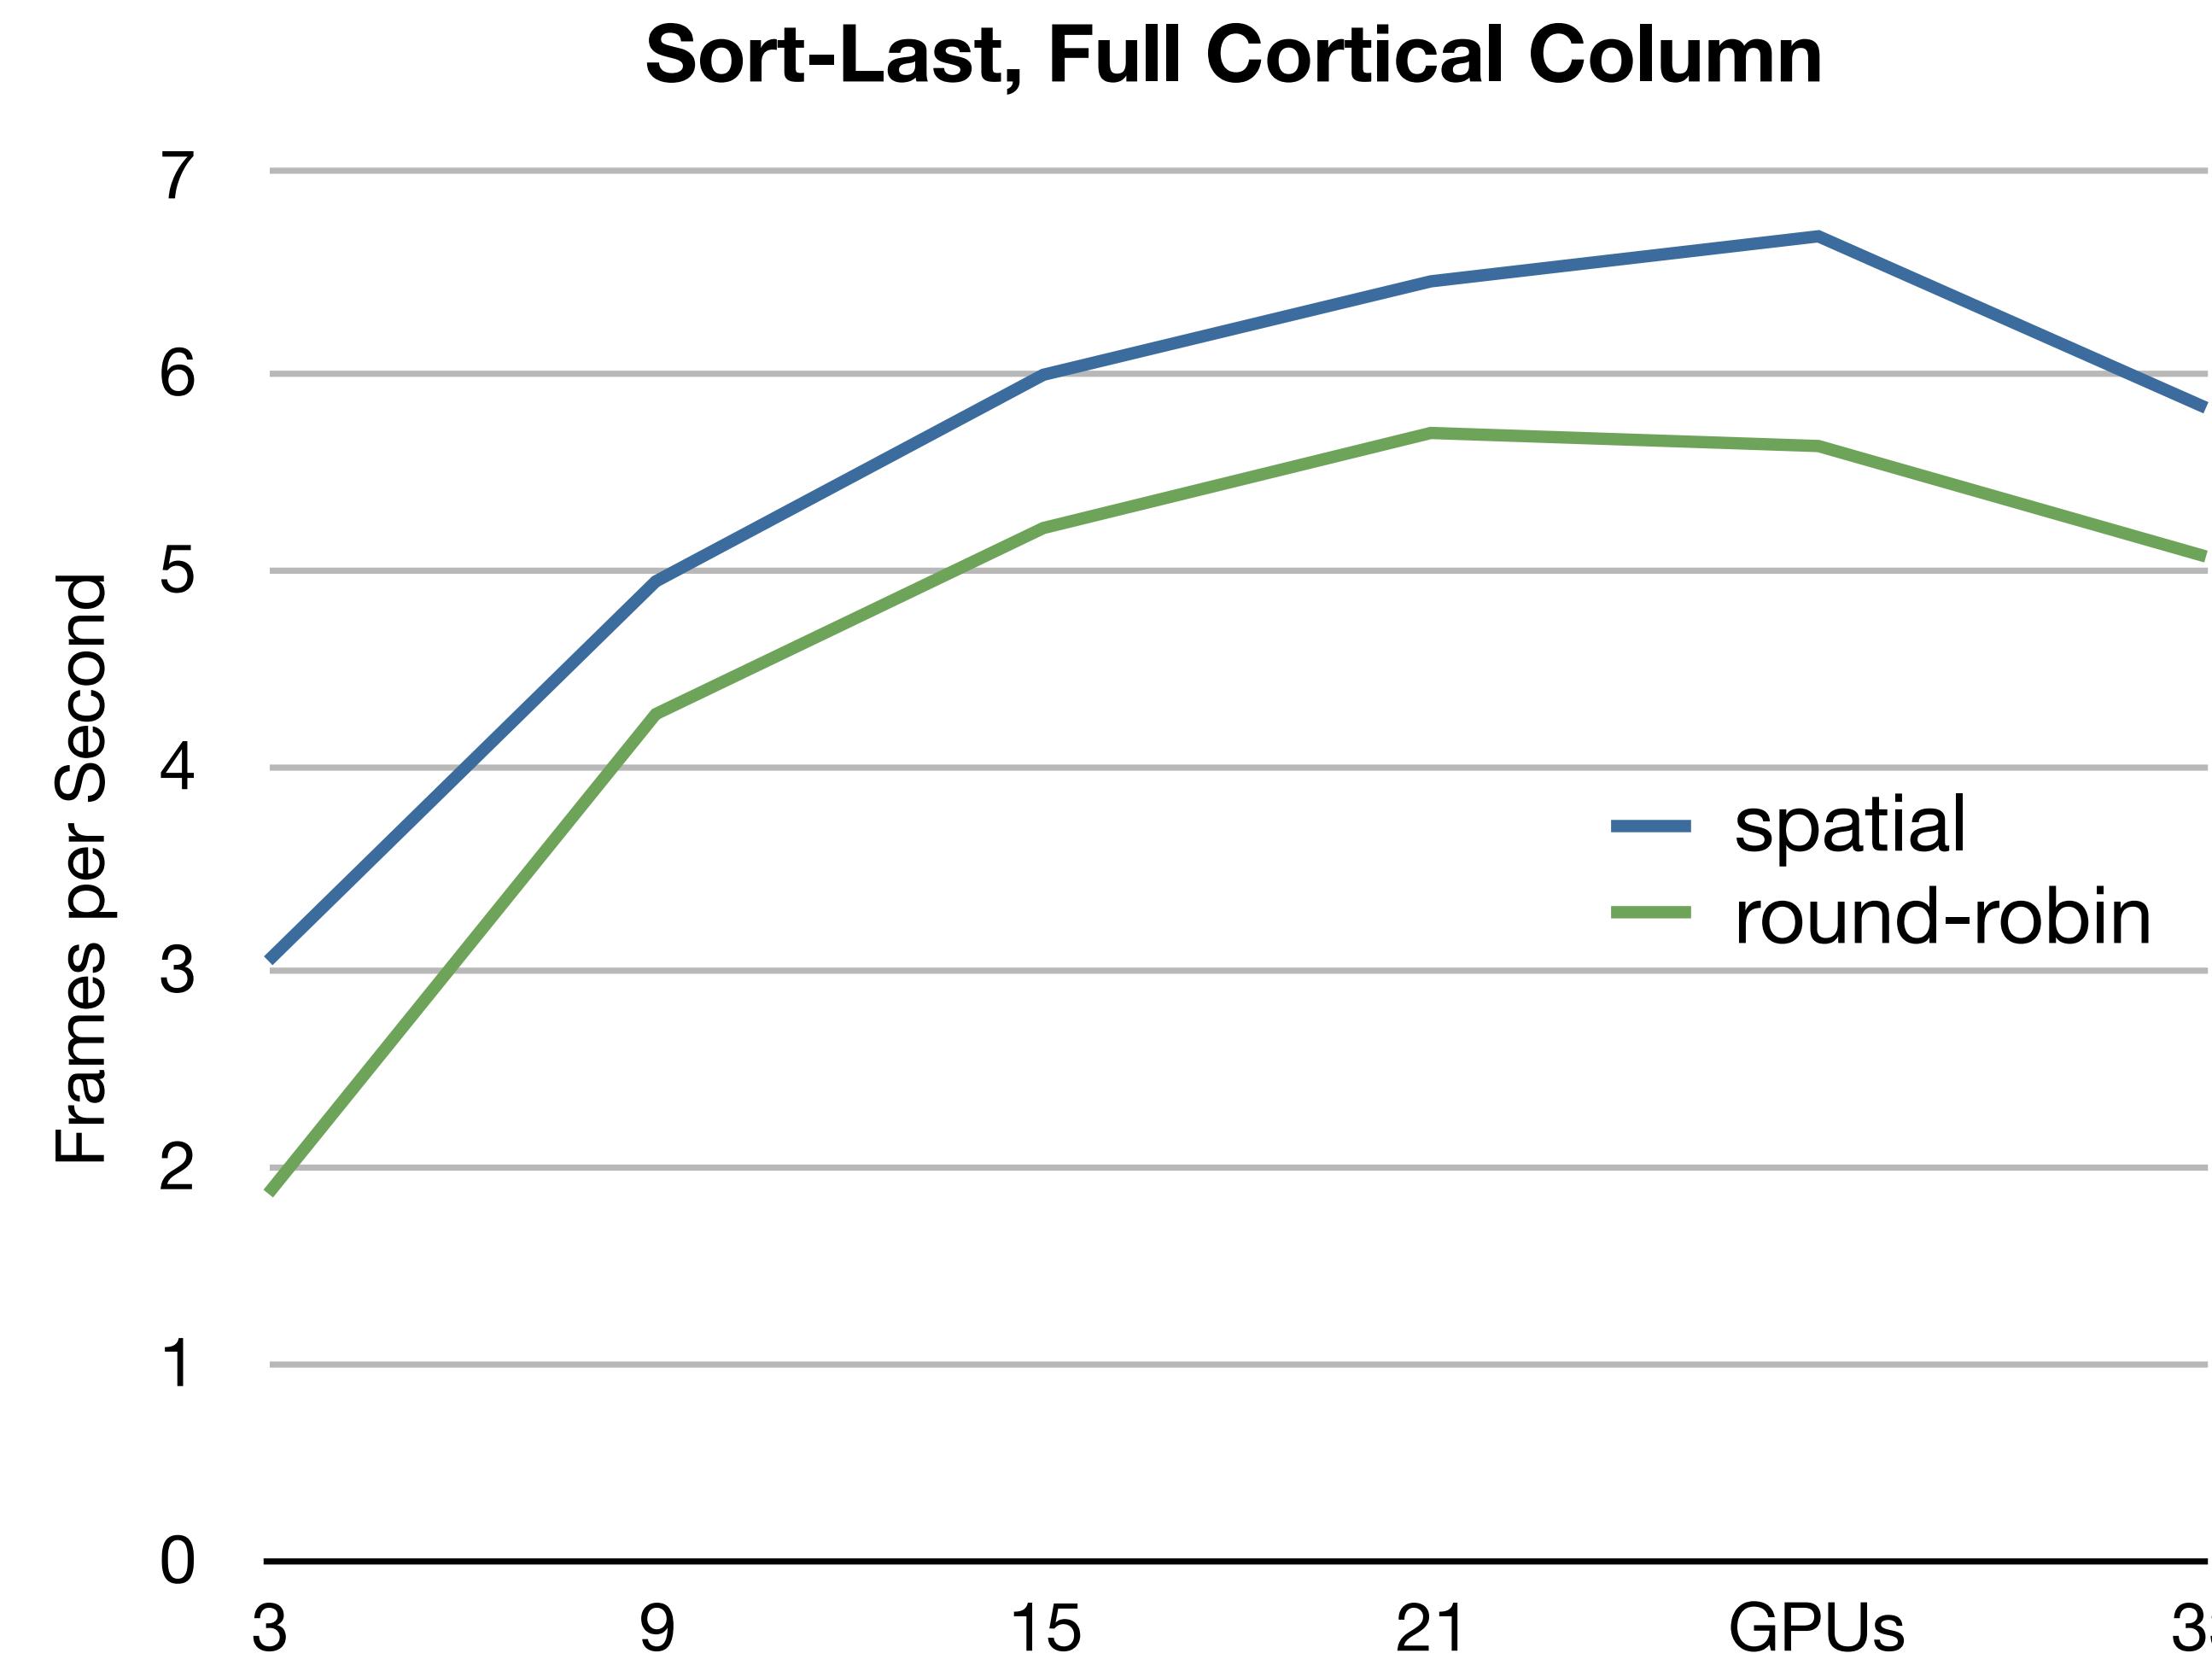
Asynchronous Execution



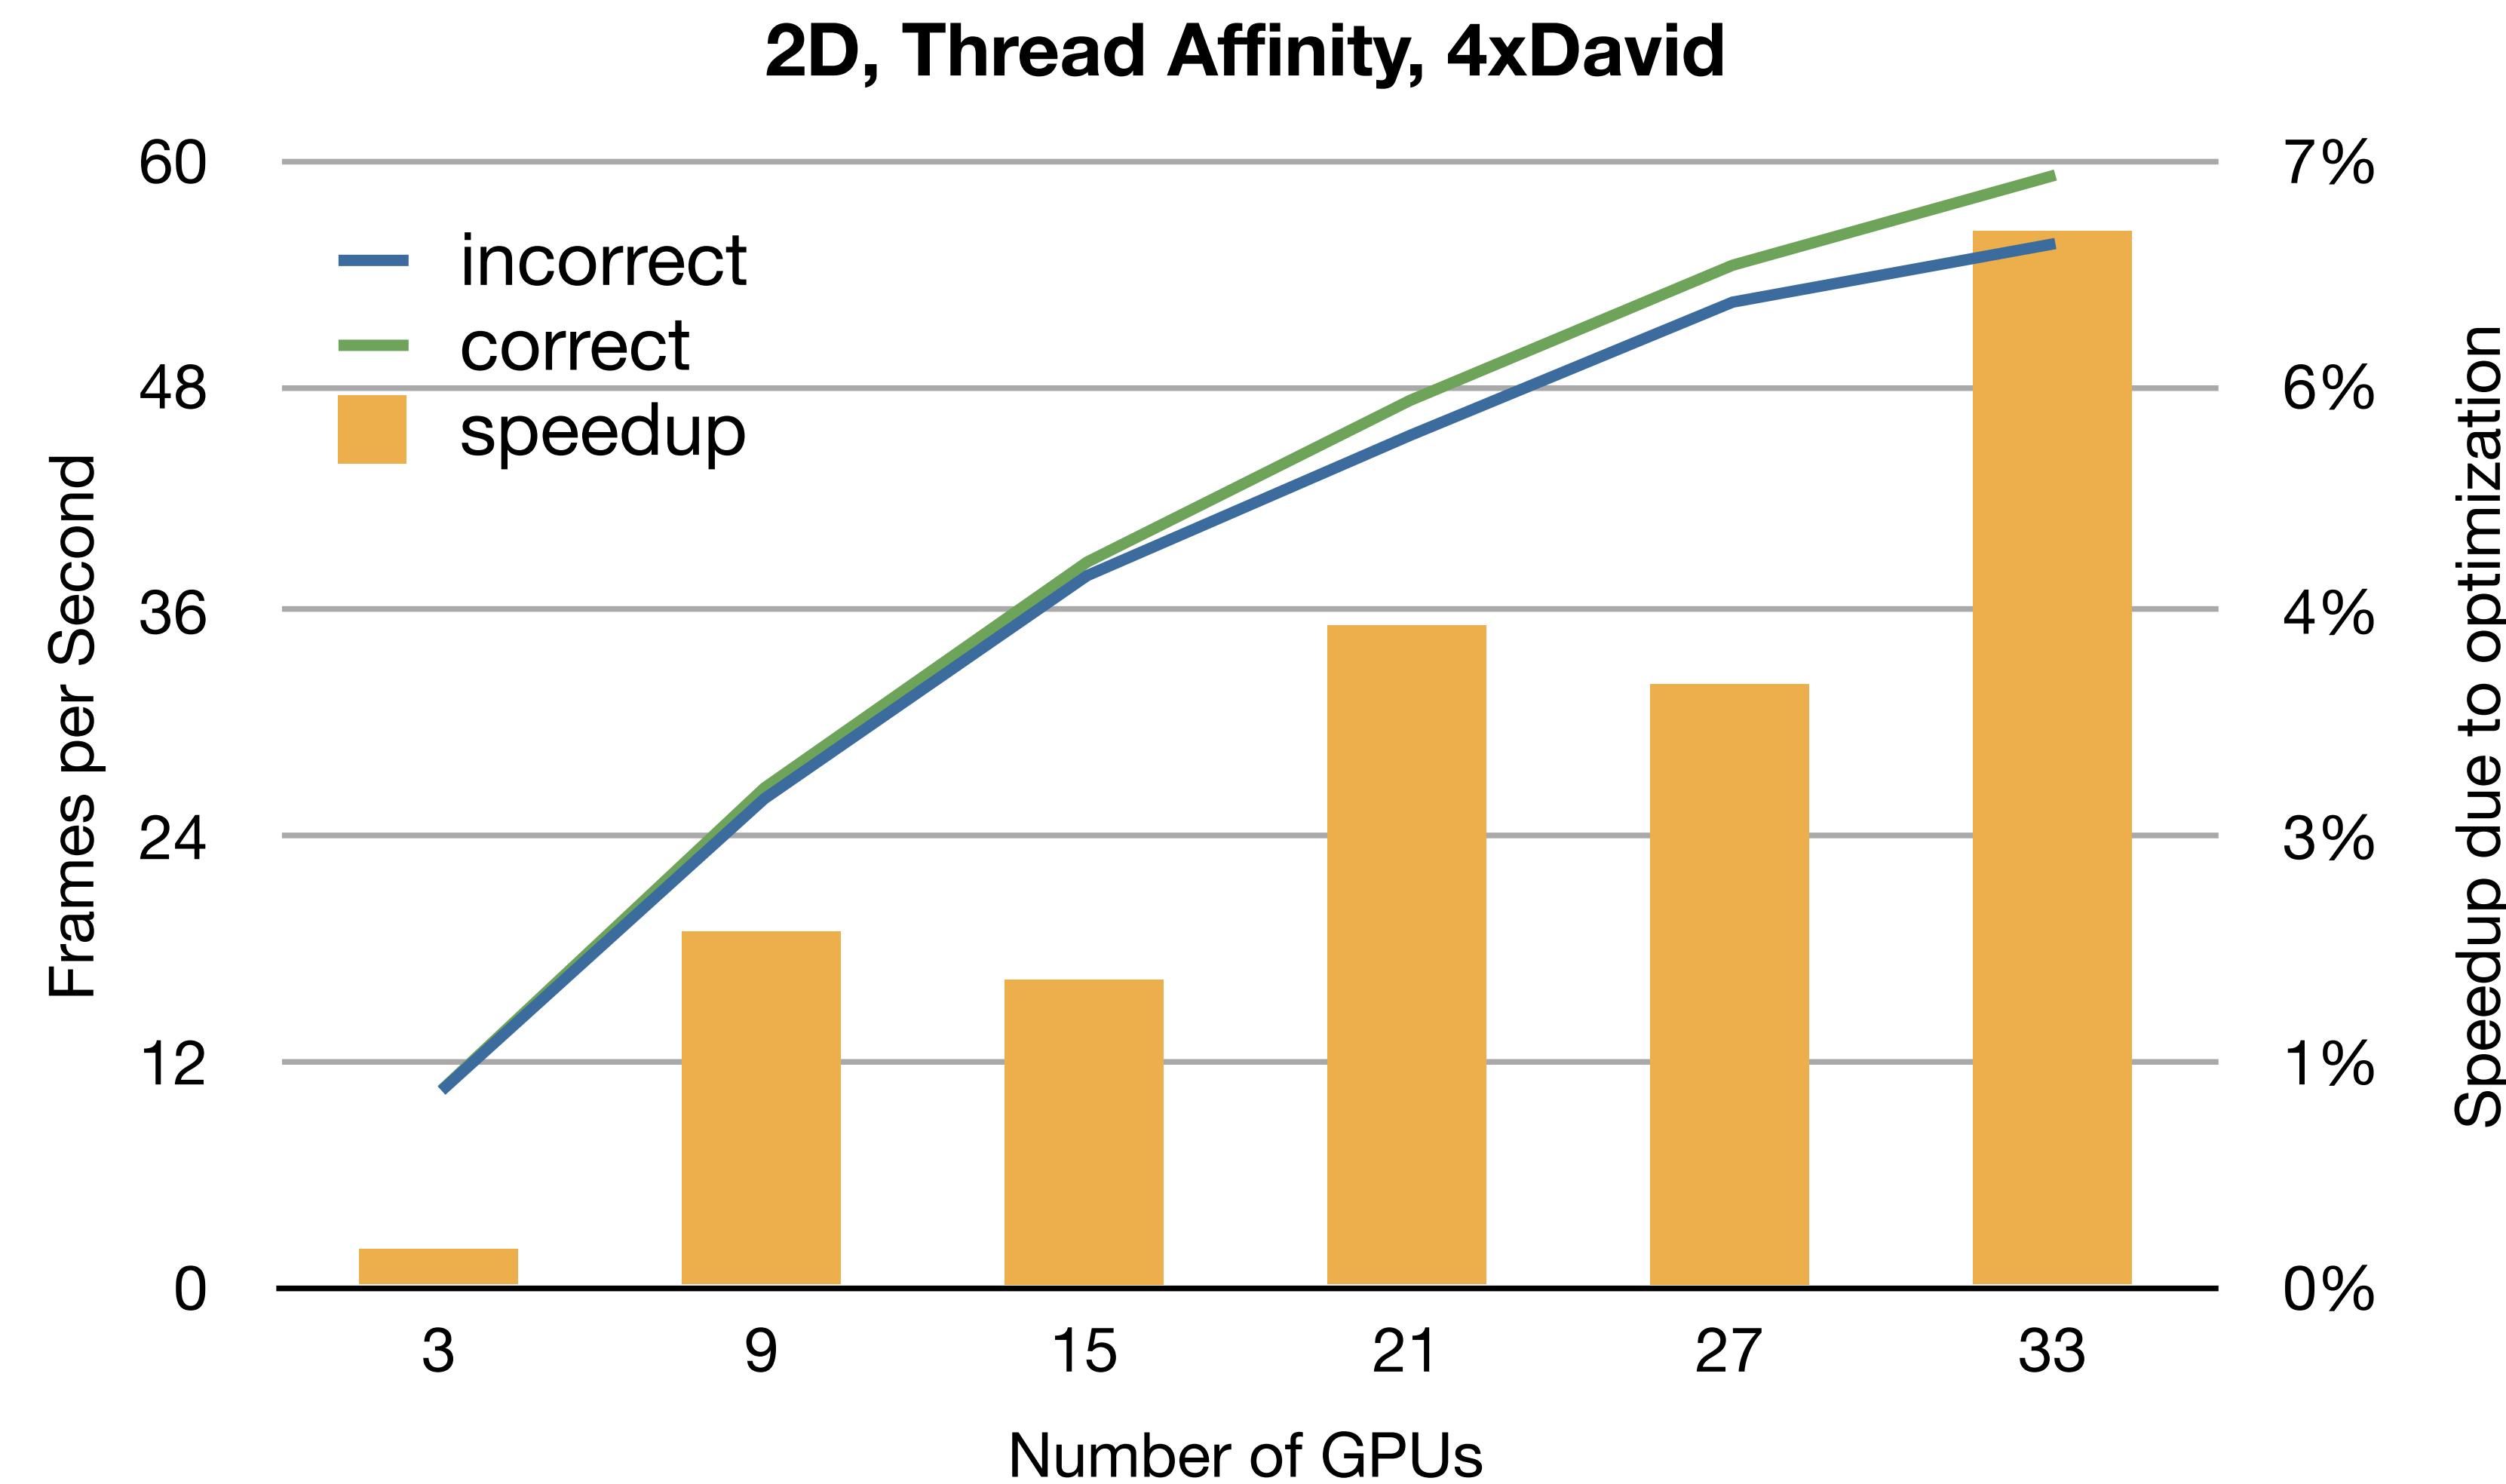
Cross-Display Load Balancing



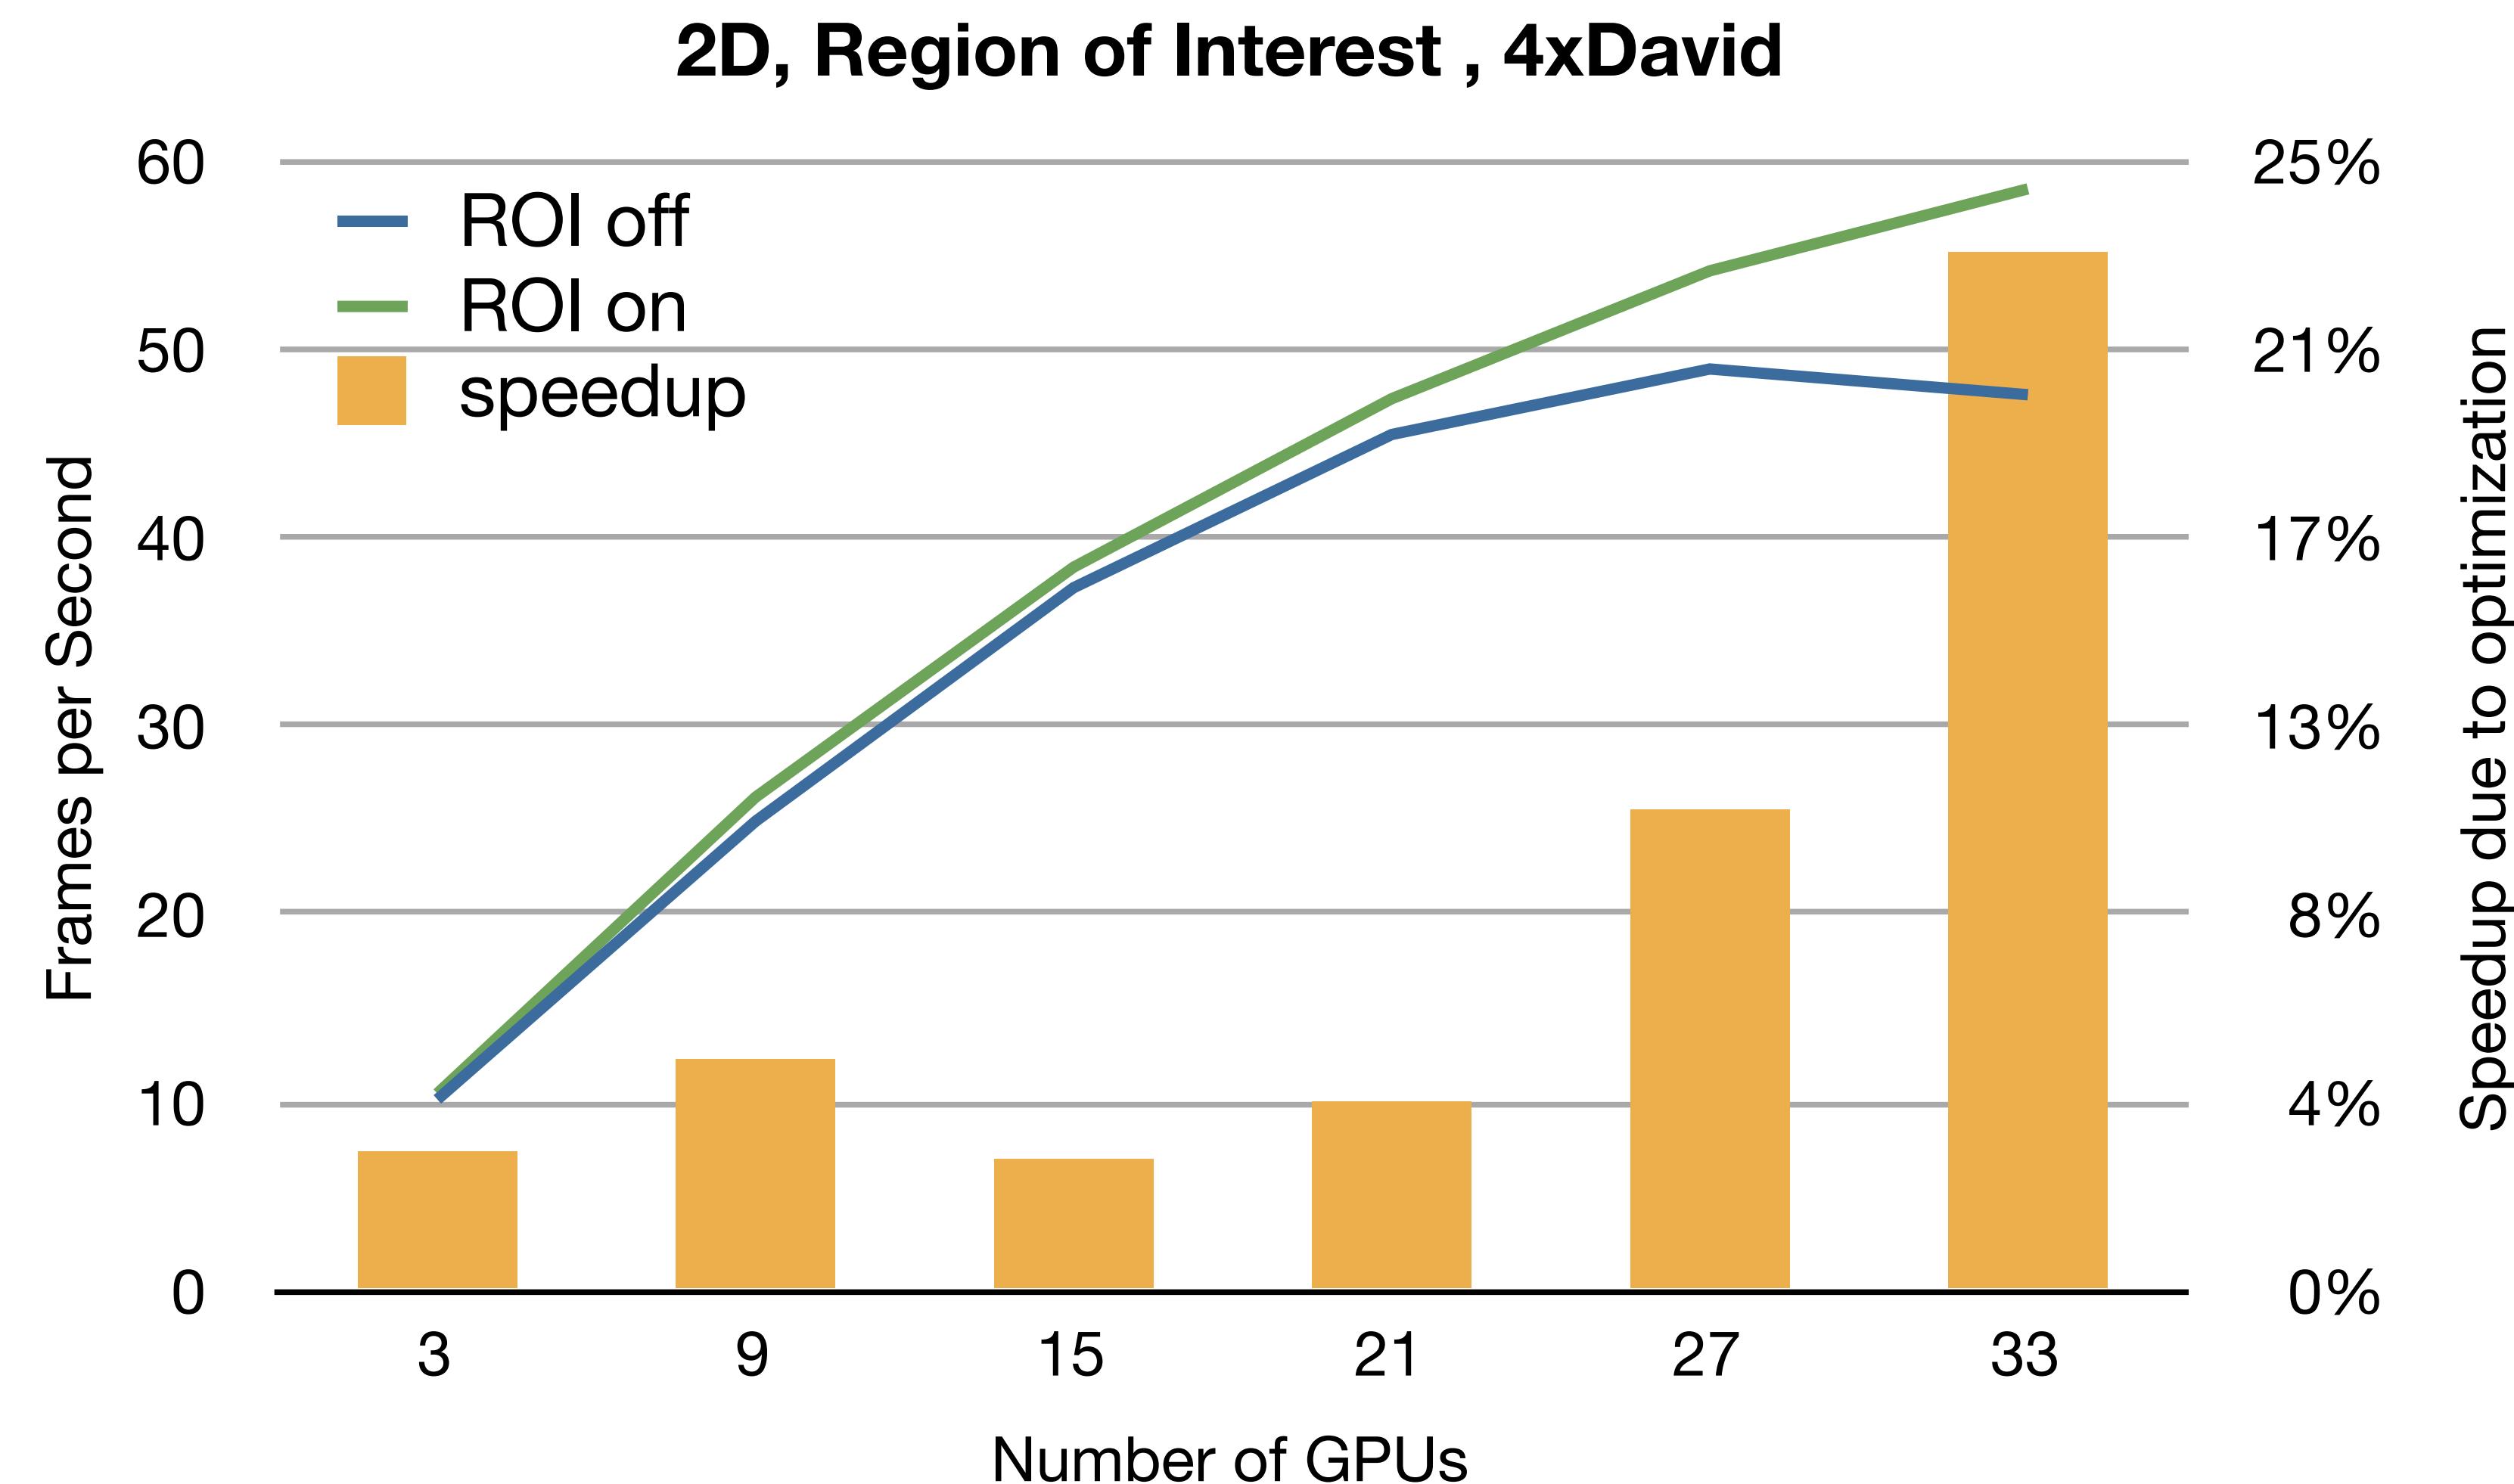
Spatial vs Random Sort-Last



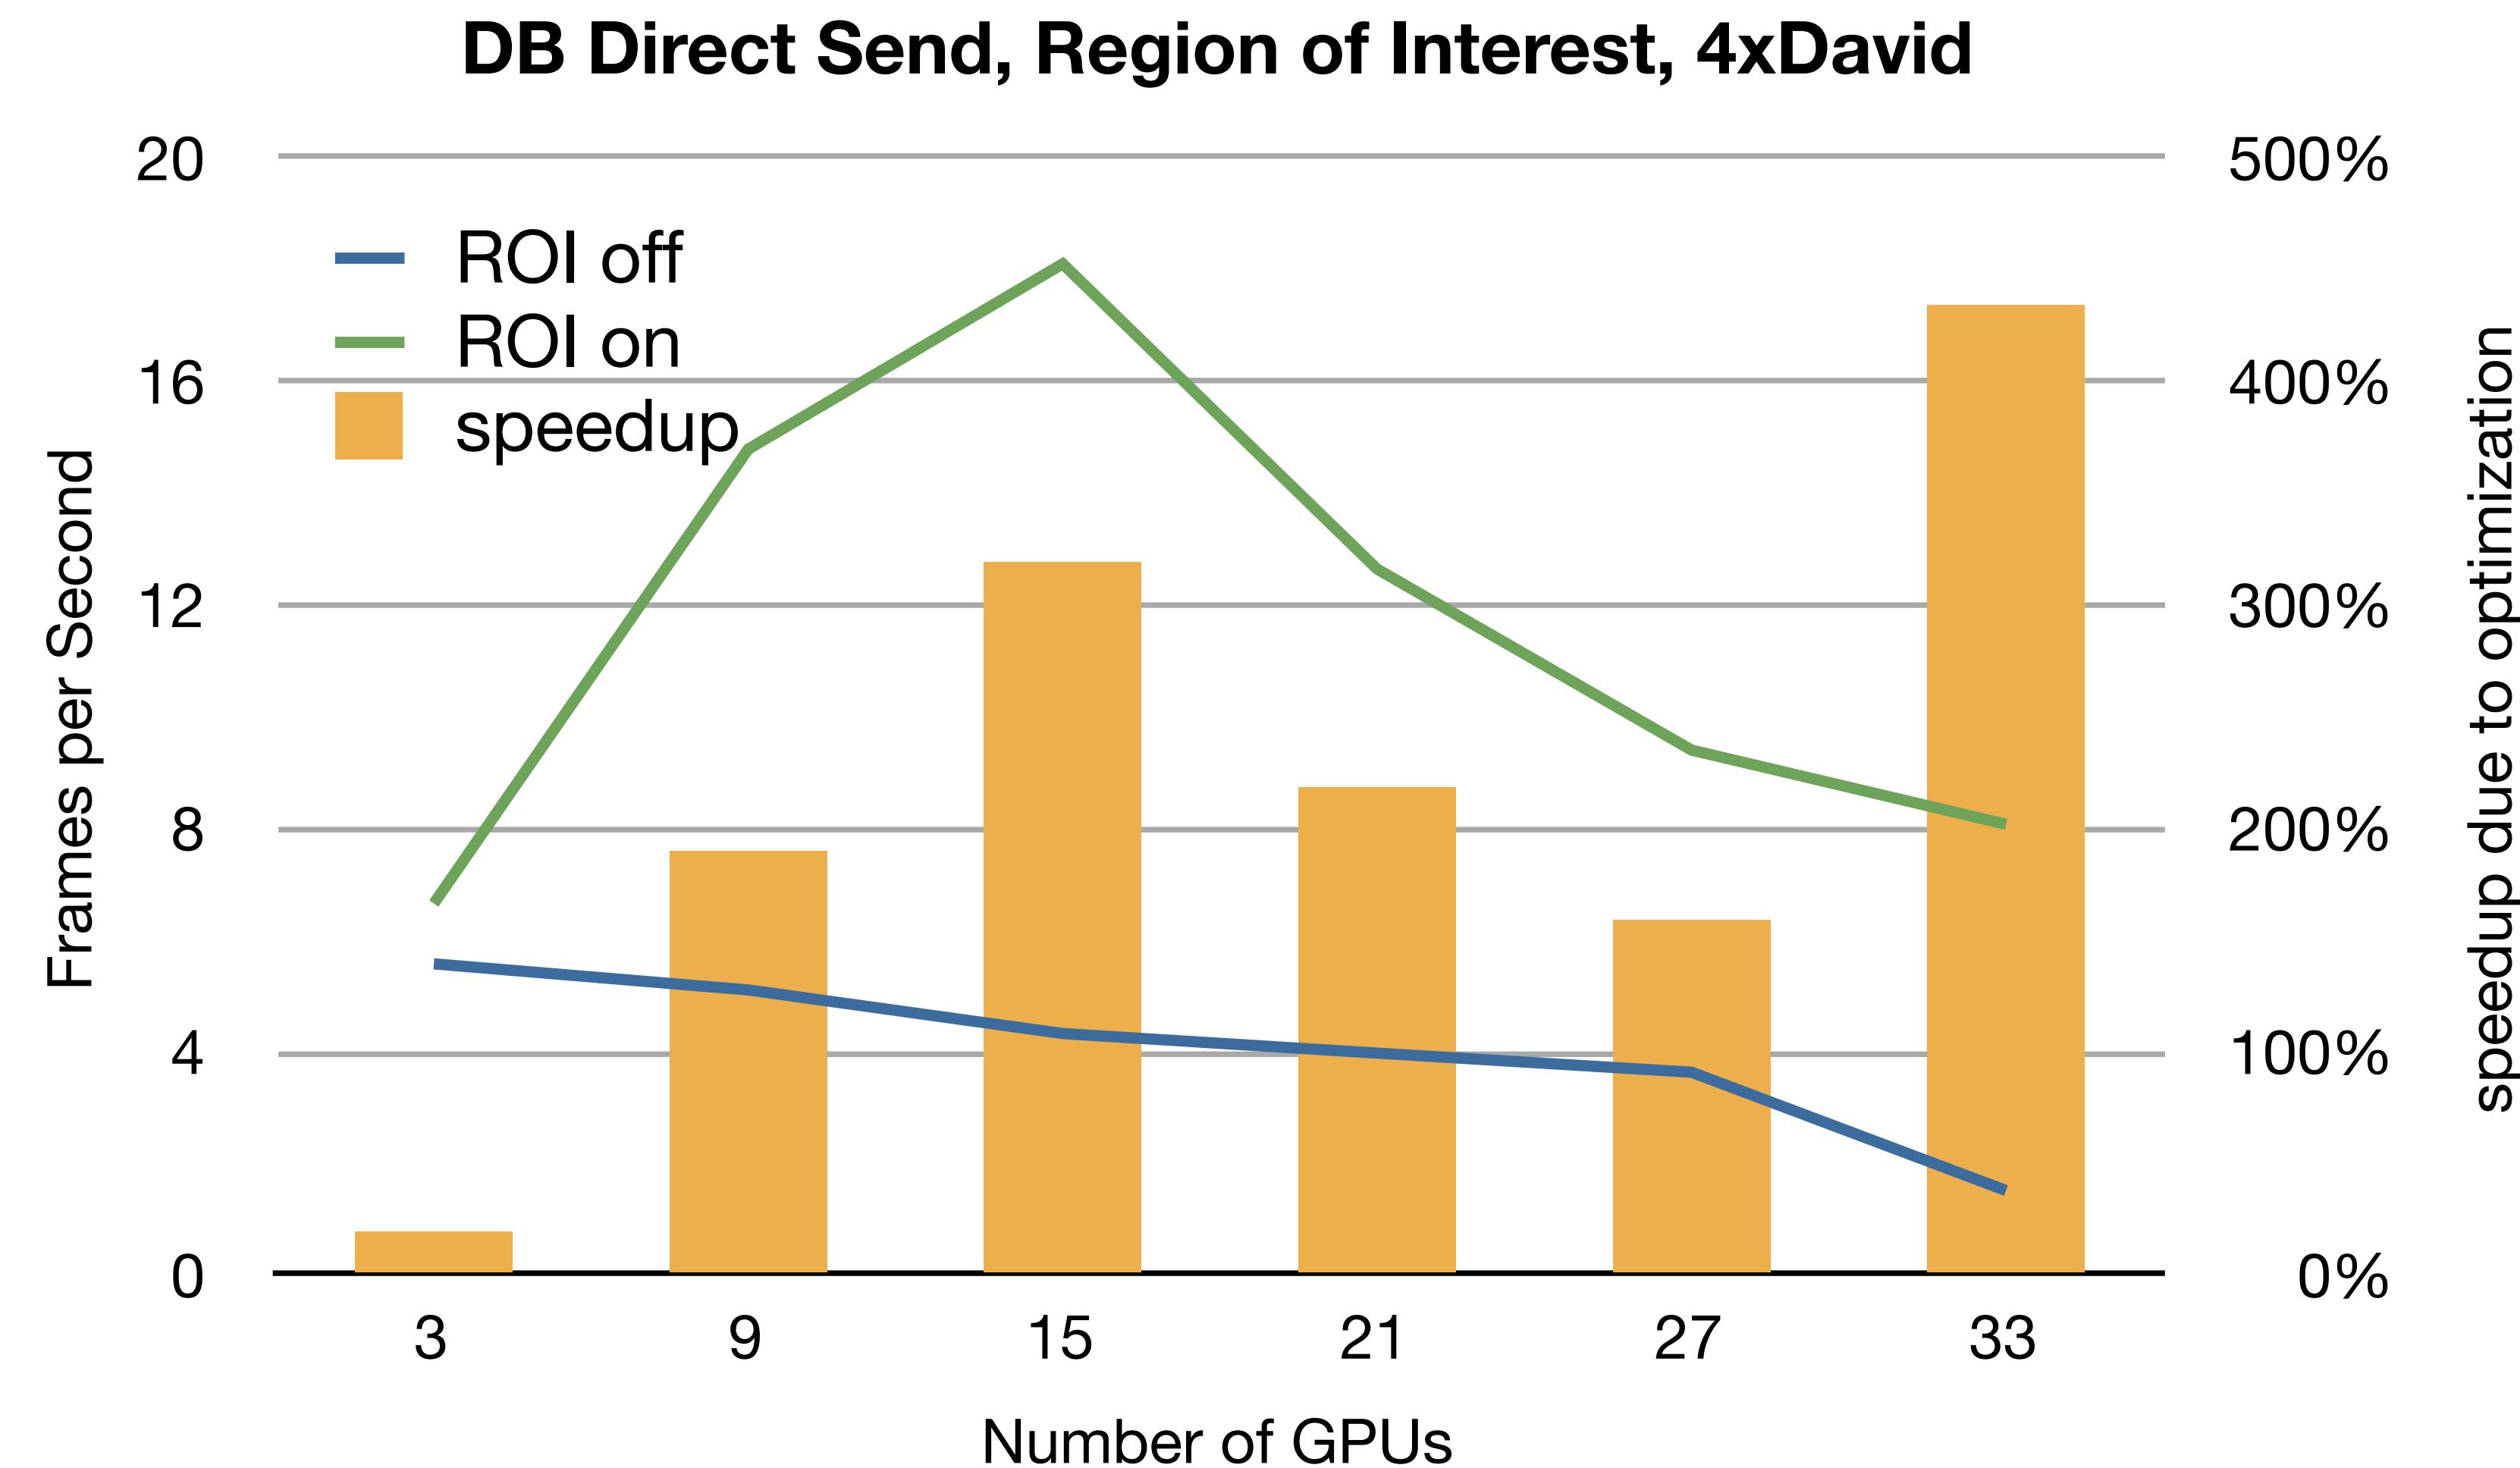
NUMA Thread Affinity



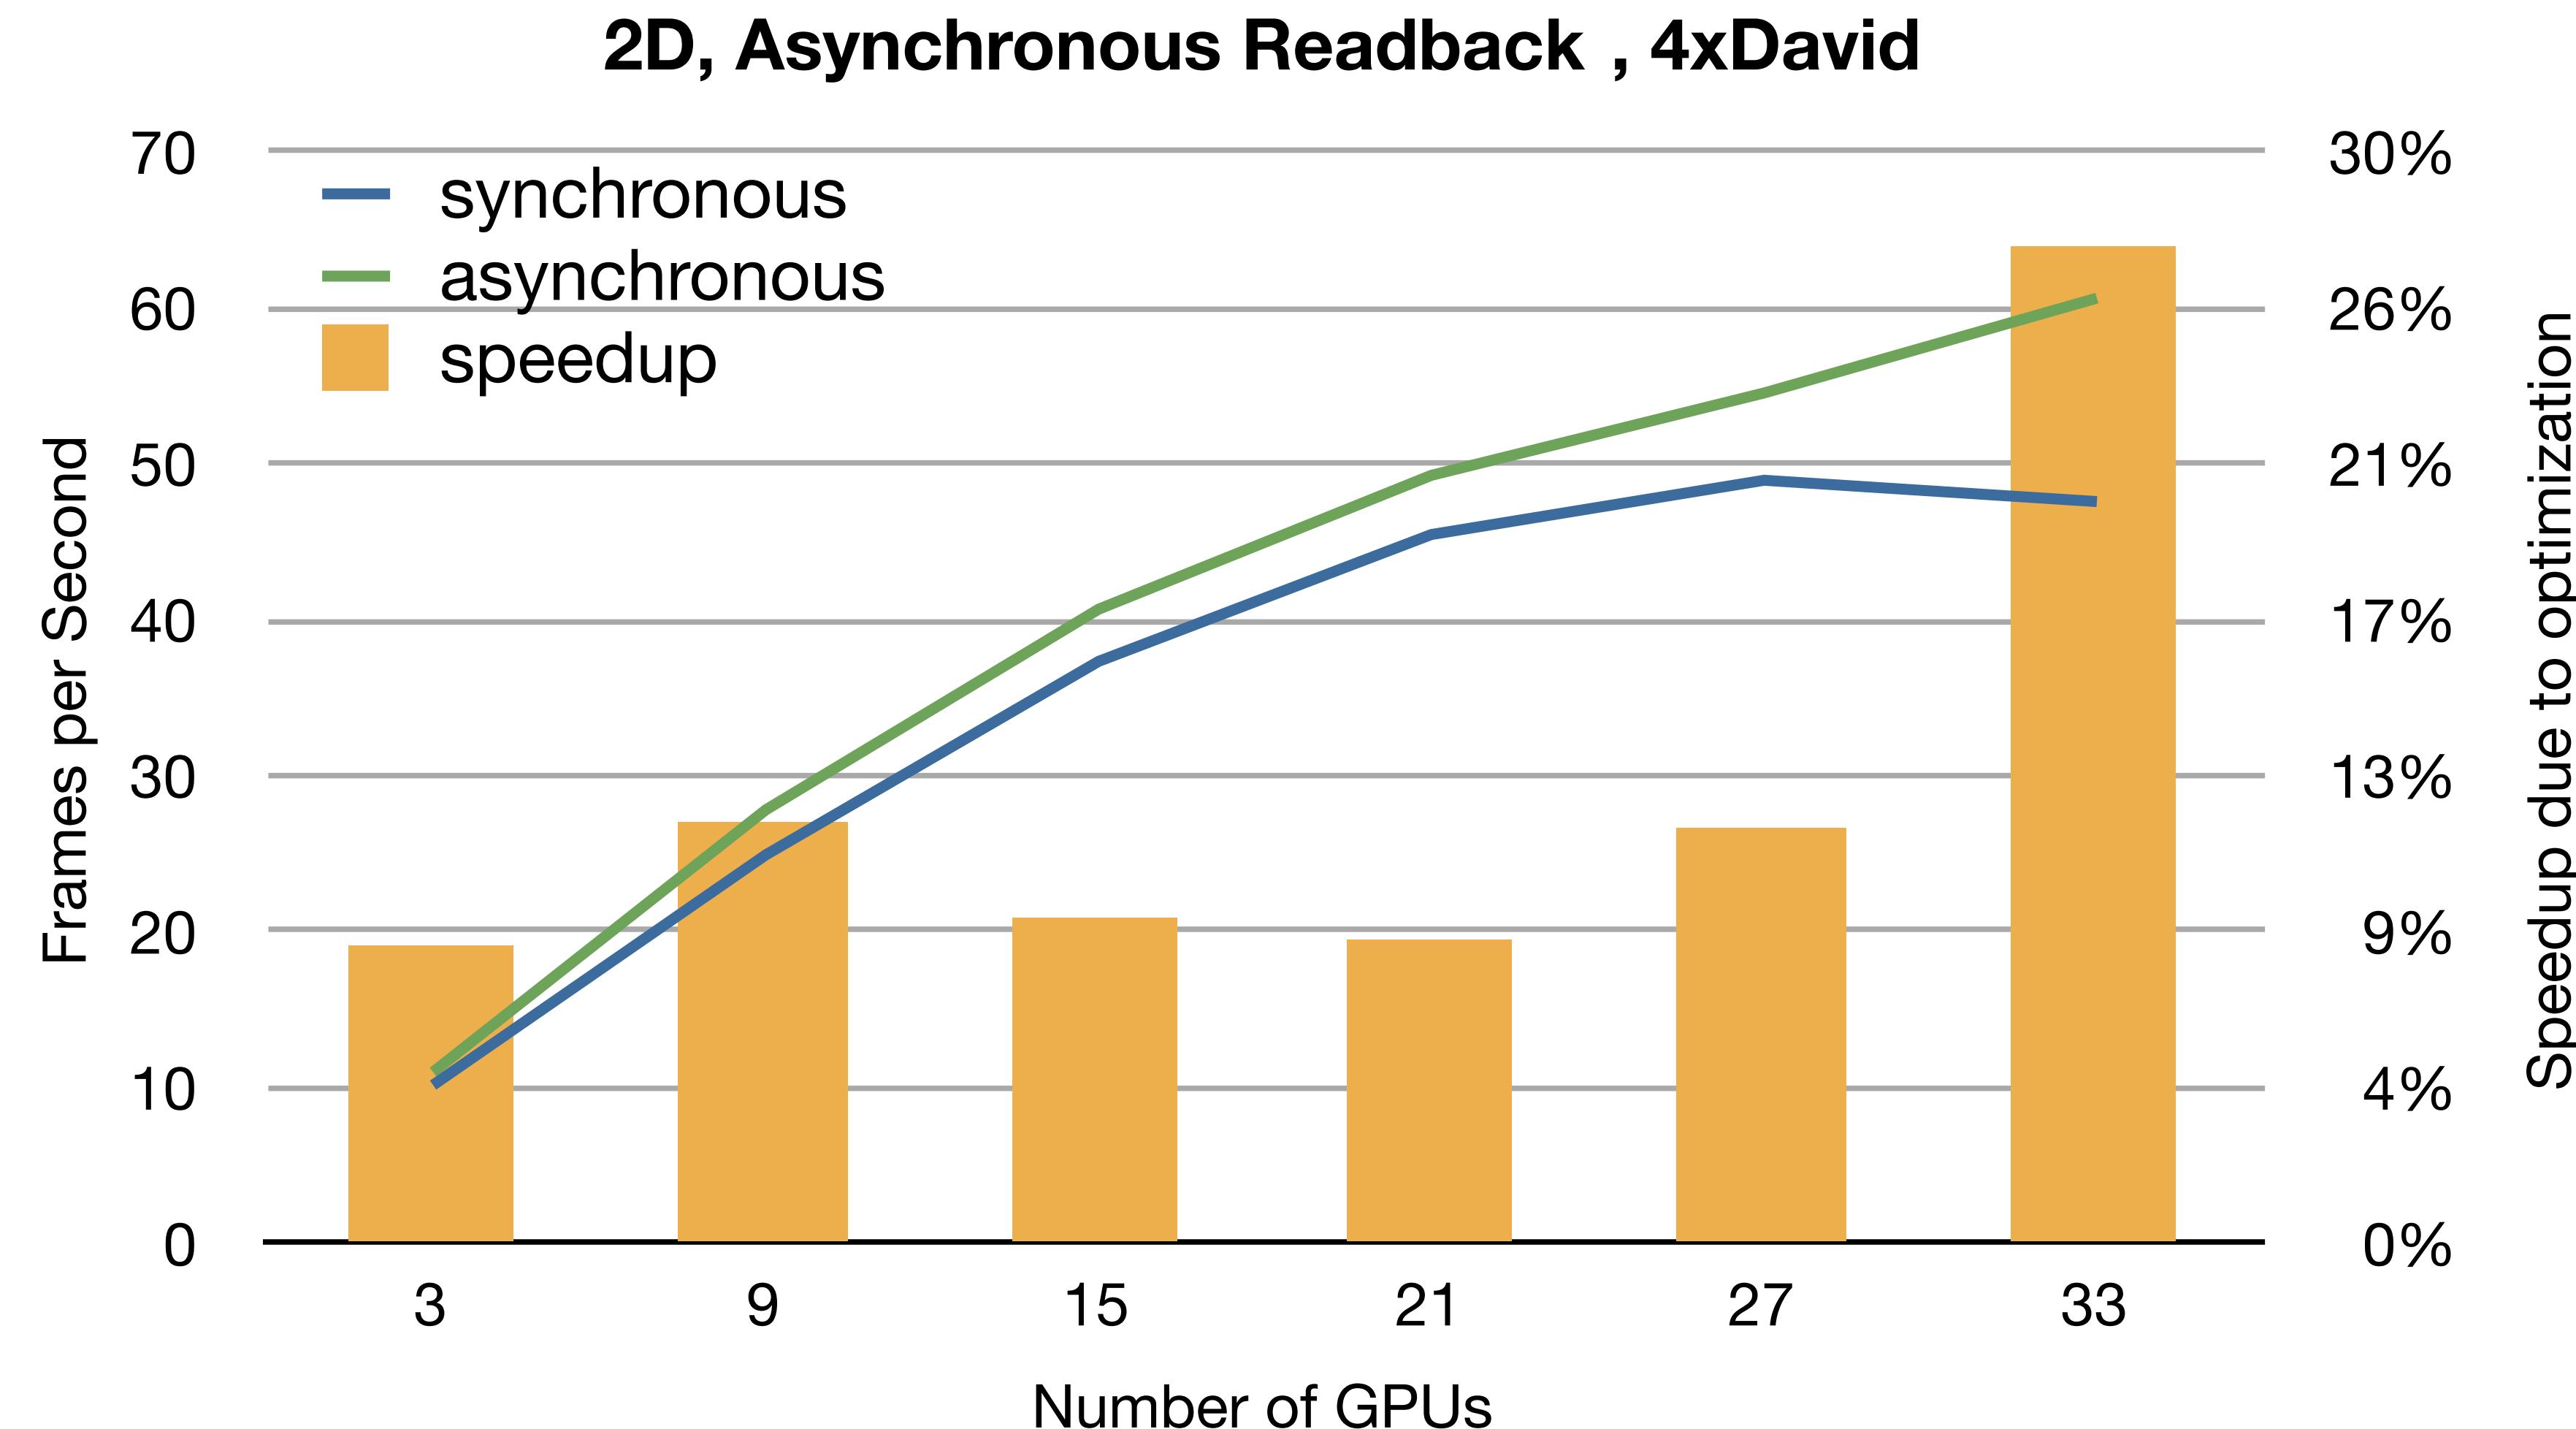
Sort-First Region of Interest



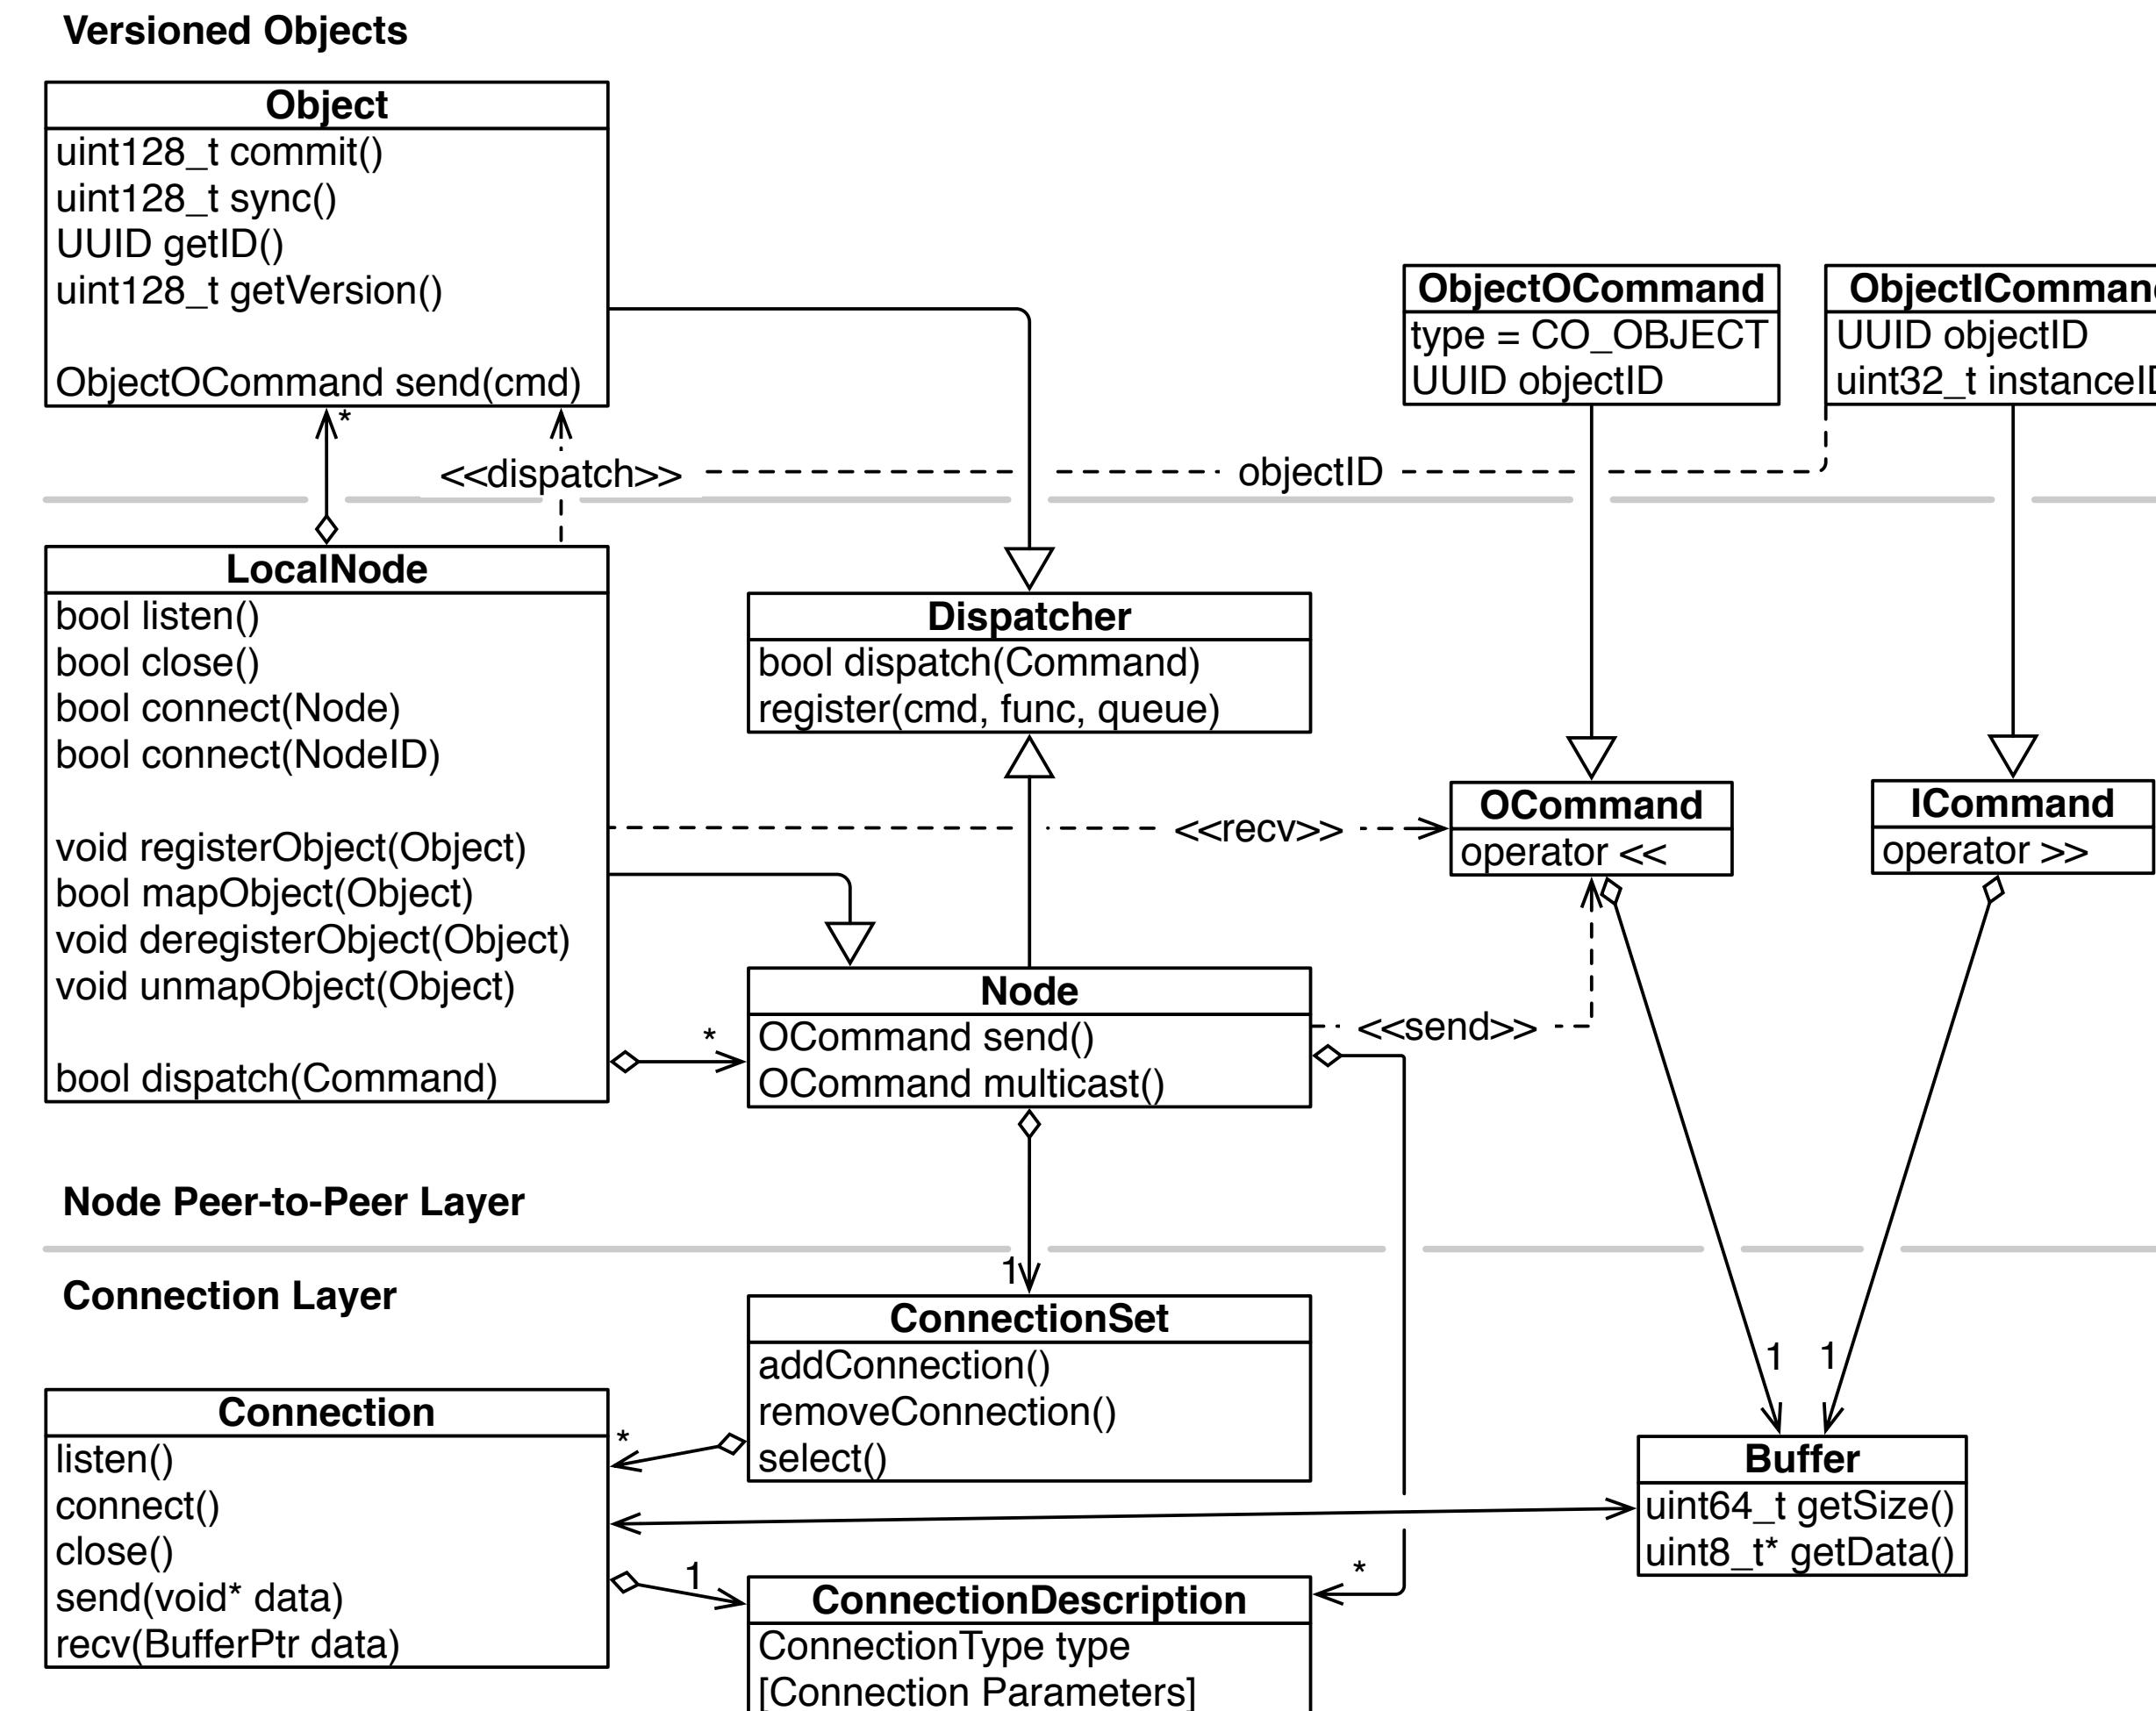
Sort-Last Region of Interest



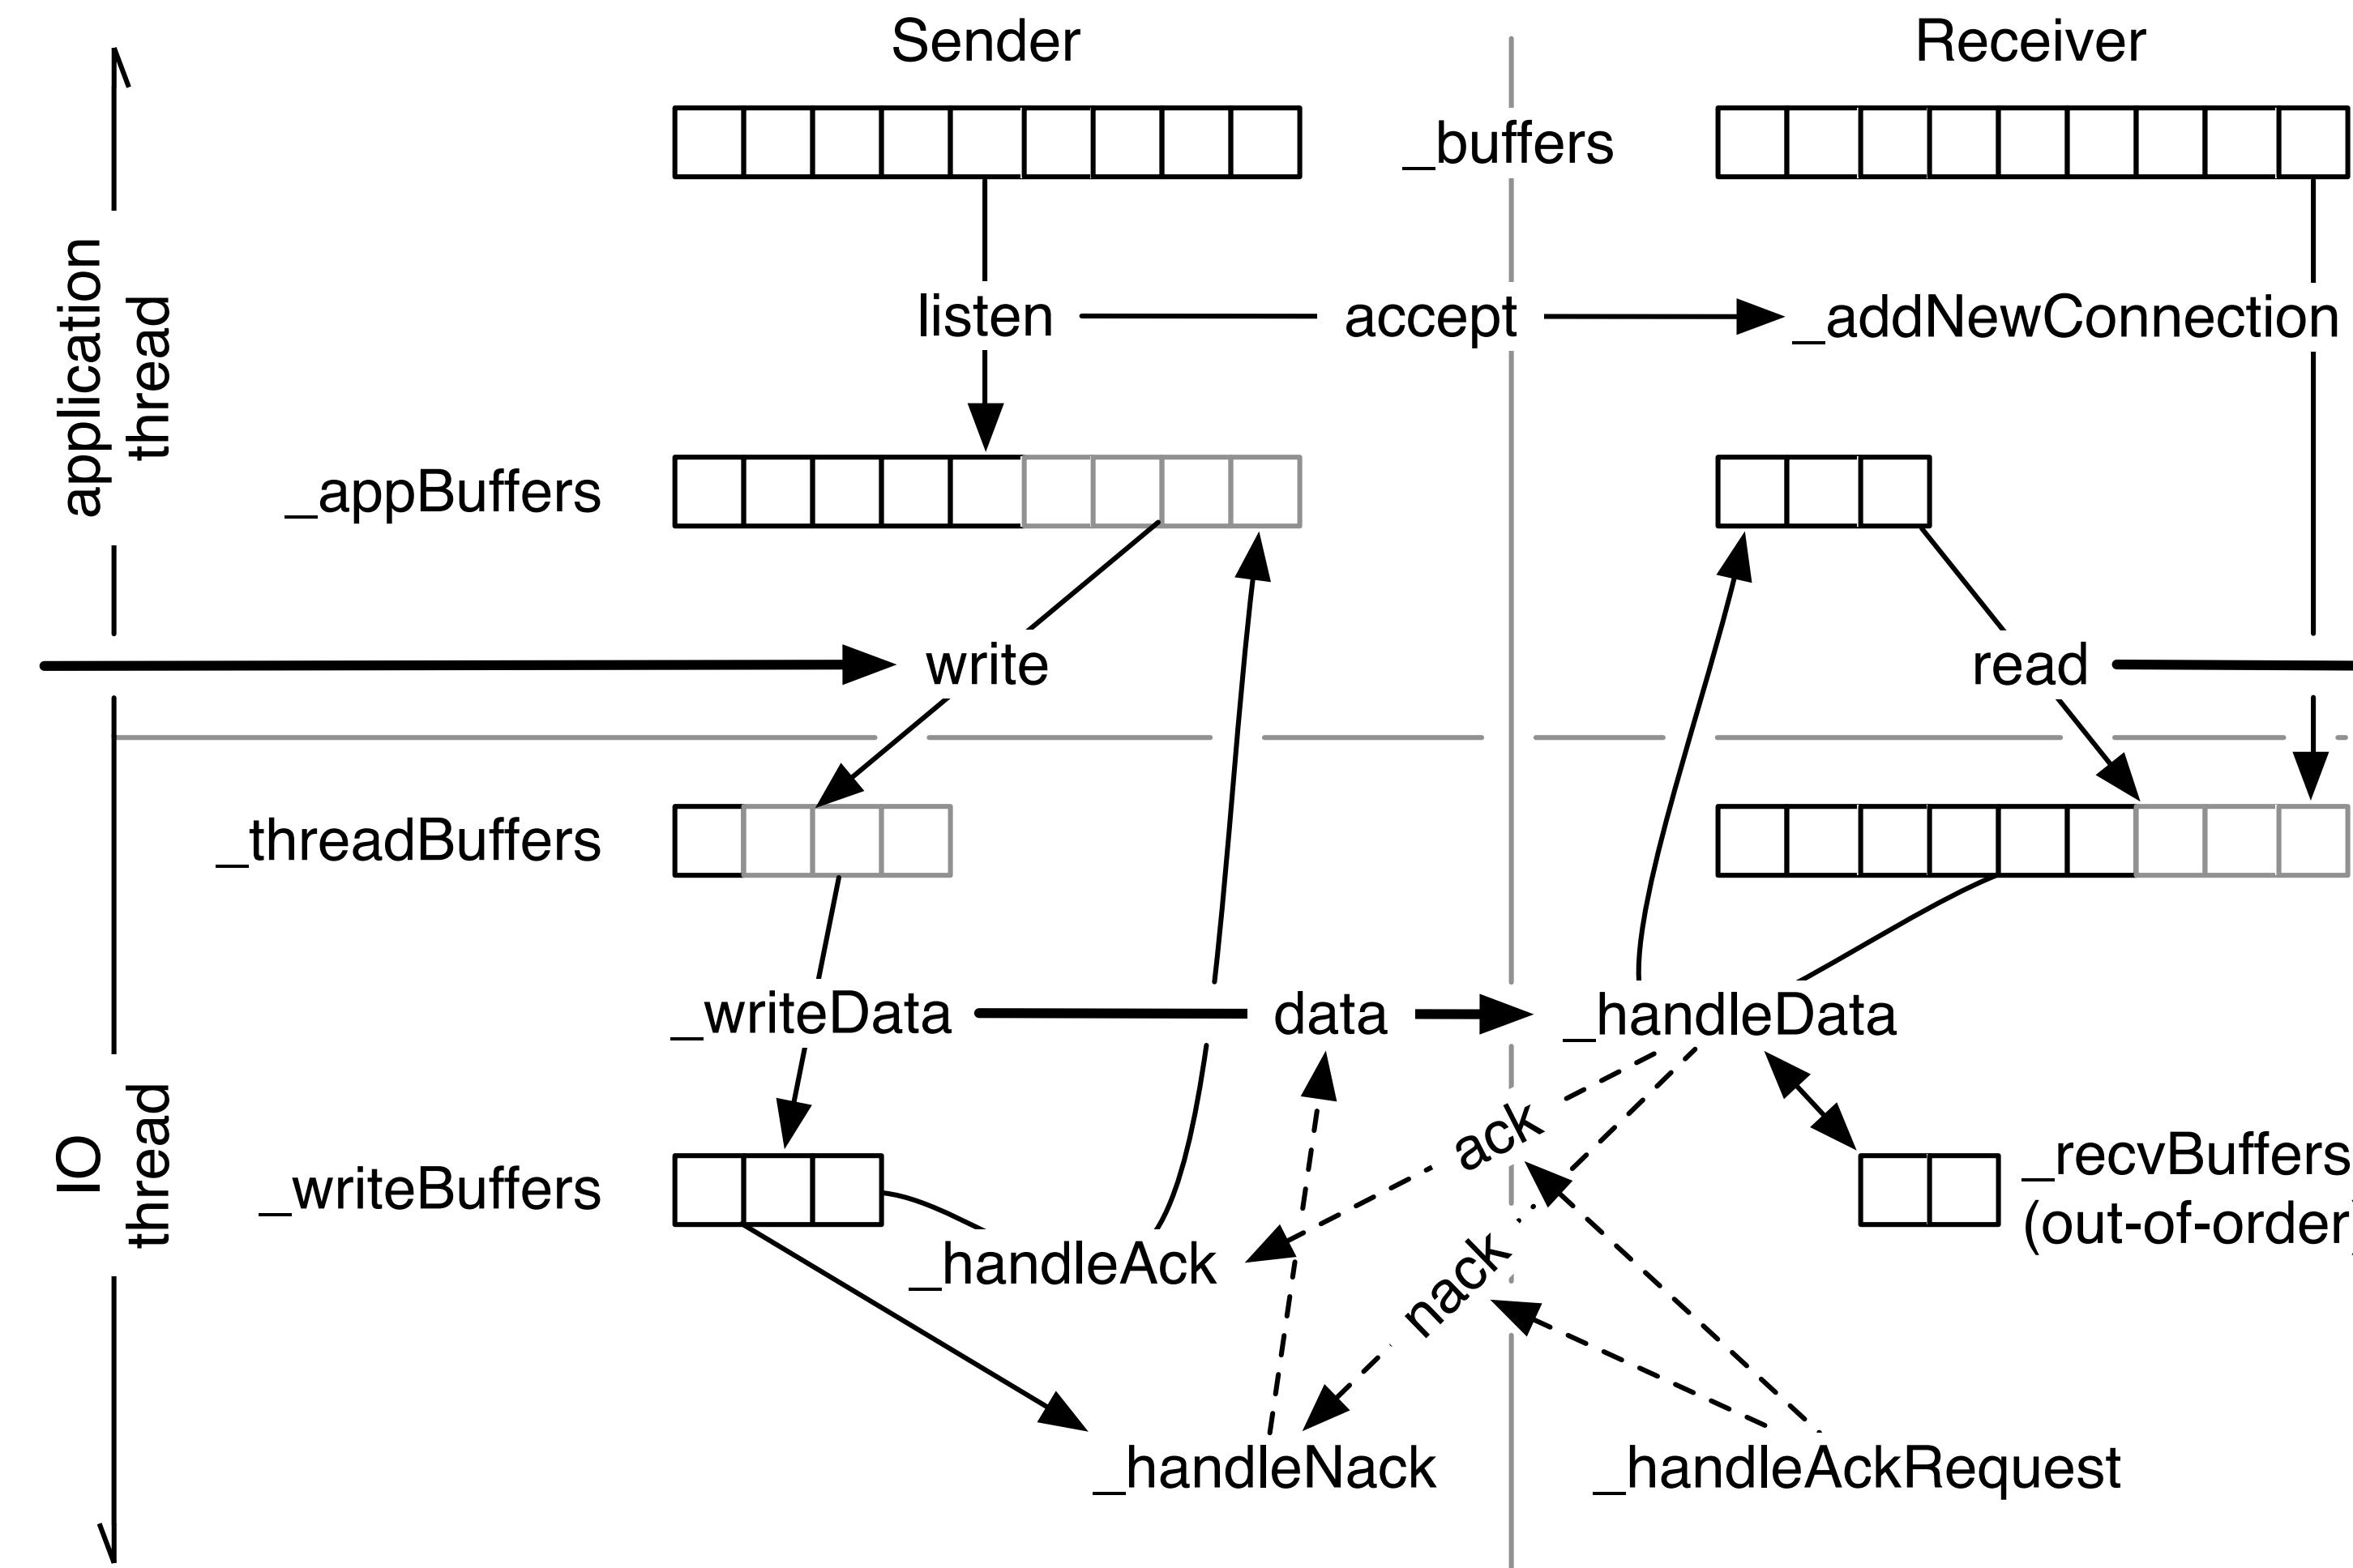
Asynchronous Readback



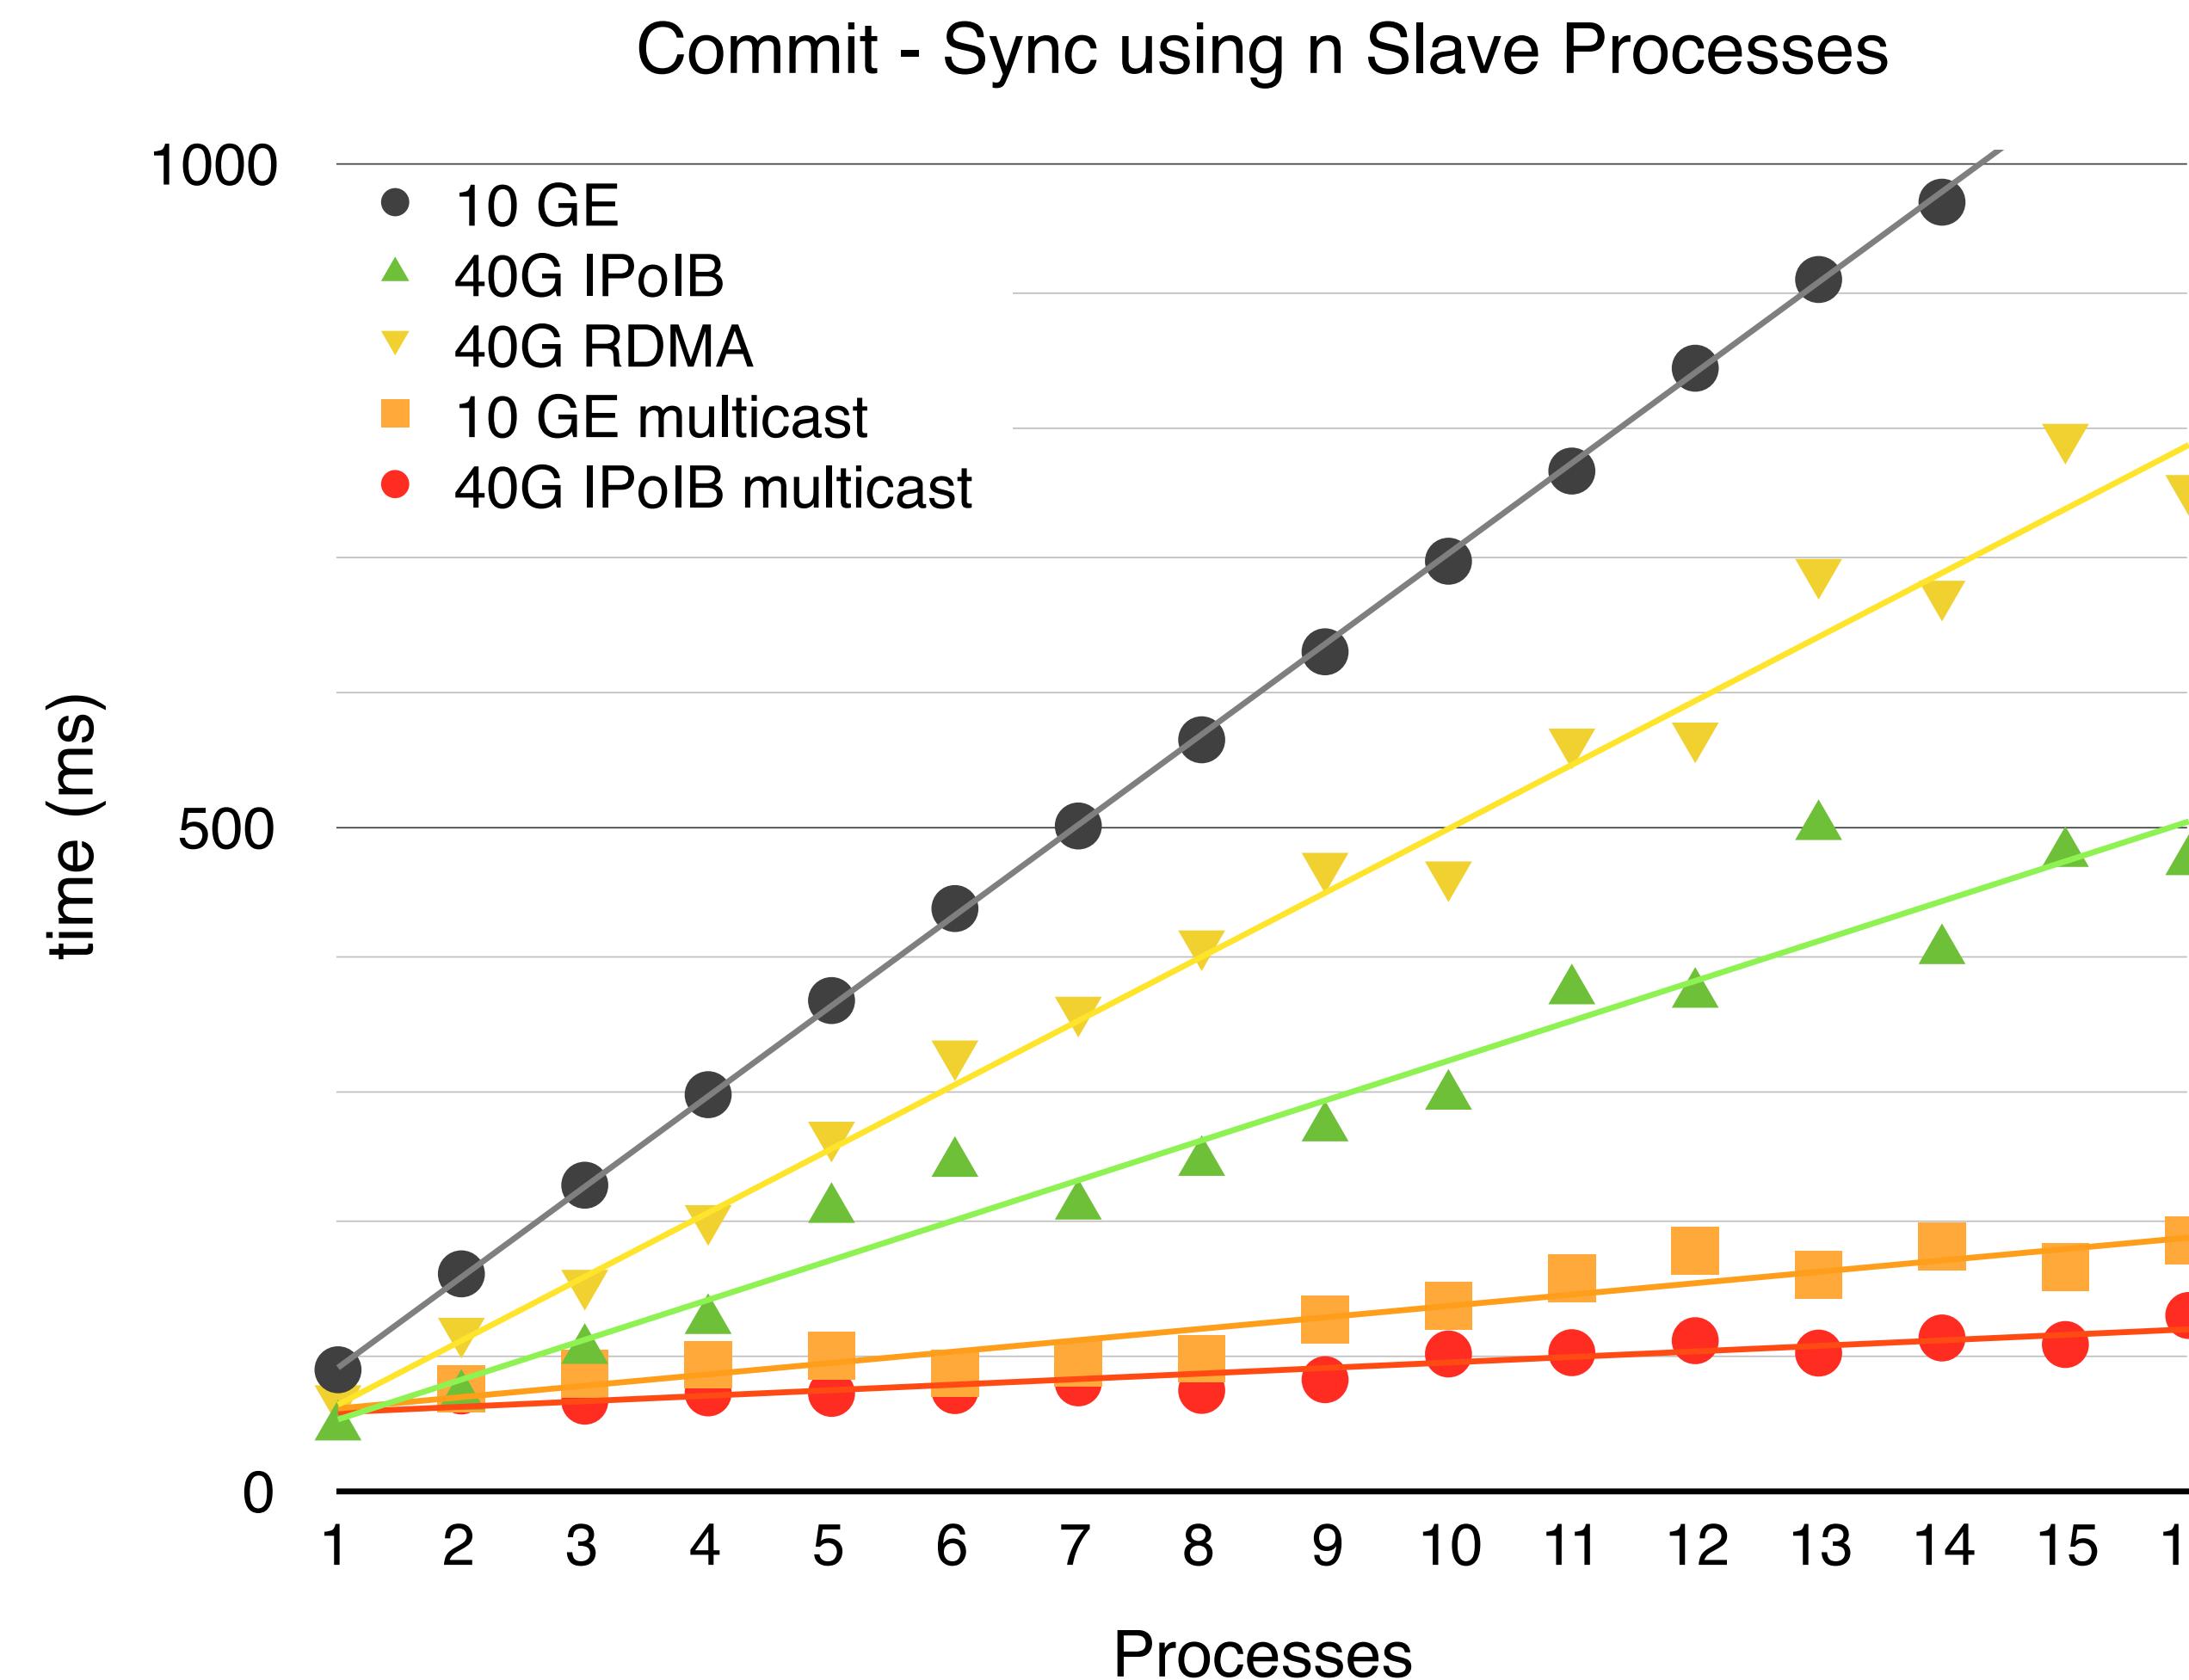
Collage Architecture



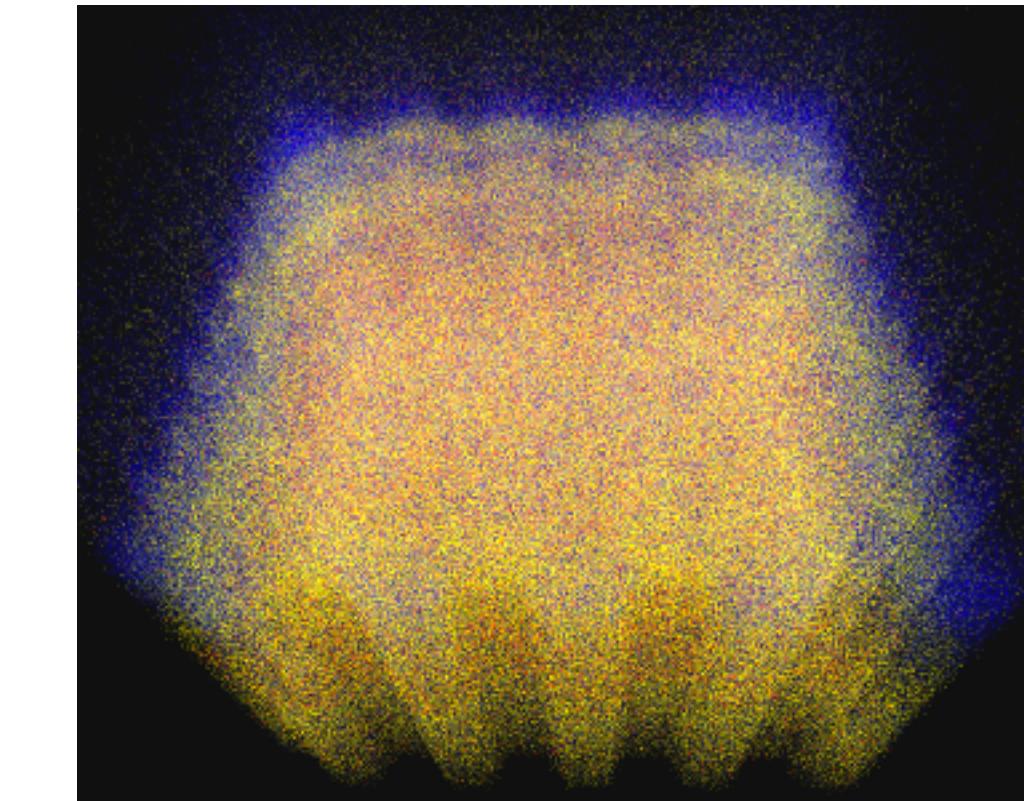
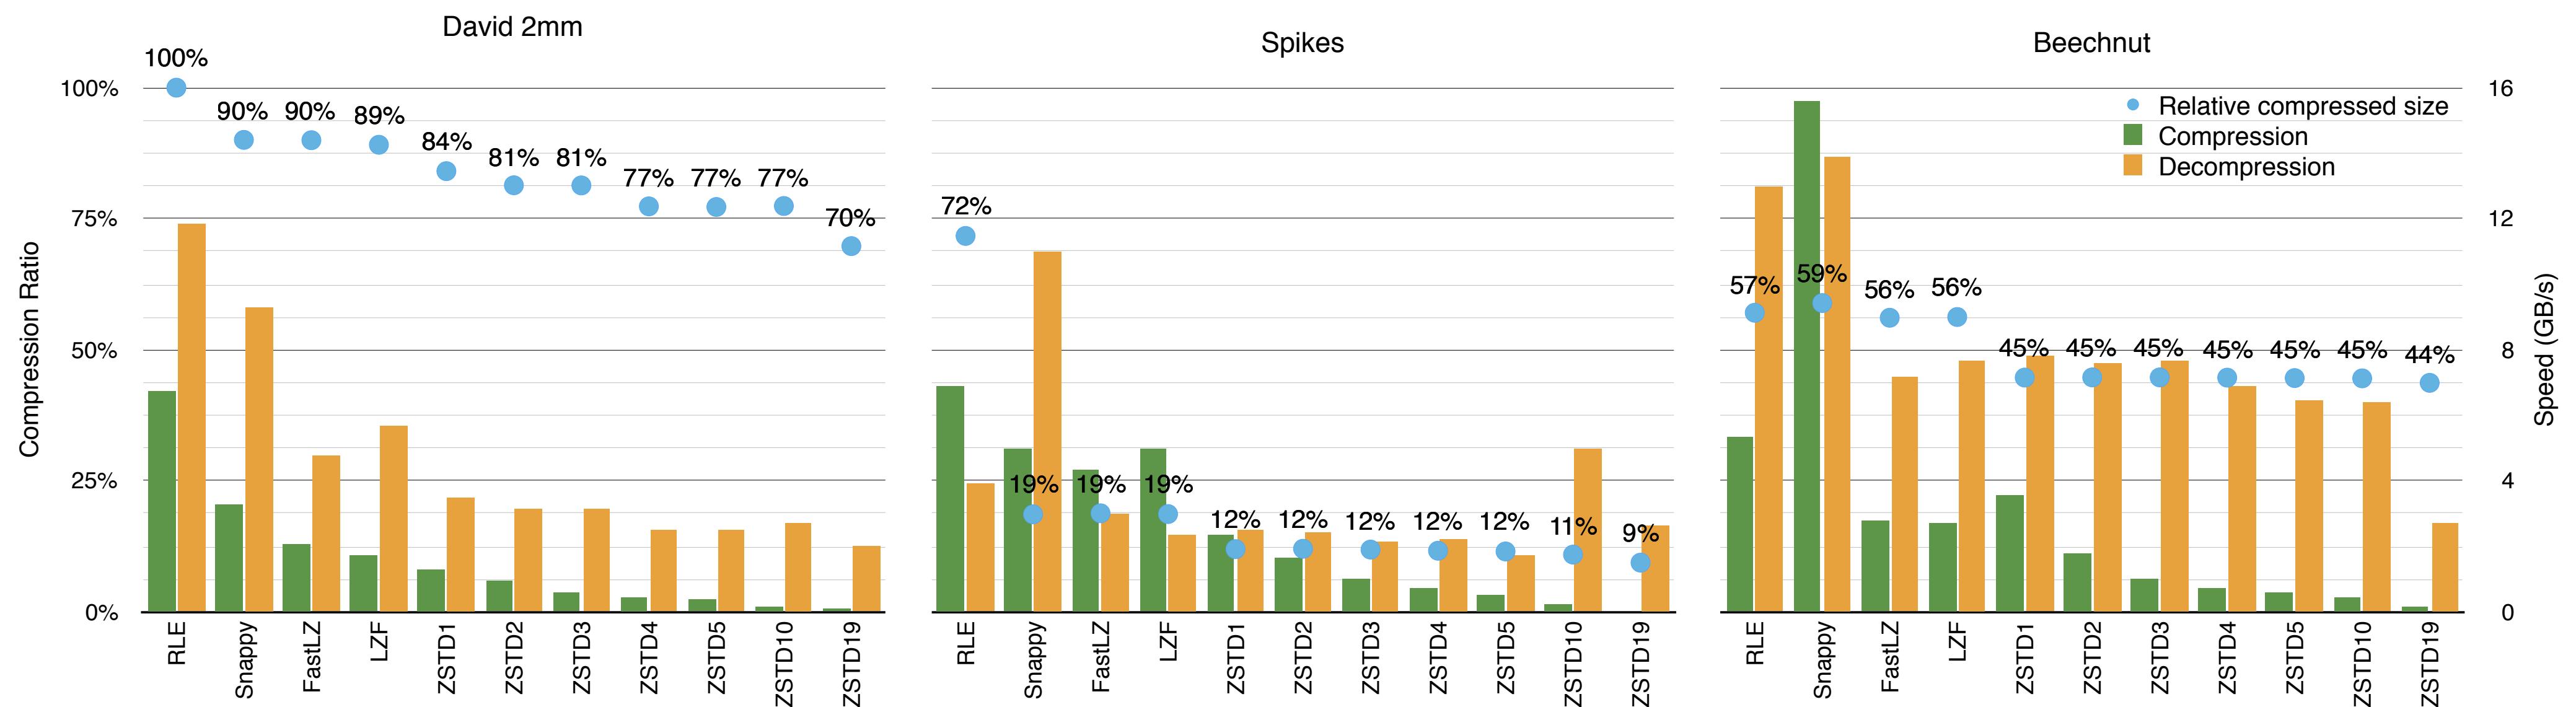
Reliable Stream Protocol



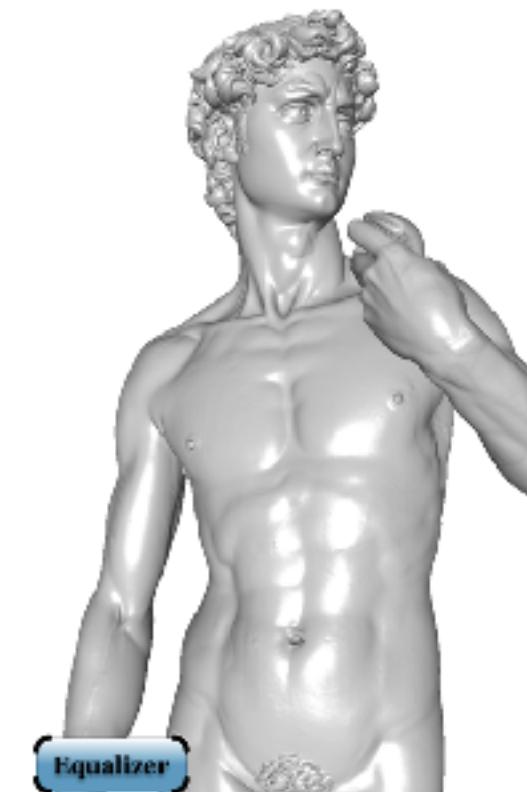
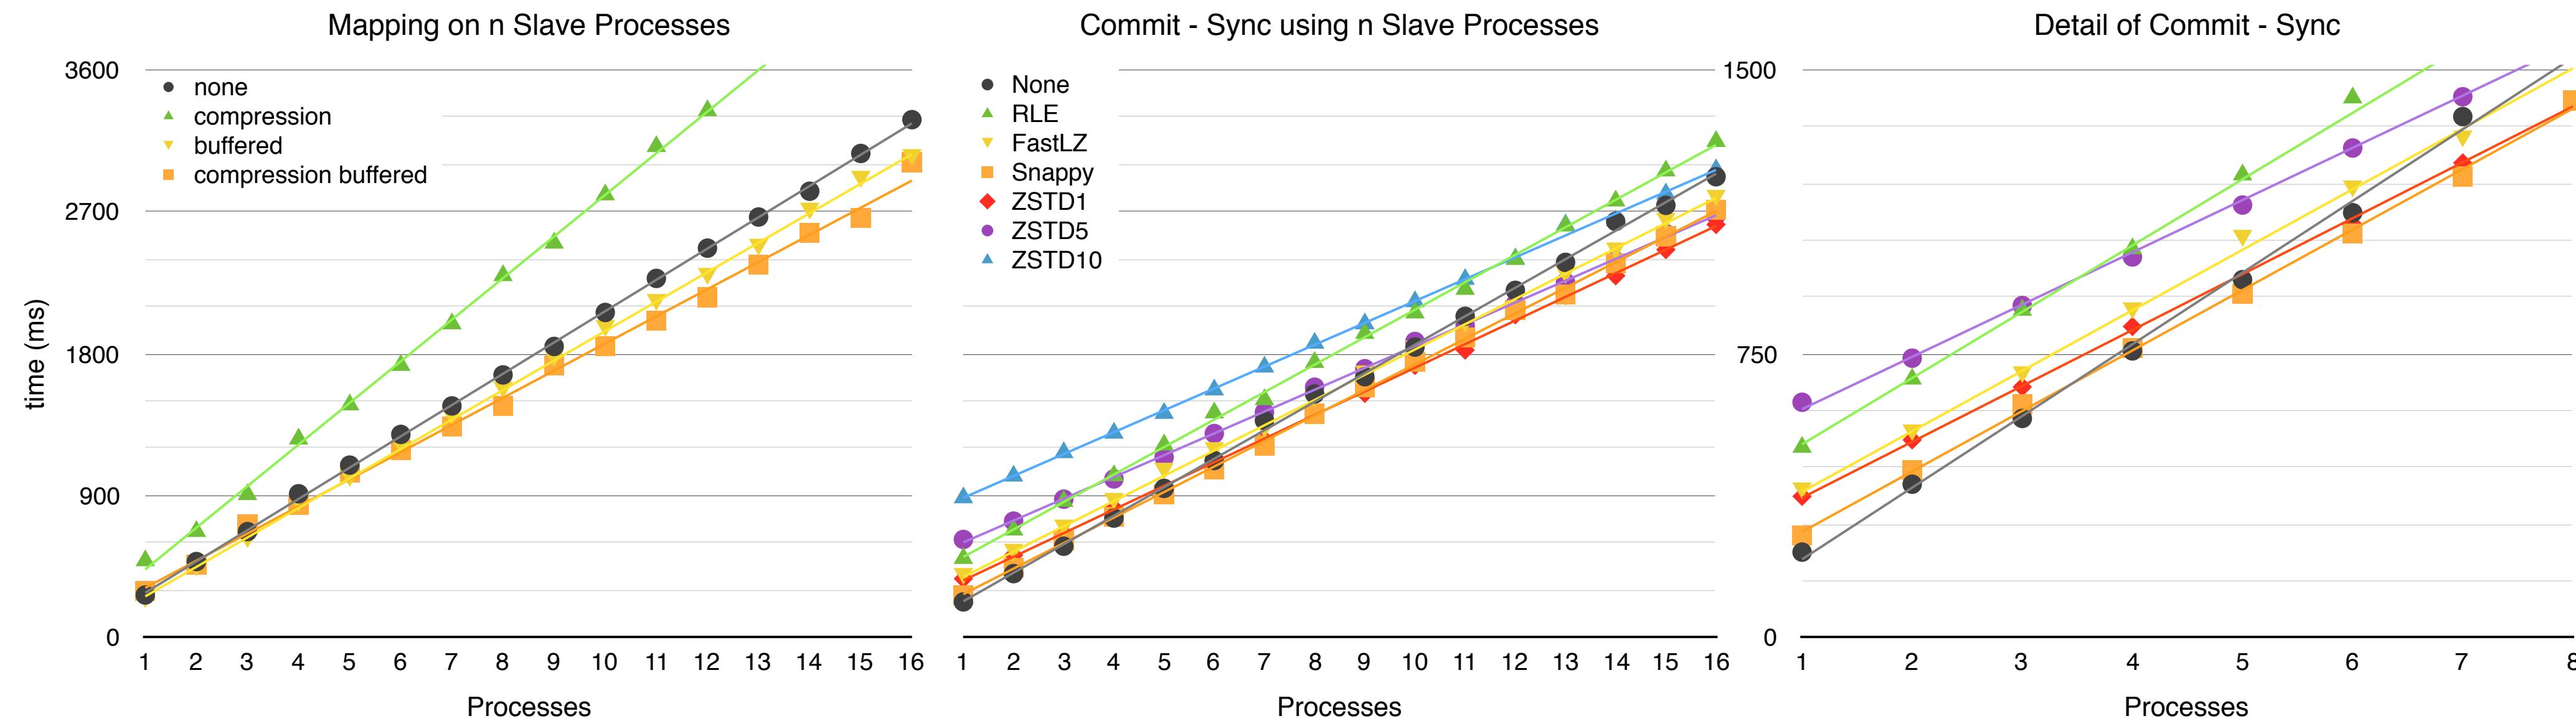
Reliable Stream Protocol



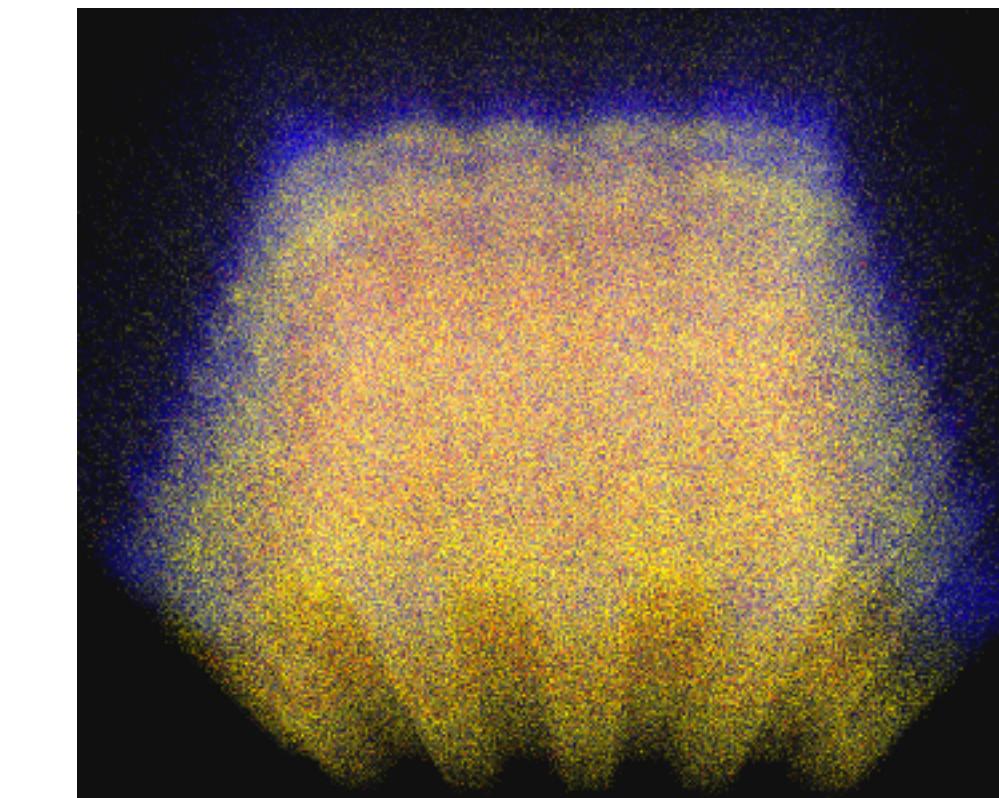
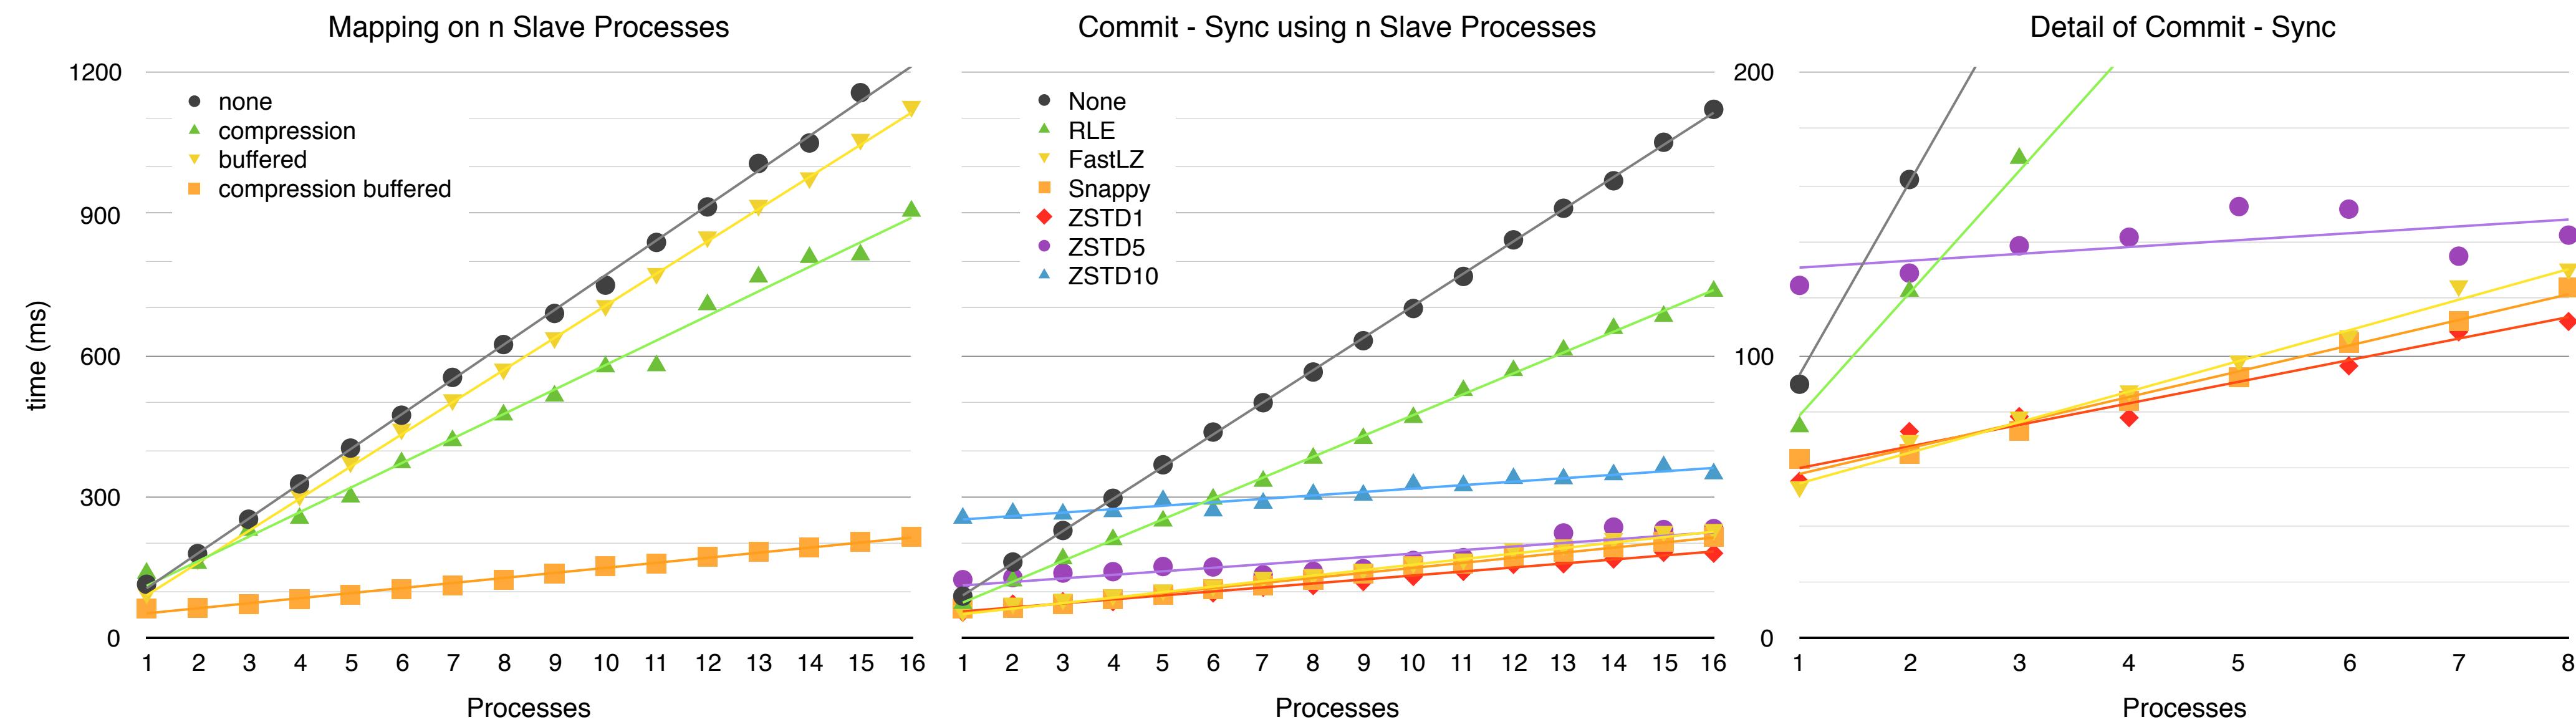
Compression



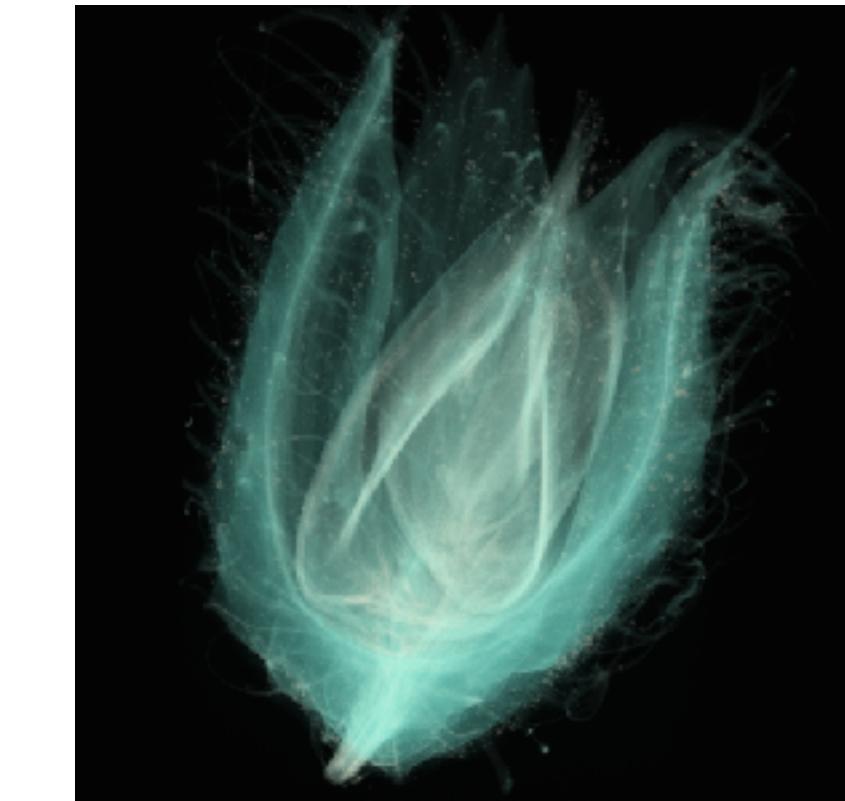
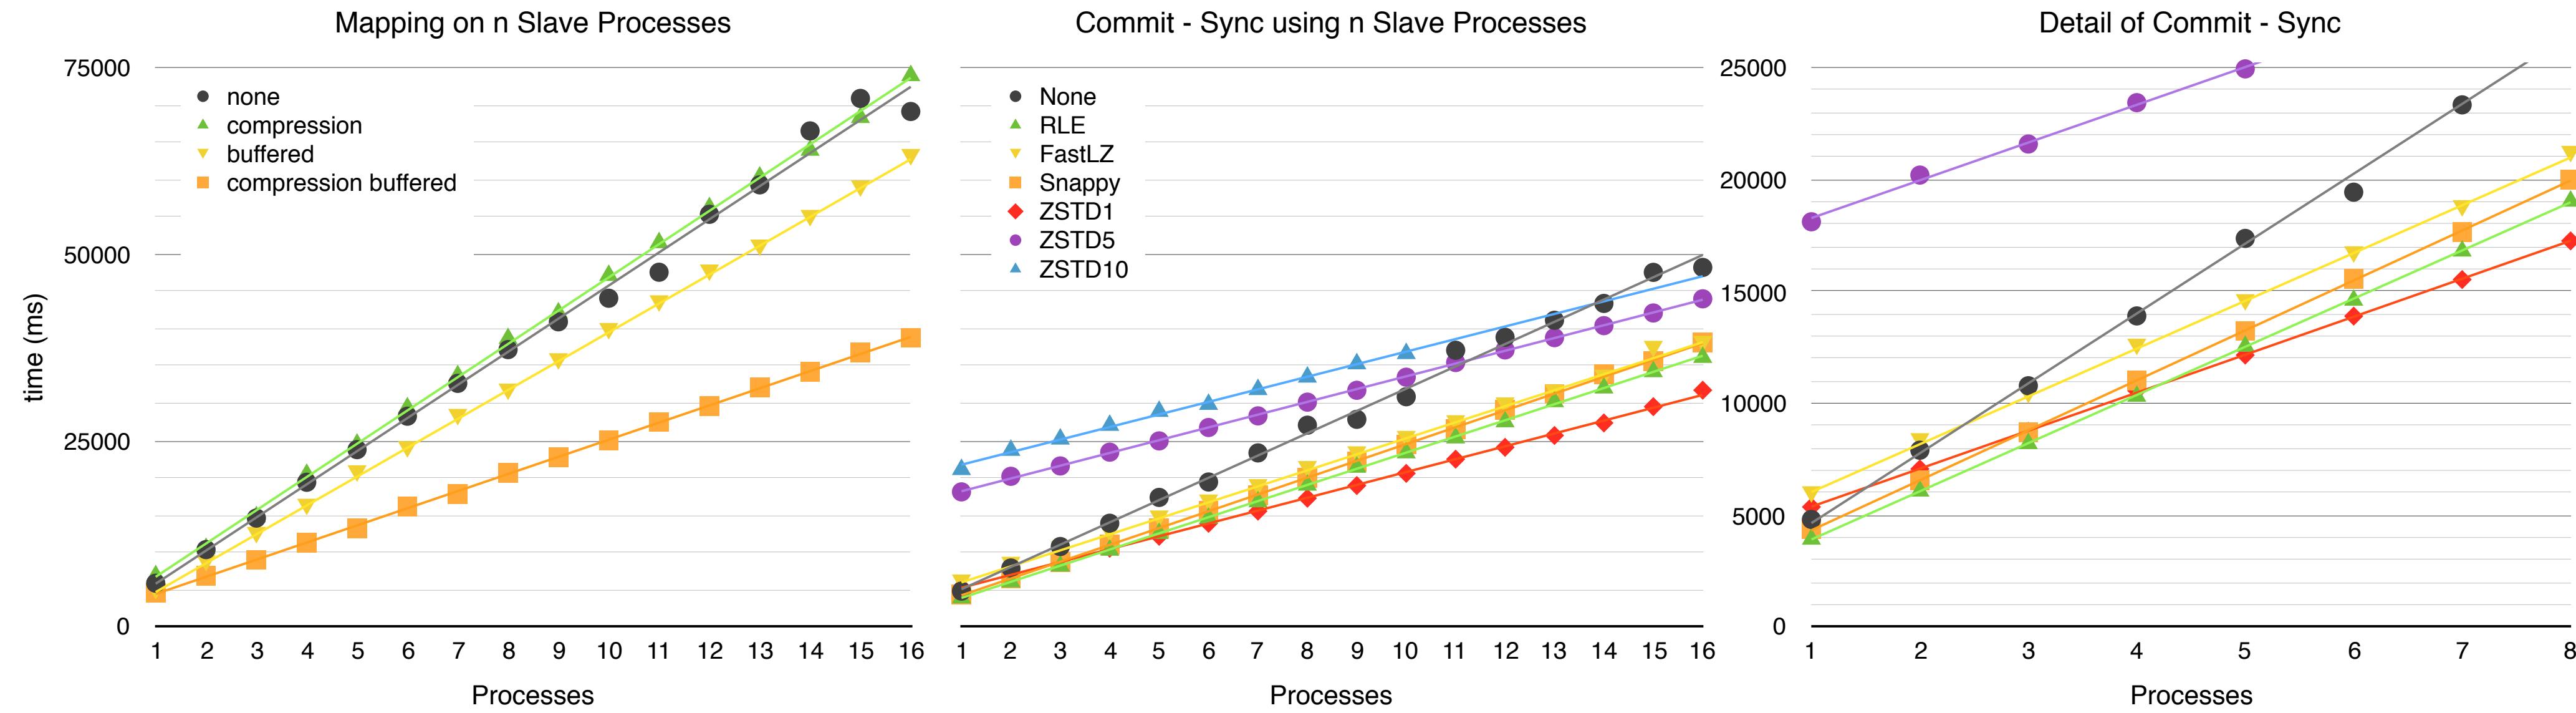
Distribution



Distribution



Distribution



Distribution

